

Computer Organization: A Programmer's Perspective

Basic Terms

Computer Science Department
Bar Ilan University, Israel



It's all just bits to me...

A view from the CPU

1. Get next instruction from memory
2. Execute it
3. Go to step 1

Why bits? Why 0/1?

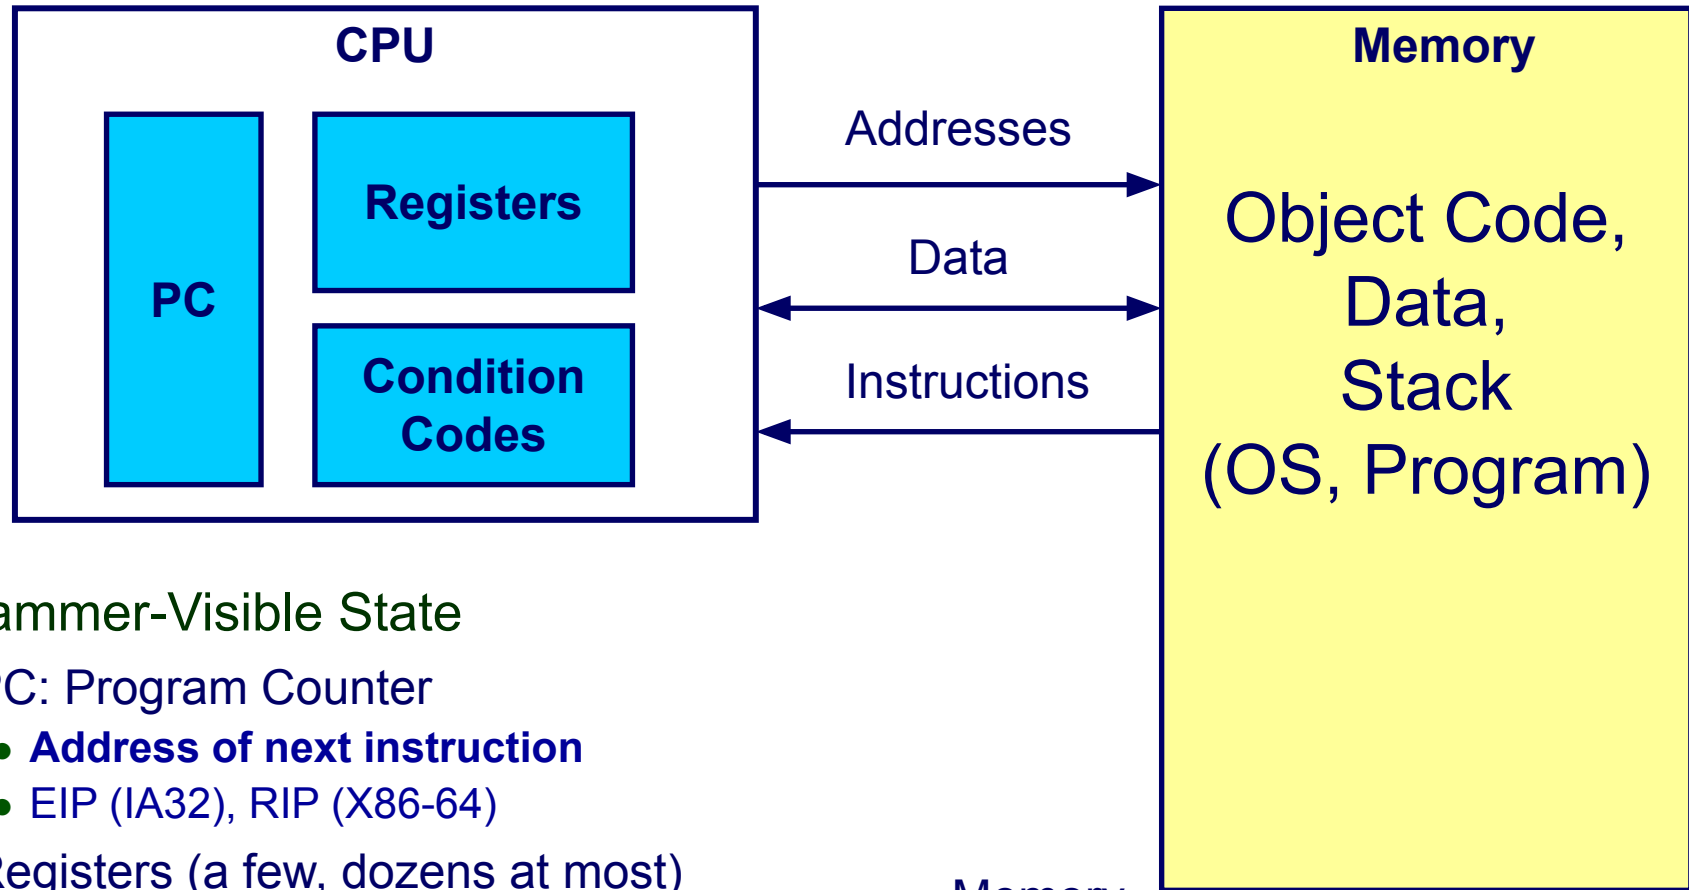
Basic terms: Bits, Bytes, Nibbles, Words

Bit-level manipulations

Boolean algebra

Expressing in C

Assembly Programmer's View of Computer



Programmer-Visible State

- **PC: Program Counter**
 - **Address of next instruction**
 - EIP (IA32), RIP (X86-64)
- **Registers (a few, dozens at most)**
 - Heavily used program data
- **Condition Codes**
 - Store status information about most recent arithmetic operation
 - Used for conditional branching
- **Memory**
 - Byte addressable array
 - Code, user data, OS data
 - Includes stack used to support procedures
 - **OS controls permissions**

Instruction Set Architecture

Assembly Language View

Processor state

Registers, memory, ...

Instructions

`addl, movl, leal, ...`

How instructions are encoded as bytes

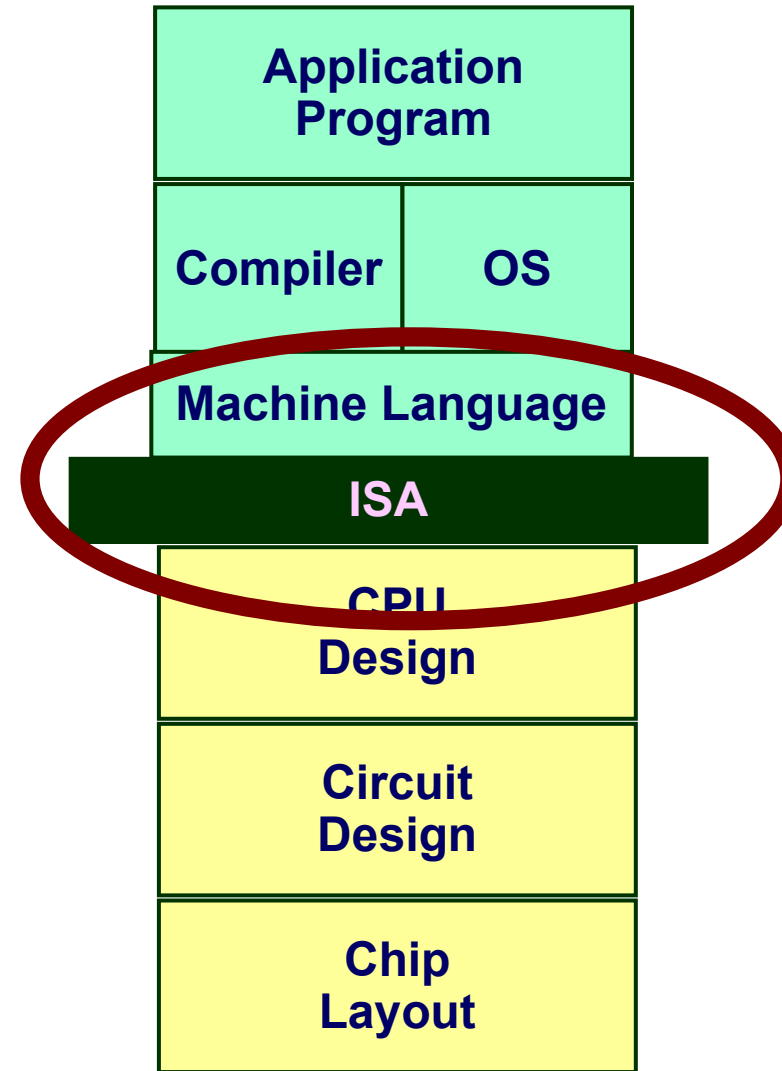
Layer of Abstraction

Above: how to program machine

Processor executes instructions in a sequence


Below: what needs to be built

Use variety of tricks to make it run fast
E.g., execute multiple instructions simultaneously



Everything is bits

- Each bit is 0 or 1
- By encoding/interpreting sets of bits in various ways
 - Computers represent numbers, sets, strings, etc...
 - Computers manipulate representations (instructions)
- Why bits? Electronic implementation is easy
 - Easy to store with bistable elements
 - Reliably transmitted on noisy and inaccurate wires



Binary Representations

- Numbers representation (base 2)
 - Represent 15213_{10} as 11101101101101_2
 - Represent 1.5213×10^4 as $1.1101101101101_2 \times 2^{13}$
 - Represent 1.20_{10} as $1.0011001100110011[0011]\dots_2$
- Characters representation
 - e.g., ASCII table assigns symbol to number
- Instruction representation
 - Example in AMD64 machine code
 - RETQ command: $C3$ (hex) = 11000011_2
 - MOV \$0, %EAX: $b8$ (hex) = $10111000_2 [00000000 \dots]$

Terms

- Bit: Single binary digit, 0 or 1
- Byte: 8 bits.
 - Smallest unit of memory used in modern computers
- Nibble (English: small bite): 4 bits
 - 2 nibbles = 1 byte
- Word: 8-64 bits (1 to 8 bytes)
 - Depends on machine!

Encoding Byte Values

Byte = 8 bits

- Binary 00000000_2 to 11111111_2
 - in C: “`x=0b10010000;`”

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Encoding Byte Values

Byte = 8 bits

- Binary 00000000_2 to 11111111_2
 - in C: “`x=0b10010000;`”
- Decimal: 0_{10} to 255_{10}

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Encoding Byte Values

Byte = 8 bits

- Binary 00000000_2 to 11111111_2
 - in C: “x=0b10010000;”
- Decimal: 0_{10} to 255_{10}
- Hexadecimal (base 16) 00_{16} to FF_{16}
 - Use characters ‘0’ to ‘9’ and ‘A’ to ‘F’
 - Two nibbles
 - in C: “x=0xFA1D37B”, or “x=0xfa1d37b”

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Encoding Byte Values

Byte = 8 bits

- Binary 00000000_2 to 11111111_2
 - in C: “x=0b10010000;”
- Decimal: 0_{10} to 255_{10}
- Hexadecimal (base 16) 00_{16} to FF_{16}
 - Use characters ‘0’ to ‘9’ and ‘A’ to ‘F’
 - Two nibbles
 - in C: “x=0xFA1D37B”, or “x=0xfa1d37b”
- Octal (base 8): 0_8 to 377_8
 - in C as ‘x=0256’ (0 zero)
 - 3 bits

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Encoding Byte Values

Byte = 8 bits, nibble = 4 bits

- Binary 00000000_2 to 11111111_2
 - in C: “`x=0b10010000;`”
- Decimal: 0_{10} to 255_{10}
- Hexadecimal (base 16) 00_{16} to FF_{16}
 - Use characters ‘0’ to ‘9’ and ‘A’ to ‘F’
 - Two nibbles
 - in C: “`x=0xFA1D37B`”, or “`x=0xfa1d37b`”
- Octal (base 8): 0_8 to 377_8
 - in C as ‘`x=0256`’ (0 zero)
 - 3 bits

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Bit level operations

Boolean Algebra

(George Boole, 19th century)

- **Developed by George Boole in 19th Century**

- Algebraic representation of logic
- Encode “True” as 1 and “False” as 0

And

- **$A \& B = 1$ when both $A=1$ and $B=1$**

$\&$	0	1
0	0	0
1	0	1

Or

- **$A | B = 1$ when either $A=1$ or $B=1$**

$ $	0	1
0	0	1
1	1	1

Not

- **$\sim A = 1$ when $A=0$**

\sim	
0	1
1	0

Exclusive-Or (Xor)

- **$A \wedge B = 1$ when either $A=1$ or $B=1$, but not both**

\wedge	0	1
0	0	1
1	1	0

Operators applied bitwise

$$\begin{array}{r} 01101001 \\ \& 01010101 \\ \hline 01000001 \end{array}$$

$$\begin{array}{r} 01101001 \\ | 01010101 \\ \hline 01111101 \end{array}$$

$$\begin{array}{r} 01101001 \\ \wedge 01010101 \\ \hline 00111100 \end{array}$$

$$\begin{array}{r} \sim 01010101 \\ \hline 10101010 \end{array}$$

Example: Representing & Manipulating Sets

■ Representation

- Width w bit vector represents subsets of $\{0, \dots, w-1\}$
- $a_j = 1$ if $j \in A$
 - 01101001 { 0, 3, 5, 6 }
 - 76543210
 - 01010101 { 0, 2, 4, 6 }
 - 76543210

■ Operations

- & Intersection 01000001 { 0, 6 }
- | Union 01111101 { 0, 2, 3, 4, 5, 6 }
- ^ Symmetric difference 00111100 { 2, 3, 4, 5 }
- ~ Complement 10101010 { 1, 3, 5, 7 }

Bit-Level Operations in C

■ Operations $\&$, $|$, \sim , \wedge Available in C

- Apply to any “integral” data type
 - long, int, short, char, unsigned
- View arguments as bit vectors
- Arguments applied bit-wise

■ Examples (Char data type)

- $\sim 0x41 \rightarrow 0xBE$
 - $\sim 01000001_2 \rightarrow 10111110_2$
- $\sim 0x00 \rightarrow 0xFF$
 - $\sim 00000000_2 \rightarrow 11111111_2$
- $0x69 \& 0x55 \rightarrow 0x41$
 - $01101001_2 \& 01010101_2 \rightarrow 01000001_2$
- $0x69 | 0x55 \rightarrow 0x7D$
 - $01101001_2 | 01010101_2 \rightarrow 01111101_2$

Contrast: Logic Operations in C

■ Contrast to Logical Operators

- `&&`, `||`, `!`
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1
 - **Early termination**

■ Examples (char data type)

- `!0x41 → 0x00`
- `!0x00 → 0x01`
- `!!0x41 → 0x01`

- `0x69 && 0x55 → 0x01`
- `0x69 || 0x55 → 0x01`
- `p && *p` (avoids null pointer access)

Shift Operations

- **Left Shift: $x \ll y$**
 - Shift bit-vector x left y positions
 - Throw away extra bits on left
 - Fill with 0's on right
- **Right Shift: $x \gg y$**
 - Shift bit-vector x right y positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- **Undefined Behavior**
 - Shift amount < 0 or \geq word size

Argument x	01100010
$\ll 3$	00010000
Log. $\gg 2$	00011000
Arith. $\gg 2$	00011000

Argument x	10100010
$\ll 3$	00010000
Log. $\gg 2$	00101000
Arith. $\gg 2$	11101000

Cool Stuff with Xor: SWAP

- Bitwise Xor is form of addition
- With extra property that every value is its own additive inverse

$$A \oplus A = 0$$

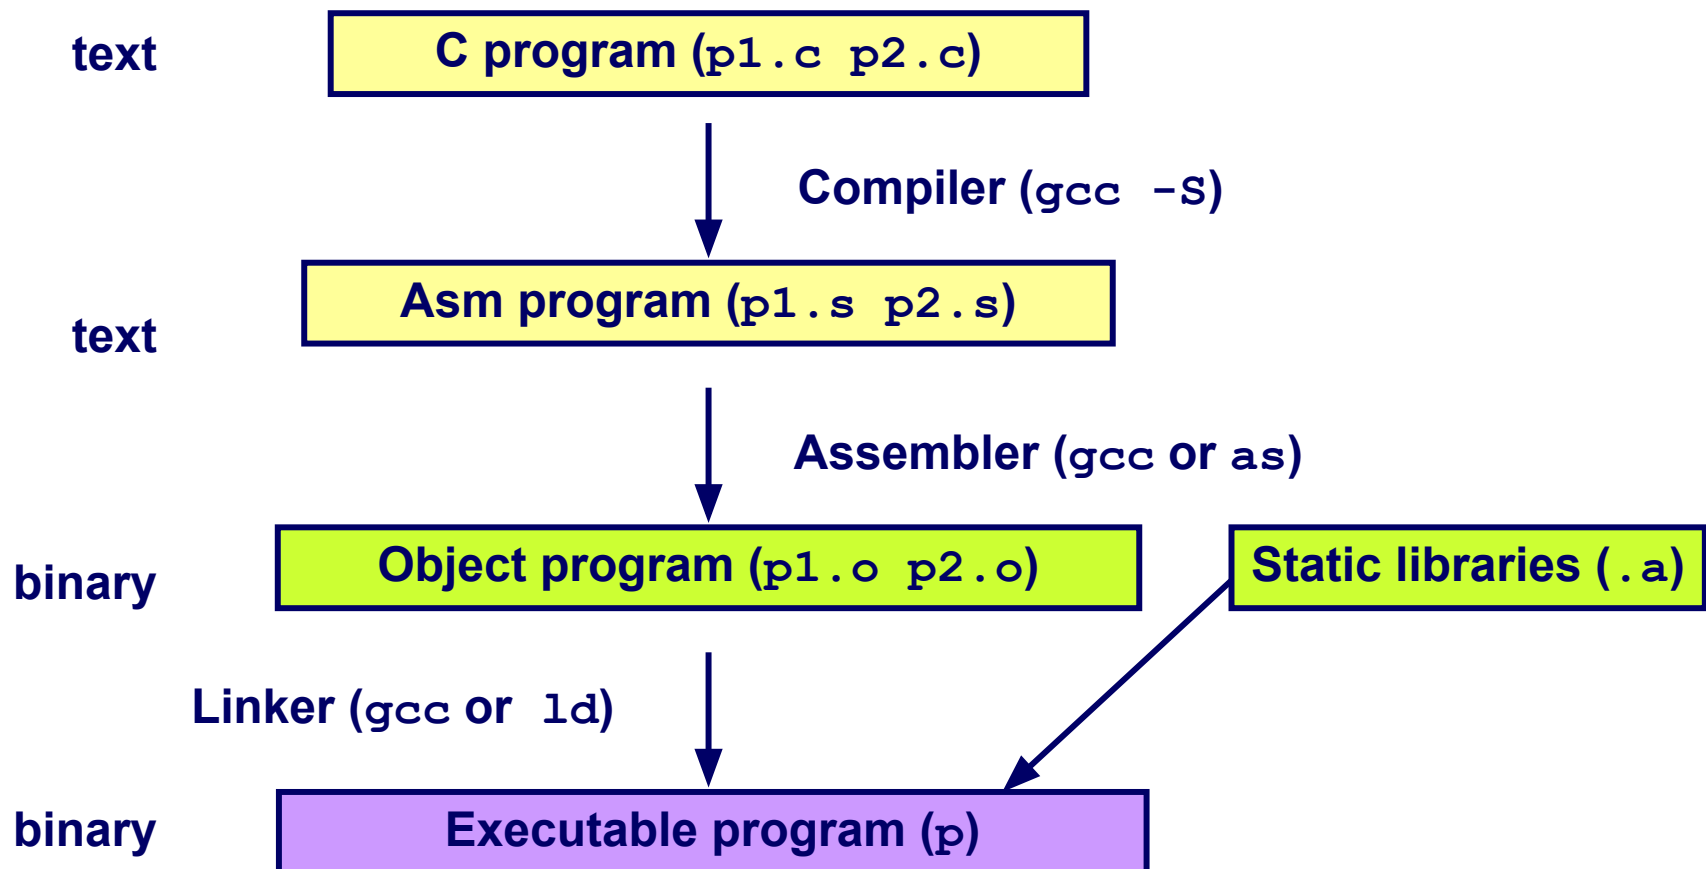
```
void funny(int *x, int *y)
{
    *x = *x ^ *y;    /* #1 */
    *y = *x ^ *y;    /* #2 */
    *x = *x ^ *y;    /* #3 */
}
```

	*x	*y
Begin	A	B
1	$A \oplus B$	B
2	$A \oplus B$	$(A \oplus B) \oplus B = A$
3	$(A \oplus B) \oplus A = B$	A
End	B	A

Instructions as Bits

Turning C into Object Code

- Code in files `p1.c p2.c`
- Compile with command: `gcc -O p1.c p2.c -o p`
 - Use optimizations (`-O`)
 - Put resulting binary in file `p`



Compiling Into Assembly

C Code (sum.c)

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

AMD64 assembly (optimized)

```
_sum:
    leal    (%rdi,%rsi), %eax
    ret
```

gcc -O1 -S sum.c

IA32 assembly (non-optimized)

```
_sum:
    pushl   %ebp
    movl    %esp,%ebp
    movl    12(%ebp),%eax
    addl    8(%ebp),%eax
    movl    %ebp,%esp
    popl    %ebp
    ret
```

gcc -m32 -O0 -S sum.c

IA32 assembly (optimized)

```
_sum:
    movl    4(%esp),%eax
    addl    8(%esp),%eax
    ret
```

gcc -m32 -O1 -S sum.c

Machine-Level Code Representation

- Encode Program as Sequence of Instructions
 - Arithmetic, logical, math operations
 - Read or write memory
 - Conditional branches, jumps
- Different machines, different instructions
 - Code not binary compatible (in general)
 - Different CPU “families” do not agree
 - PowerPC, ARM (tablets, phones) use fix-length instructions
 - PC’s (AMD, Intel) use variable length instructions
 - Follow different design approaches (RISC vs CISC)

Representing Instructions

```
int sum(int x, int y)
{
    return x+y;
}
```

- For this example, Alpha & Sun use two 4-byte instructions
 - Use differing numbers of instructions in other cases
- PC uses 7 instructions with lengths 1, 2, and 3 bytes

Alpha

00
00
30
42
01
80
FA
6B

Sun

81
C3
E0
08
90
02
00
09

Intel 32

55
89
E5
8B
45
0C
03
45
08
89
EC
5D
C3

Different CPUs use totally different instructions and encodings

Storing bits (data)

Machines have Words

- Imprecise definitions
 - Nominal size of integer-valued register
 - Sometimes size of address, memory bus width
- Current desktop machines are 64 bits (8 bytes)
 - Potentially address $\approx 1.8 \times 10^{19}$ bytes
 - 32-bit machines phasing out (but in phones, tablets)
- Low-end use 8- or 16-bit words
- Machines support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

Byte-Oriented Memory Organization

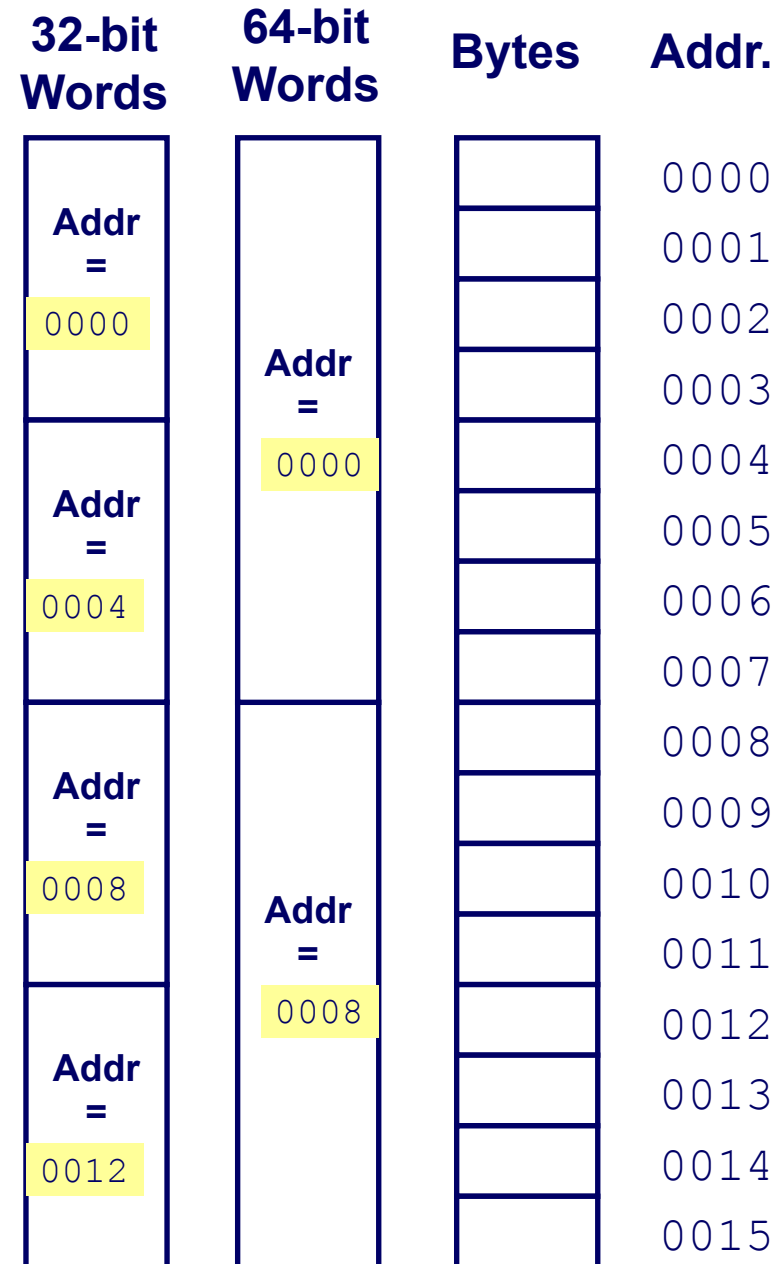
- Programs Refer to Virtual Addresses
 - Conceptually very large array of bytes
 - Implemented with hierarchy of different memory types
 - SRAM, DRAM, disk
 - In Unix and Windows NT, address space private to “process”
 - Program can clobber its own data, but not that of others
 - You will see this again, in much more detail
- Compiler + Run-Time System Control Allocation
 - Where different program objects should be stored
 - Multiple mechanisms: static, stack, and heap
 - In any case, all allocation within single virtual address space
 - You will see this again, in much more detail

Word-Oriented Memory Organization

Addresses:

Specify Byte Locations

- Address of first byte in word
- Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



Byte Ordering

How should bytes within multi-byte word be ordered in memory?

Conventions

- Sun's, Mac's are "Big Endian" machines
 - Least significant byte has highest address
- ARM, Intel are "Little Endian" machines
 - Least significant byte has lowest address
 - Some ARM CPUs have "big endian" mode

Byte Ordering Example

Big Endian

- Least significant byte has highest address


Little Endian

- Least significant byte has lowest address


Example

- Variable `x` has 4-byte representation `0x01234567`
- Address given by `&x` is `0x100`

Big Endian



Little Endian



Reading Byte-Reversed Listings

Disassembly

- Text representation of binary machine code
- Generated by program that reads the machine code

Example Fragment

Address	Instruction	Code	Assembly	Rendition
8048365:	5b		pop	%ebx
8048366:	81 c3	ab 12 00 00	add	\$0x12ab, %ebx
804836c:	83 bb	28 00 00 00 00	cmpl	\$0x0, 0x28(%ebx)

Deciphering Numbers

- Value: 0x12ab
- Pad to 4 bytes: 0x000012ab
- Split into bytes: 00 00 12 ab
- Reverse: ab 12 00 00

Examining Data Representations

Code to Print Byte Representation of Data

- Casting pointer to `unsigned char *` creates byte array

```
typedef unsigned char *pointer;

void show_bytes(pointer start, int len)
{
    int i;
    for (i = 0; i < len; i++)
        printf("0x%p\t0x%.2x\n",
               start+i, start[i]);
    printf("\n");
}
```

Printf directives:

%p: Print pointer

%x: Print Hexadecimal

show_bytes Execution Example

```
int a = 15213;  
printf("int a = 15213;\n");  
show_bytes((pointer) &a, sizeof(int));
```




Result (on 32bit, little endian machine):

```
int a = 15213;  
0x11ffffcb8    0x6d  
0x11ffffcb9    0x3b  
0x11ffffcba    0x00  
0x11ffffcbb    0x00
```

Representing Integers

```
int A = 15213;
int B = -15213;
long int C = 15213;
```

Decimal:	15213			
Binary:	0011	1011	0110	1101
Hex:	3	B	6	D



Two's complement representation

Representing Pointers

```
int B = -15213;
int *P = &B;
```

Alpha Address

Hex: ... 0 1 F F F F F C A 0

Binary: 0000 0001 1111 1111 1111 1111 1111 1100 1010 0000

Sun Address

Hex: E F F F F B 2 C

Binary: 1110 1111 1111 1111 1111 1011 0010 1100

Linux Address

Hex: B F F F F 8 D 4

Binary: 1011 1111 1111 1111 1111 1000 1101 0100

Alpha P

A0

FC

FF

FF

01

00

00

00

Sun P

EF

FF

FB

2C

Linux P

D4

F8

FF

BF

Different compilers & machines assign different locations to objects

Main Points

- Boolean Algebra is Mathematical Basis
 - Basic form encodes “false” as 0, “true” as 1
 - General form like bit-level operations in C
 - Good for representing & manipulating sets
- It's All About Bits & Bytes
 - Numbers, text, programs
- Different Machines Follow Different Conventions
 - Representations
 - Word size
 - Byte ordering