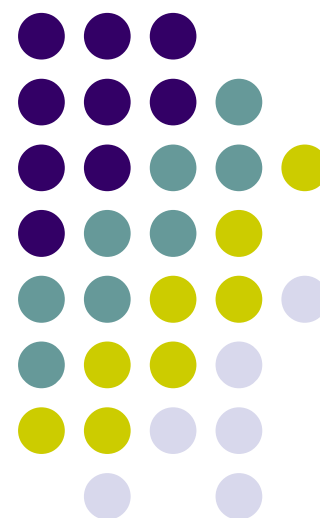



Computer Organization: A Programmer's Perspective

Machine-Level Programming (1: Introduction)

Gal A. Kaminka
galk@cs.biu.ac.il



Instruction Set Architecture



Assembly Language View

Processor state

Registers, memory, ...

Instructions

`addl, movl, leal, ...`

How instructions are encoded as bytes

Layer of Abstraction

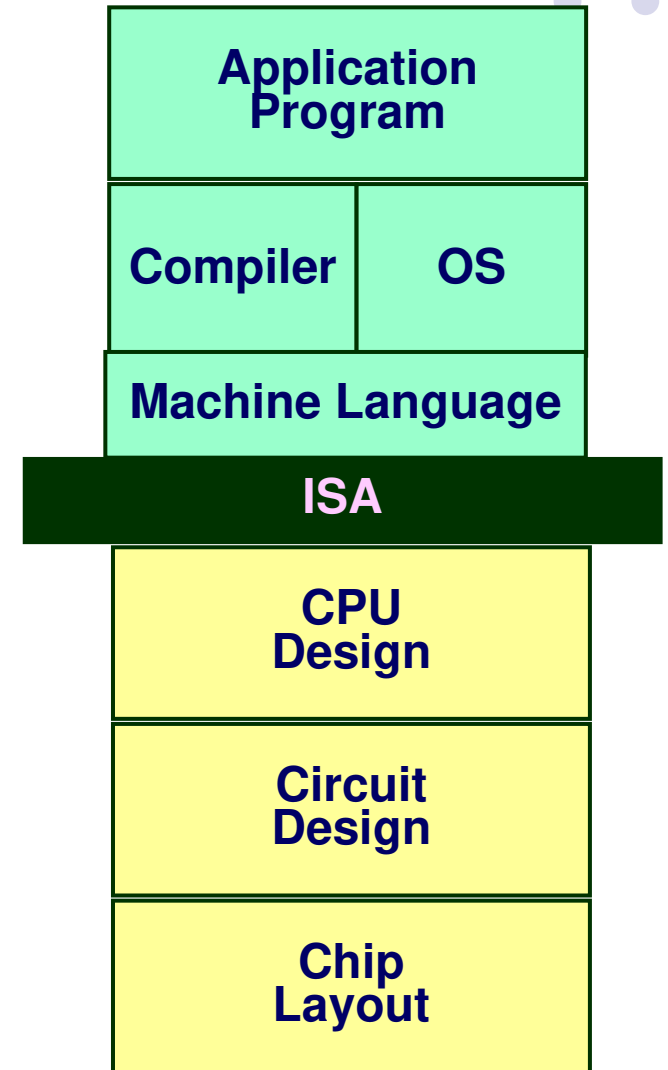
Above: how to program machine

Processor executes instructions in a sequence


Below: what needs to be built

Use variety of tricks to make it run fast

E.g., execute multiple instructions simultaneously



Instruction Set Architecture



Assembly Language View

Processor state

Registers, memory, ...

Instructions

`addl, movl, leal, ...`

How instructions are encoded as bytes

Layer of Abstraction

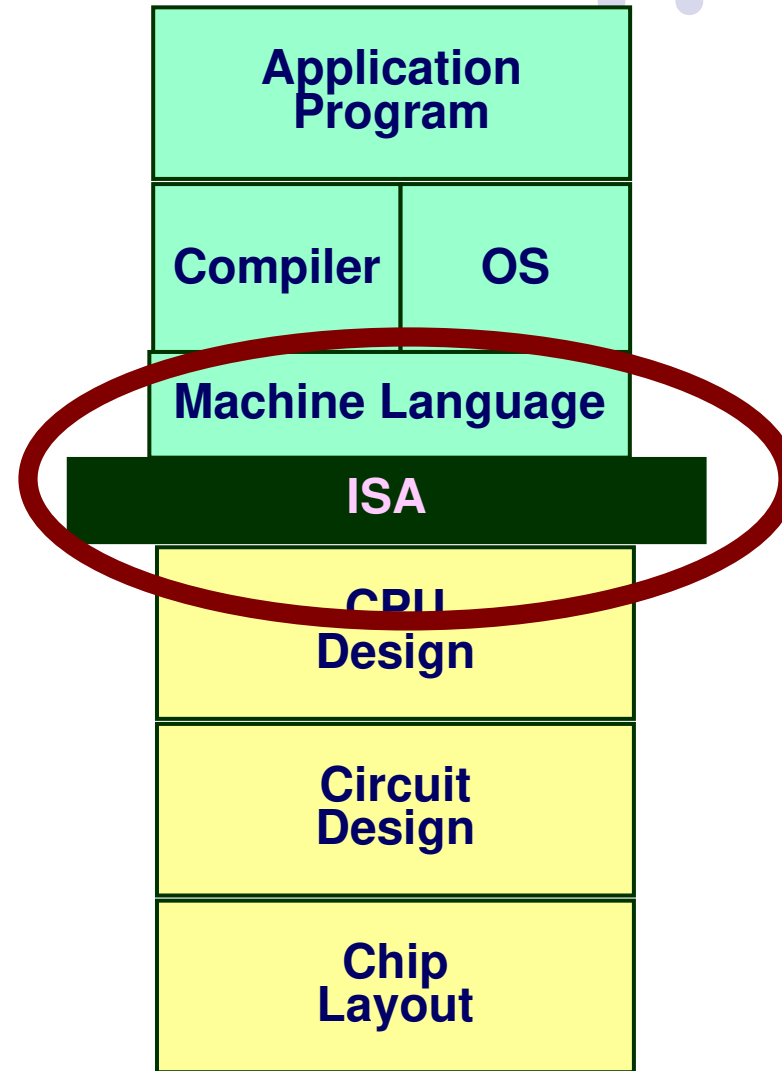
Above: how to program machine

Processor executes instructions in a sequence

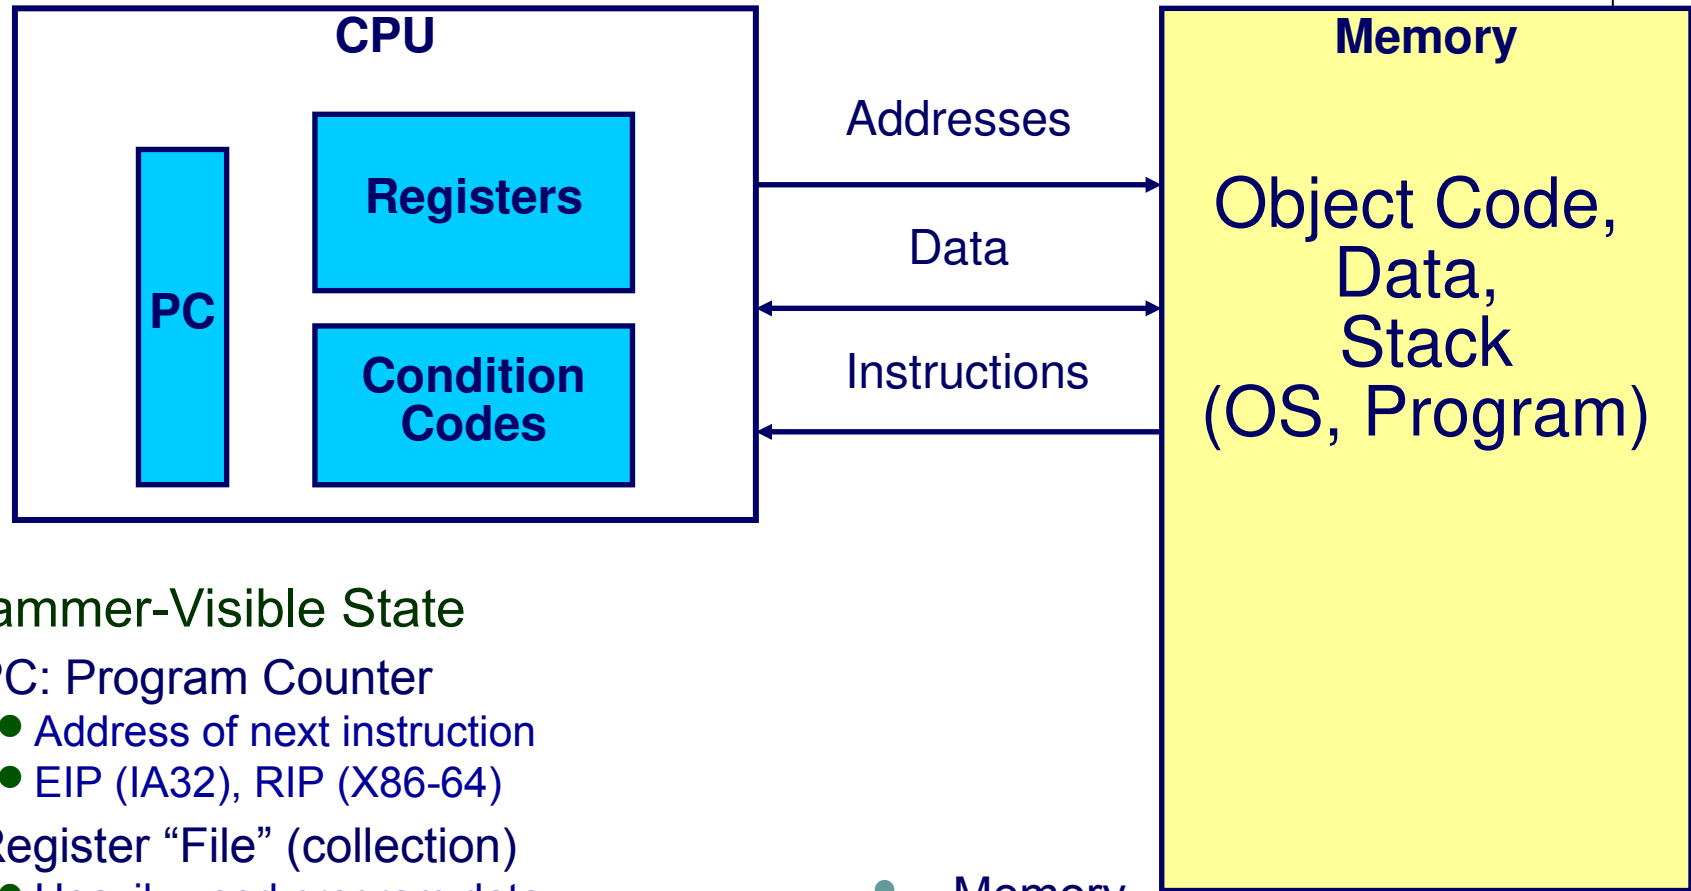

Below: what needs to be built

Use variety of tricks to make it run fast

E.g., execute multiple instructions simultaneously



Assembly Programmer's View

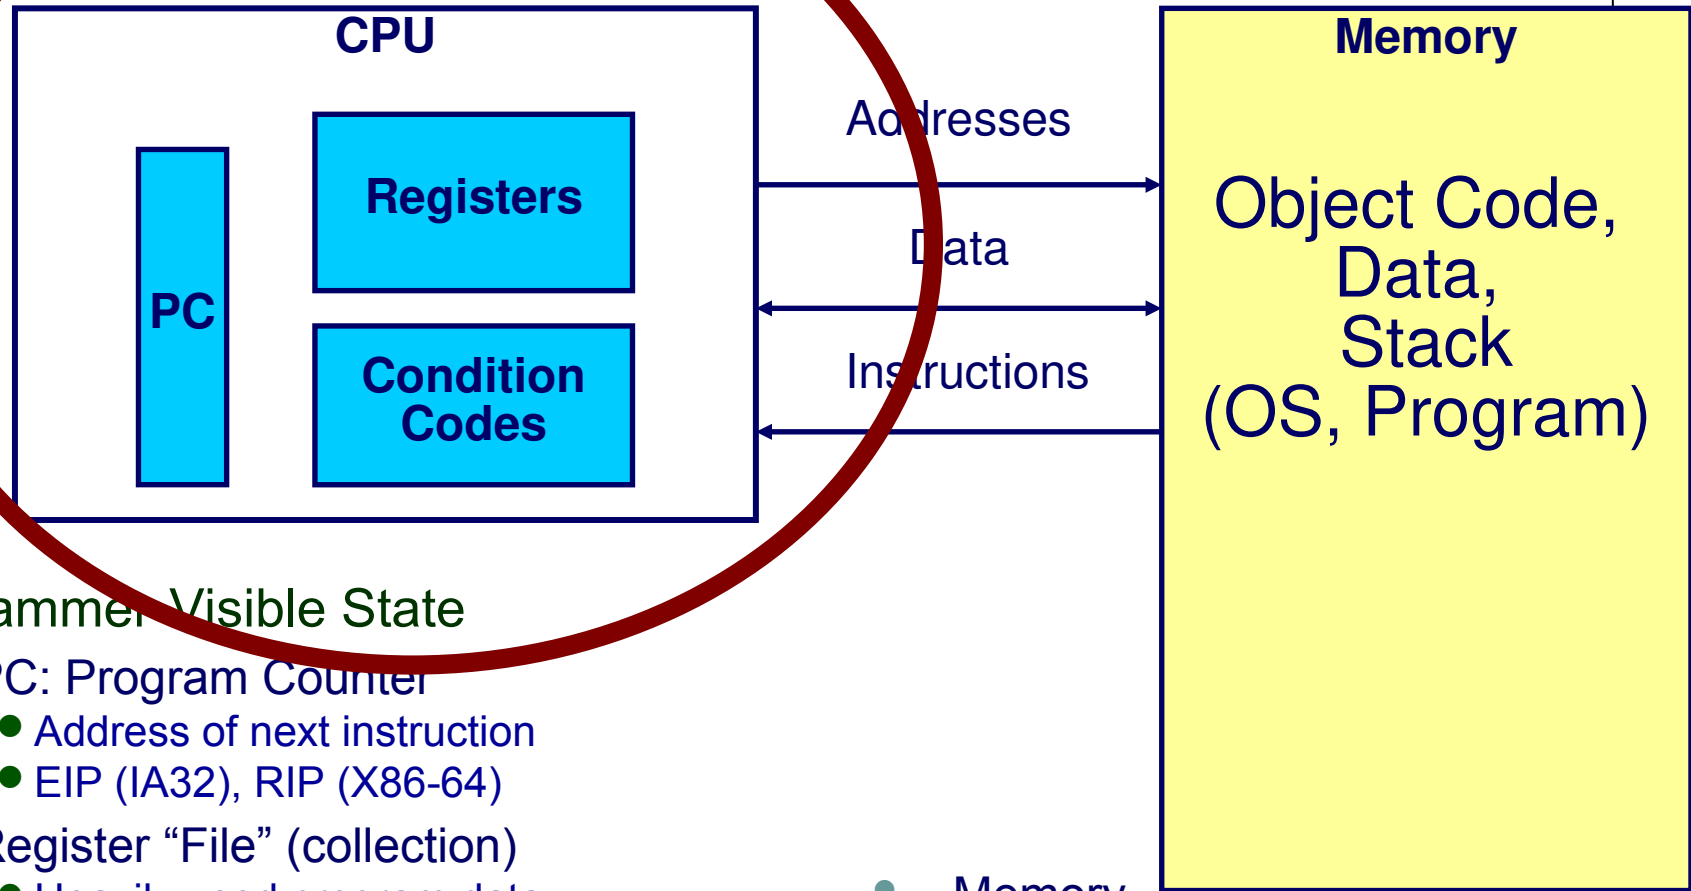



Programmer-Visible State

- PC: Program Counter
 - Address of next instruction
 - EIP (IA32), RIP (X86-64)
- Register "File" (collection)
 - Heavily used program data
- Condition Codes
 - Store status information about most recent arithmetic operation
 - Used for conditional branching

- Memory
 - Byte addressable array
 - Code, user data, OS data
 - Includes stack used to support procedures
 - OS controls permissions

Assembly Programmer's View

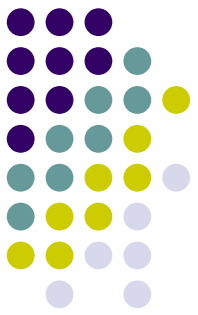


Programmer Visible State

- **PC: Program Counter**
 - Address of next instruction
 - EIP (IA32), RIP (X86-64)
- **Register "File" (collection)**
 - Heavily used program data
- **Condition Codes**
 - Store status information about most recent arithmetic operation
 - Used for conditional branching

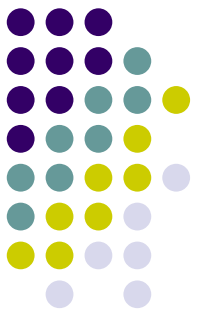
- **Memory**
 - Byte addressable array
 - Code, user data, OS data
 - Includes stack used to support procedures
 - **OS controls permissions**

Abstract View of CPU



```
void cpu(void) {  
    int iregs[ ] ; // word registers  
    bit flags[ ] ; // 1-bit flags: carry, zero, overflow, ...  
    float fregs[ ] ; // floating point – not all CPUs  
  
    address register IP = FIRST_ADDRESS  
    while (true) {  
        (instruction) ← get_next_instruction(IP)  
        (iregs, fregs, flags, IP) ← execute(instruction,iregs,fregs,flags)  
    }  
}
```

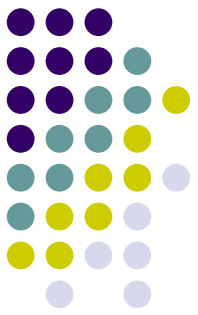
Abstract View of CPU



```
void cpu(void) {  
    int iregs[ ] ; // word registers  
    bit flags[ ] ; // 1-bit flags: carry, zero, overflow, ...  
    float fregs[ ] ; // floating point – not all CPUs  
  
    address register IP = FIRST_ADDRESS  
    while (true) {  
        (instruction) ← get_next_instruction(IP)  
        (iregs, fregs, flags, IP) ← execute(instruction, iregs, fregs, flags)  
    }  
}
```

ISA changes available registers, available instructions

Compiling Into Assembly



C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

Obtain with command

```
gcc -O -S sum.c
```

Produces file `sum.s`

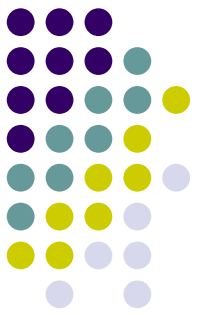
Generated Assembly

```
_sum:
    pushl %ebp
    movl %esp,%ebp
    movl 12(%ebp),%eax
    addl 8(%ebp),%eax
    movl %ebp,%esp
    popl %ebp
    ret
```



IA32 assembly!

Compiling Into Assembly



C Code

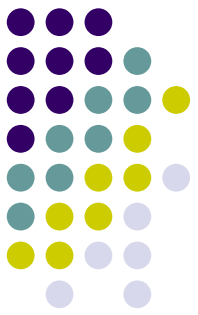
```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

Generated Assembly

```
_sum:
    pushl %ebp
    movl %esp, %ebp
    movl 12(%ebp), %eax
    addl 8(%ebp), %eax
    movl %ebp, %esp
    popl %ebp
    ret
```

registers: ebp, esp, eax

Compiling Into Assembly



C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

Generated Assembly

```
_sum:
    pushl %ebp
    movl %esp,%ebp
    movl 12(%ebp),%eax
    addl 8(%ebp),%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

instructions: `movl, addl, pushl, popl, ret`

Compiling Into Assembly

C Code (sum.c)

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

Generated Assembly

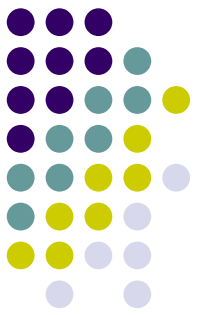
```
sum:
    leal    (%rdi,%rsi), %eax
    ret
```

Different ISA

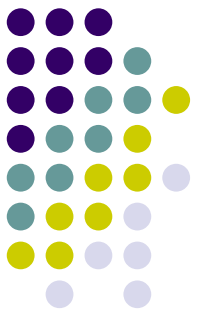
- Different registers (rdi, rsi)
- Some commands different

AMD64/x86-64
assembly!

ISA vs processor (CPU)



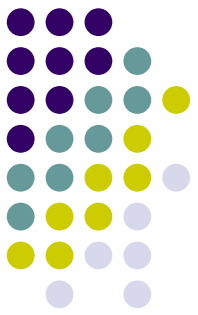
- ISA defines the expected behavior and capabilities
 - How many registers? What type? What size?
 - What instructions are available?
 - What addressing modes (memory access) available?
 - What condition codes available? How can they be checked?
 - ...
- CPUs implement ISA
 - For example, Intel and AMD compete on same ISA implementation
 - The competition includes other companies, too
 - CPUs may offer better speed on some instructions or on all
 - May use less or more energy
 - Be more robust to radiation, temperature, etc.



X86 ISA Evolution

| Name | Date | Word | # General Regs |
|--|------|---------|-----------------------------------|
| X86-16 | 1978 | 16 bits | 4 (+4 limited) |
| <ul style="list-style-type: none">■ e.g., Intel 8086 16-bit processor. Basis for IBM PC & DOS<ul style="list-style-type: none">● Limited to 1MB address space. DOS only gives you 640K■ CPUs produced by Intel AMD, VIA Technologies, NEC■ 8088, 80186, 80286, ... all extending. # of general registers same■ Registers: AX, BX, CX, DX, ... | | | |
| IA32(1) | 1985 | 32 bits | 8 (+6 limited) |
| <ul style="list-style-type: none">■ Intel 386, AMD Am386■ Extended to 32 bits. Added “flat addressing”, modern OS support: capable of running Unix | | | |
| IA32(2) | 1989 | 32 bits | 8 (+6 limited) +8 FP (80 bits) |
| <ul style="list-style-type: none">■ Floating Point in hardware!■ Intel 486, but also by IBM, AMD, TI, ... Combined 386+387 co-processor | | | |


X86 ISA Evolution



| Name | Date | Word | # General Regs |
|--------------|-------|------|---|
| IA32(3) | 1997+ | 32 | 8 (+6 limited) + 8 FP (80 bits) + 8 64bit integers (XMM) <ul style="list-style-type: none">• SSE: SIMD instructions for operating on 64-bit vectors of 1, 2, or 4 byte integer data• Later: SSE2, SSE3, ... and other instruction extensions |
| X86-64/AMD64 | 2000 | 64 | 16 (+ 6 limited) + 16 128bit (XMM) + 8 FP, ... <ul style="list-style-type: none">• Capable of running 16/32 bit binaries of older architecture (backwards compatible)• CPUs by AMD, Intel, VIA |

Later additions: multiple cores (2006-), virtualization (2007-), AES, ...

2015 State of the Art: Core i7 Broadwell



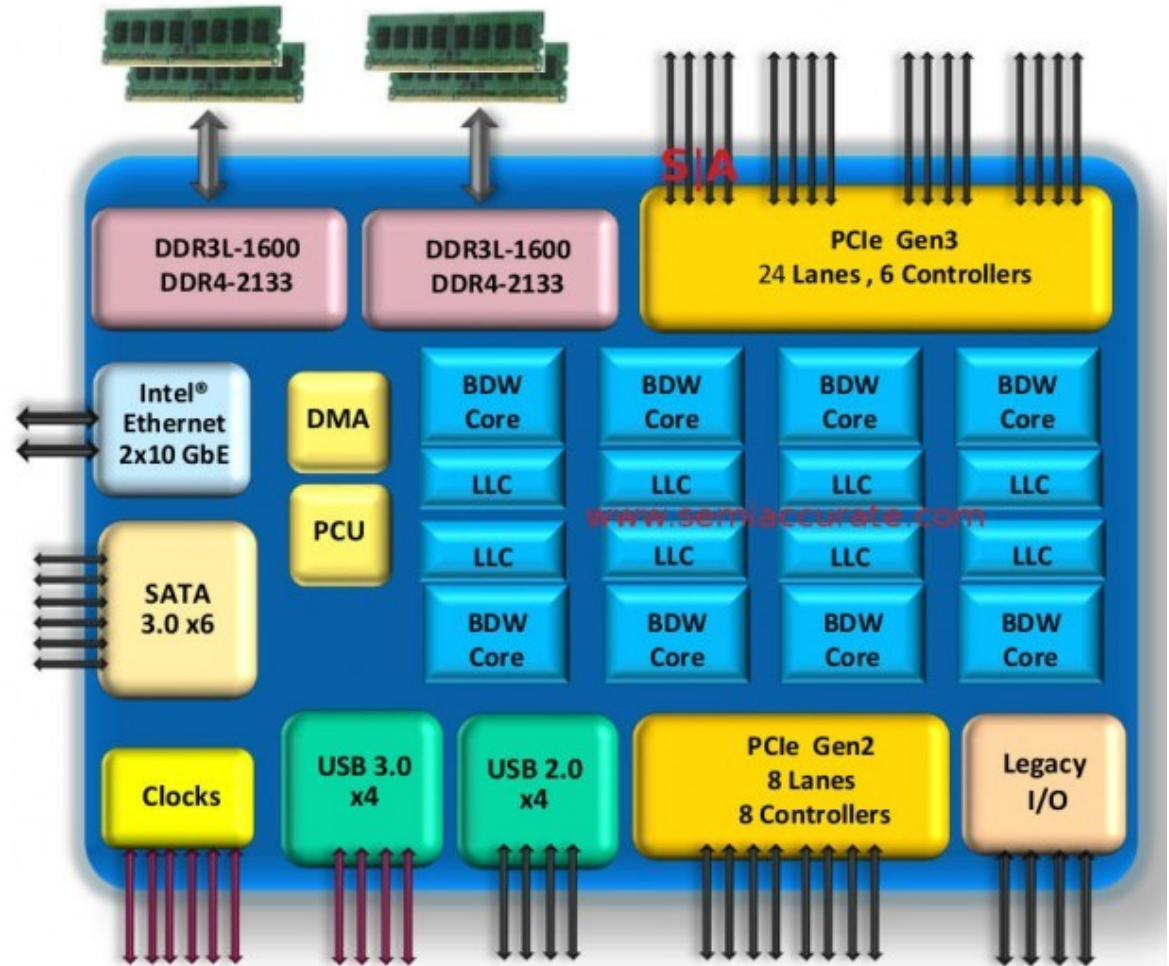
■ Desktop Model


- 4 cores
- Integrated graphics
- 3.3-3.8 GHz
- 65W

■ Server Model

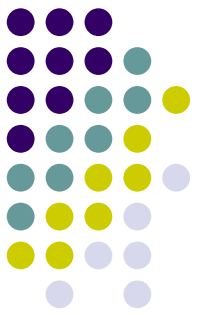
- 8 cores
- Integrated I/O
- 2-2.6 GHz
- 45W

Same ISA





Assembly Characteristics



Minimal Data Types in Registers

- “Integer” data of 1, 2, 4, 8 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- No aggregate types such as arrays or structures
 - Exception: SIMD registers (XMM in X86)


Limited Number of Registers

- Typically a few dozen (or less)

Code: Byte sequences that encode instructions

- Choice of instruction determines type of data!

Assembly Characteristics: Operations



- **Perform arithmetic function on register or memory data**
- **Transfer data between memory and register**
 - Load data from memory into register
 - Store register data into memory
- **Transfer control**
 - Unconditional jumps to/from procedures
 - Conditional branches

Compiling Into Assembly

C Code (sum.c)

```
long plus(long x, long y);

void sumstore(long x, long y,
              long *dest)
{
    long t = plus(x, y);
    *dest = t;
}
```

Generated Assembly

```
sumstore:
    pushq    %rbx
    movq     %rdx, %rbx
    call     plus
    movq     %rax, (%rbx)
    popq     %rbx
    ret
```

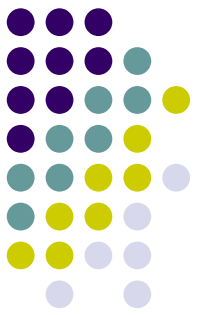
Obtain with command

```
gcc -S sum.c
```

Produces file `sum.s`

AMD64/x86-64
assembly!

Object Code from Assembly



Machine Code example

0x0400595:

0x53

0x48

0x89

0xd3

0xe8

0xf2

0xff

0xff

0xff

0x48

0x89

0x03

0x5b

0xc3

- **Total of 14 bytes**
- **Each instruction 1, 3, or 5 bytes**
- **Starts at address 0x0400595**

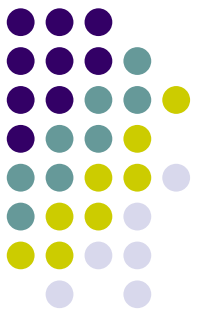
■ Assembler

- Translates .s into .o
- Binary encoding of each instruction
- Image of executable code (almost)
- Missing linkages between code in different files

■ Run 'gcc -c sum.s'

■ Examine result using objdump

Machine Instruction Example



```
*dest = t;
```

```
movq %rax, (%rbx)
```

```
0x40059e:  48 89 03
```

■ C Code

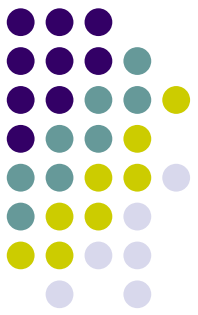
- Store value `t` where designated by `dest`

■ Assembly

- Move 8-byte value to memory
 - Quad words in x86-64 parlance
- Operands:
 - `t`: Register `%rax`
 - `dest`: Register `%rbx`
 - `*dest`: Memory `M[%rbx]`

■ Object Code

- 3-byte instruction
- Stored at address `0x40059e`



Disassembling Object Code

`objdump -d sumstore.o`

Disassembled

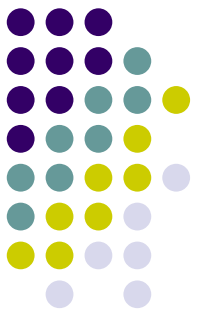
```
0000000000400595 <sumstore>:
 400595: 53                push    %rbx
 400596: 48 89 d3          mov     %rdx,%rbx
 400599: e8 f2 ff ff ff    callq   400590 <plus>
 40059e: 48 89 03          mov     %rax, (%rbx)
 4005a1: 5b                pop     %rbx
 4005a2: c3                retq
```

Disassembler

`objdump -d sum`

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can be run on either a `.out` (complete executable) or `.o` file

Alternate Disassembly



Object

0x0400595:

0x53

0x48

0x89

0xd3

0xe8

0xf2

0xff

0xff

0xff

0x48

0x89

0x03

0x5b

0xc3

Disassembled

Dump of assembler code for function sumstore:

```
0x0000000000400595 <+0>: push    %rbx
0x0000000000400596 <+1>: mov     %rdx,%rbx
0x0000000000400599 <+4>: callq   0x400590 <plus>
0x000000000040059e <+9>: mov     %rax, (%rbx)
0x00000000004005a1 <+12>: pop     %rbx
0x00000000004005a2 <+13>: retq
```

Within gdb Debugger

`gdb sum`

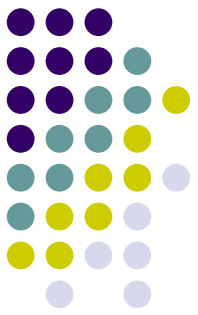
`disassemble sumstore`

- Disassemble procedure

`x/14xb sumstore`

- Examine the 14 bytes starting at `sumstore`

What Can be Disassembled?



- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

**Reverse engineering forbidden by
Microsoft End User License
Agreement**

```
% objdump -d WINWORD.EXE
```

```
WINWORD.EXE:      file format pei-i386
```


```
No symbols in "WINWORD.EXE".
```

```
Disassembly of section .text:
```

```
30001000 <.text>:
```

```
30001000:  55                push    %ebp
30001001:  8b ec            mov     %esp,%ebp
30001003:  6a ff            push    $0xffffffff
30001005:  68 90 10 00 30    push    $0x30001090
3000100a:  68 91 dc 4c 30    push    $0x304cdc91
```


x86-64 Integer Registers

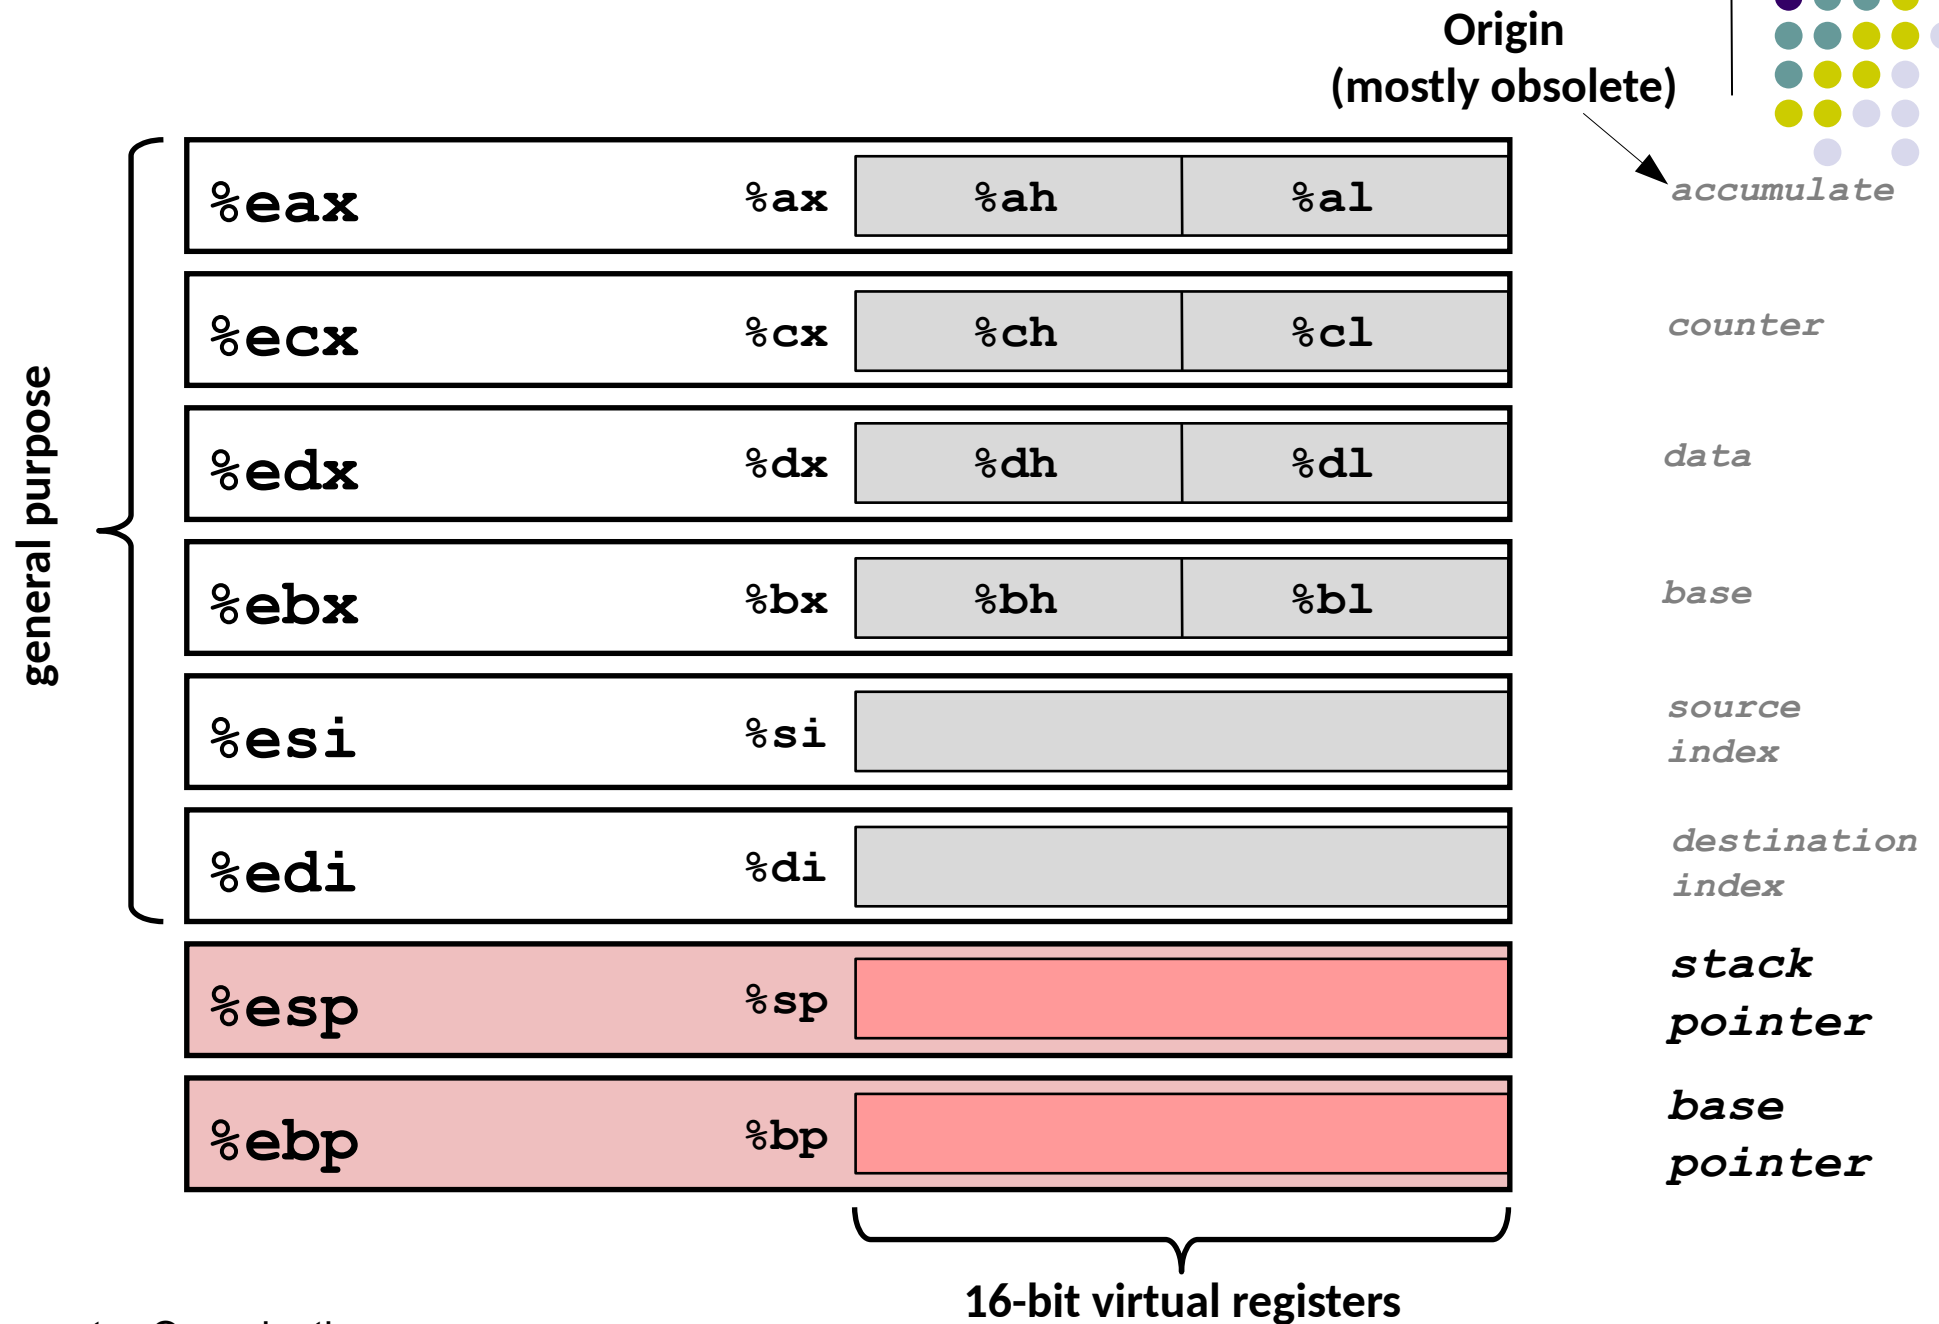


| | |
|-------------|-------------|
| %rax | %eax |
| %rbx | %ebx |
| %rcx | %ecx |
| %rdx | %edx |
| %rsi | %esi |
| %rdi | %edi |
| %rsp | %esp |
| %rbp | %ebp |


| | |
|-------------|--------------|
| %r8 | %r8d |
| %r9 | %r9d |
| %r10 | %r10d |
| %r11 | %r11d |
| %r12 | %r12d |
| %r13 | %r13d |
| %r14 | %r14d |
| %r15 | %r15d |

Can reference low-order 4 bytes (also low-order 1 & 2 bytes)

Some History: IA32 Registers



Moving Data




■ Moving Data

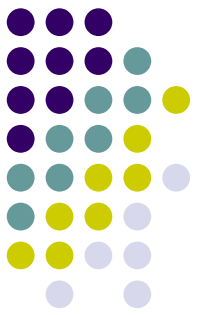
`movq Source, Dest:`

■ Operand Types

- **Immediate:** Constant integer data
 - Example: `$0x400`, `$-533`
 - Like C constant, but prefixed with ``$'`
 - Encoded with 1, 2, or 4 bytes
- **Register:** One of 16 integer registers
 - Example: `%rax`, `%r13`
 - But `%rsp` reserved for special use
 - Others have special uses for particular instructions
- **Memory:** 8 consecutive bytes of memory at address given by register
 - Simplest example: `(%rax)`
 - Various other “address modes”



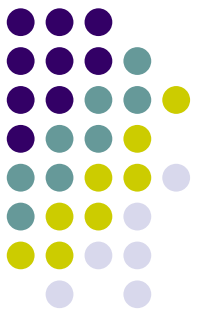
movq Operand Combinations



Cannot do memory-memory transfer with a single instruction

| | Source | Dest | Src, Dest | C Analog |
|------|--------|------|---------------------|----------------|
| movq | Imm | Reg | movq \$0x4, %rax | temp = 0x4; |
| | | Mem | movq \$-147, (%rax) | *p = -147; |
| | Reg | Reg | movq %rax, %rdx | temp2 = temp1; |
| | | Mem | movq %rax, (%rdx) | *p = temp; |
| | Mem | Reg | | |
| | | | movq (%rax), %rdx | temp = *p; |

Simple Memory Addressing Modes



■ Normal (R) Mem[Reg[R]]

- Register R specifies memory address
- Pointer dereferencing in C: “(%rcx)” is really “ *rcx ” in C notation


```
movq (%rcx), %rax
```

■ Displacement D(R) Mem[Reg[R]+D]

- Register R specifies start of memory region
- Constant displacement D specifies offset (think of “*(rbp+8)” in C)

```
movq 8(%rbp), %rdx
```

Example of Simple Addressing Modes



```
void swap
(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap:
    movq    (%rdi), %rax
    movq    (%rsi), %rdx
    movq    %rdx, (%rdi)
    movq    %rax, (%rsi)
    ret
```

Understanding Swap()




```
void swap
(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Registers

| | |
|------|--|
| %rdi | |
| %rsi | |
| %rax | |
| %rdx | |

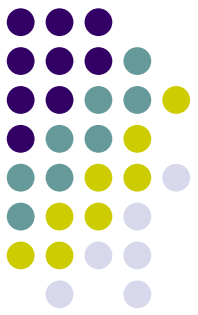
Memory



| Register | Value |
|----------|-------|
| %rdi | xp |
| %rsi | yp |
| %rax | t0 |
| %rdx | t1 |

swap:

```
movq    (%rdi), %rax    # t0 = *xp
movq    (%rsi), %rdx    # t1 = *yp
movq    %rdx, (%rdi)    # *xp = t1
movq    %rax, (%rsi)    # *yp = t0
ret
```



Complete Memory Addressing Modes

■ Most General Form

$D(Rb, Ri, S) \text{ Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri] + D]$


- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for `%rsp`
- S: Scale: 1, 2, 4, or 8 (*why these numbers?*)

■ Special Cases

$(Rb, Ri) \text{ Mem}[\text{Reg}[Rb] + \text{Reg}[Ri]]$

$D(Rb, Ri) \text{ Mem}[\text{Reg}[Rb] + \text{Reg}[Ri] + D]$


$(Rb, Ri, S) \text{ Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri]]$



Address Computation Examples

| | |
|------|--------|
| %rdx | 0xf000 |
| %rcx | 0x0100 |

| Expression | Address Computation | Address |
|---------------|----------------------|---------|
| 0x8(%rdx) | $0xf000 + 0x8$ | 0xf008 |
| (%rdx,%rcx) | $0xf000 + 0x100$ | 0xf100 |
| (%rdx,%rcx,4) | $0xf000 + 4 * 0x100$ | 0xf400 |
| 0x80(,%rdx,2) | $2 * 0xf000 + 0x80$ | 0x1e080 |



Address Computation Instruction

■ **leaq Src, Dst**

- Src is address mode expression
- Set Dst to address denoted by expression

■ **Uses**

- Computing addresses without a memory reference
 - E.g., translation of `p = &x[i];`
- Computing arithmetic expressions of the form $x + k * y$
 - $k = 1, 2, 4, \text{ or } 8$


■ **Example**

```
long m12(long x)
{
    return x*12;
}
```

Converted to ASM by compiler:

```
leaq (%rdi,%rdi,2), %rax # t <- x+x*2
salq $2, %rax           # return t<<2
```

%rdi contains x



Some Arithmetic Operations


■ Two Operand Instructions:

FormatComputation

| | | | |
|-------|-----------|---------------------------------------|------------------|
| addq | Src, Dest | Dest = Dest + Src | |
| subq | Src, Dest | Dest = Dest $\tilde{}$ Src | |
| imulq | Src, Dest | Dest = Dest * Src | |
| salq | Src, Dest | Dest = Dest << Src | Also called shlq |
| sarq | Src, Dest | Dest = Dest >> Src | Arithmetic |
| shrq | Src, Dest | Dest = Dest >> Src | Logical |
| xorq | Src, Dest | Dest = Dest ^ Src | |
| andq | Src, Dest | Dest = Dest & Src | |
| orq | Src, Dest | Dest = Dest Src | |

■ Watch out for argument order!

■ No distinction between signed and unsigned int (why?)




Some Arithmetic Operations

■ One operand instructions

FormatComputation

| | | |
|-------------------|-------------------|---------------------------------|
| <code>incq</code> | <code>Dest</code> | $\text{Dest} = \text{Dest} + 1$ |
| <code>decq</code> | <code>Dest</code> | $\text{Dest} = \text{Dest} - 1$ |
| <code>negq</code> | <code>Dest</code> | $\text{Dest} = -\text{Dest}$ |
| <code>notq</code> | <code>Dest</code> | $\text{Dest} = \sim\text{Dest}$ |

■ Lots more



Arithmetic Expression Example (64 bit)

```
long arith
(long x, long y, long z)
{
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```


```
arith:
    leaq    (%rdi,%rsi), %rax
    addq    %rdx, %rax
    leaq    (%rsi,%rsi,2), %rdx
    salq    $4, %rdx
    leaq    4(%rdi,%rdx), %rcx
    imulq   %rcx, %rax
    ret
```

Interesting Instructions

- **leaq**: address computation
- **salq**: shift
- **imulq**: multiplication
 - But, only used once

Understanding Arithmetic Expression

Example




```
long arith
(long x, long y, long z)
{
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```

arith:

```
leaq    (%rdi,%rsi), %rax    # t1
addq    %rdx, %rax          # t2
leaq    (%rsi,%rsi,2), %rdx  # rdx = 3y
salq    $4, %rdx            # t4 (3y*16)
leaq    4(%rdi,%rdx), %rcx   # t5
imulq    %rcx, %rax          # rval
ret
```

| Register | Use(s) |
|----------|----------------|
| %rdi | Argument x |
| %rsi | Argument y |
| %rdx | Argument z, 3y |
| %rax | t1, t2, rval |
| %rdx | t4 |
| %rcx | t5 |

Another Example



```
int logical(int x, int y)
{
    int t1 = x^y;
    int t2 = t1 >> 17;
    int mask = (1<<13) - 7;
    int rval = t2 & mask;
    return rval;
}
```

$2^{13} = 8192, 2^{13} - 7 = 8185$

logical:

```
pushl %ebp
movl %esp,%ebp
```

} Set Up

```
movl 8(%ebp),%eax
xorl 12(%ebp),%eax
sarl $17,%eax
andl $8185,%eax
```

} Body

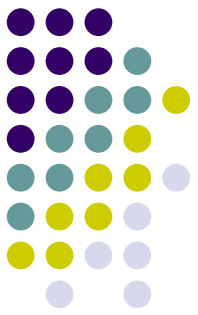
```
movl %ebp,%esp
popl %ebp
ret
```

} Finish

```
movl 8(%ebp),%eax
xorl 12(%ebp),%eax
sarl $17,%eax
andl $8185,%eax
```

```
eax = x
eax = x^y      (t1)
eax = t1>>17  (t2)
eax = t2 & 8185
```

Whose Assembler?



Intel/Microsoft Format

```
lea    eax, [ecx+ecx*2]
sub     esp, 8
cmp     dword ptr [ebp-8], 0
mov     eax, dword ptr [eax*4+100h]
```

GAS/Gnu Format

```
leal    (%ecx,%ecx,2),%eax
subl    $8,%esp
cmpl    $0,-8(%ebp)
movl    $0x100(,%eax,4),%eax
```

Intel/Microsoft Differs from GAS

- Operands listed in opposite order

`mov Dest, Src` `movl Src, Dest`

- Constants not preceded by '\$', Denote hex with 'h' at end

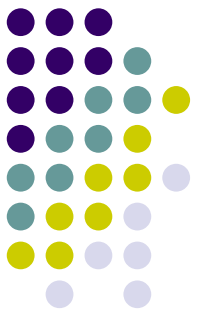
`100h` `$0x100`

- Operand size indicated by operands rather than operator suffix

`sub` `subl`

- Addressing format shows effective address computation

`[eax*4+100h]` `$0x100(,%eax,4)`



Halt and Catch Fire (HCF)

- פקודות אסמבלי ידועות בשמות מוזרים (LEA? באמת?)
- במחשב IBM SYSTEM/360, התגלה באג במעבד שגרם לו להתקע
- בצחוק, המתכנתים החליטו לסמן את הפקודה ב HCF
- יותר מאוחר, היא הפכה לפקודה רשמית
 - למטרות בדיקה עצמית – לא עושה שום דבר הרסני
- ועוד אחר כך, לכינוי כללי לפקודות מכונה שגורמות למחשב להתקע
 - קרה מאז בהרבה מעבדים, כולל מעבדי X86
- ועוד אחר כך: סדרת טלוויזיה מרתקת על תחילת עידן המחשב האישי
 - אבל זה לאחרי הסמסטר, בסדר?

