


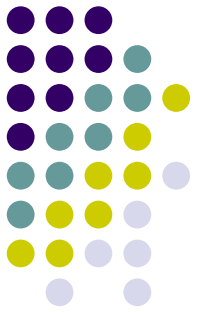
# Computer Organization: A Programmer's Perspective

---

## Machine-Level Programming (4: Data Structures)



# Today



## ■ Arrays

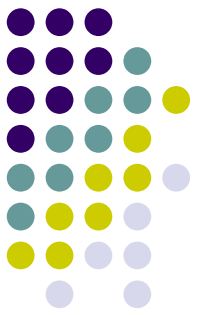
- One-dimensional
- Multi-dimensional (nested)
- Multi-level

## ■ Structures

- Allocation
- Access
- Alignment

## ■ Optional: Floating Point

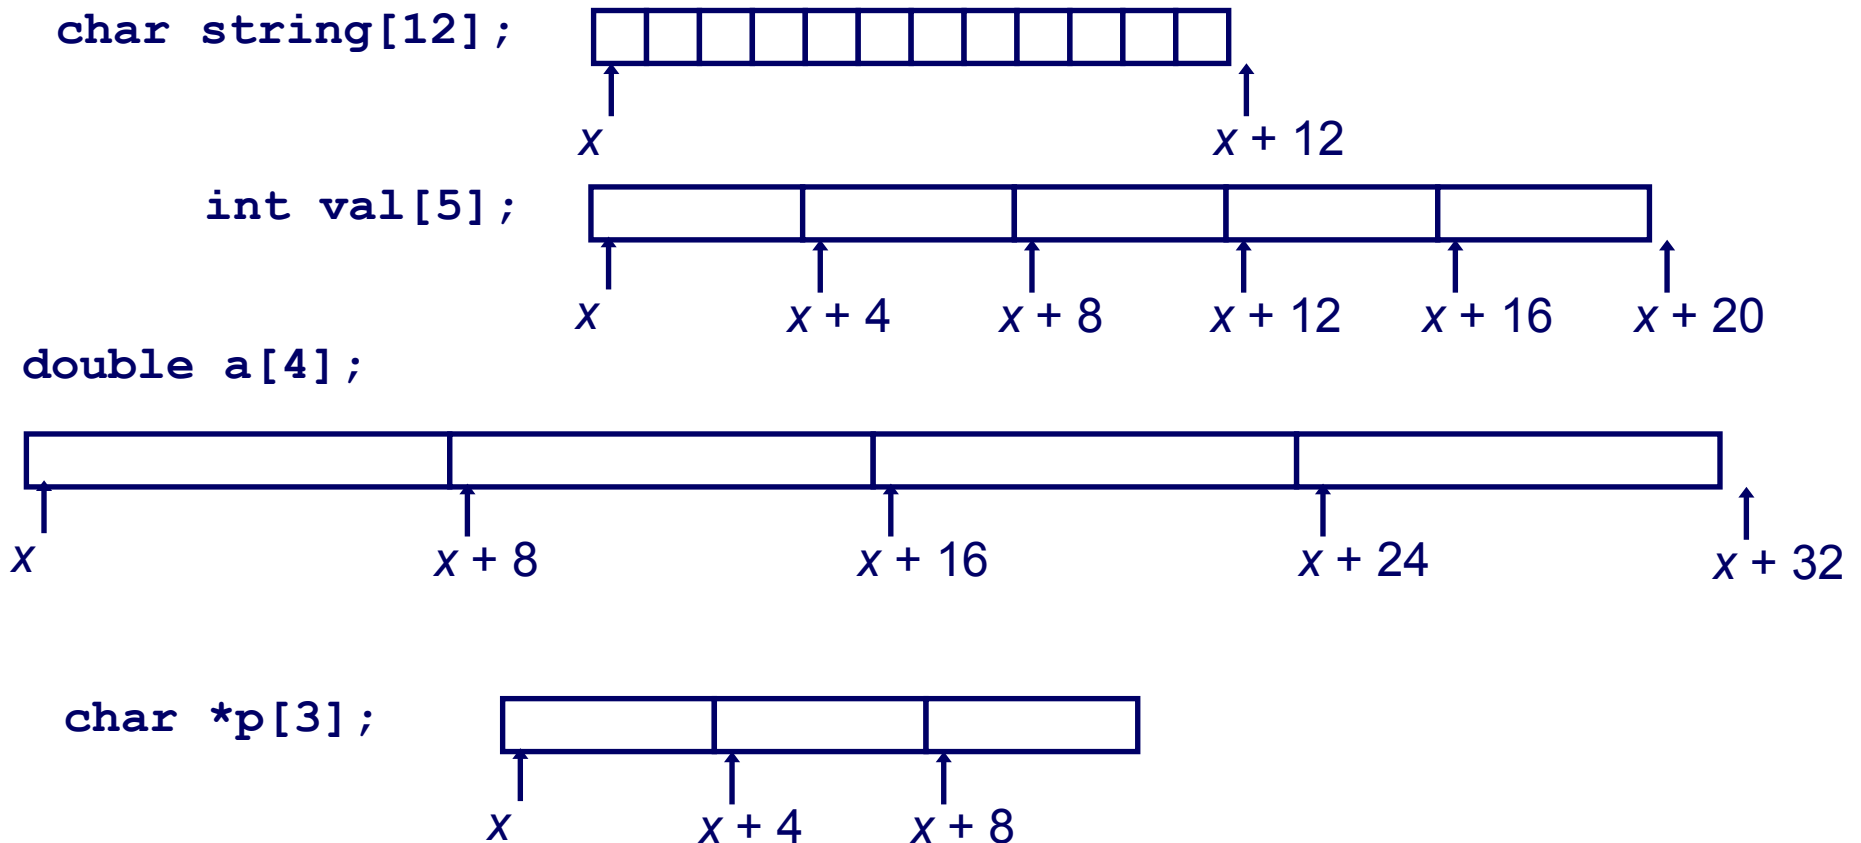
# Array Allocation



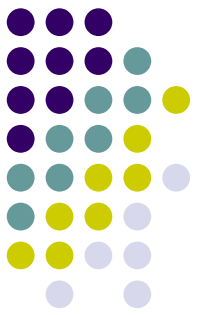
## Basic Principle

$T \ A[L];$

- Array of data type  $T$  and length  $L$
- Contiguously allocated region of  $L * \text{sizeof}(T)$  bytes




# Array Access



- Basic Principle

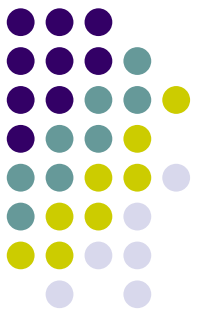
- $T \ A[L]; \implies$  Array of data type  $T$  and length  $L$
- Identifier  $A$  can be used as a pointer to array element 0



- 

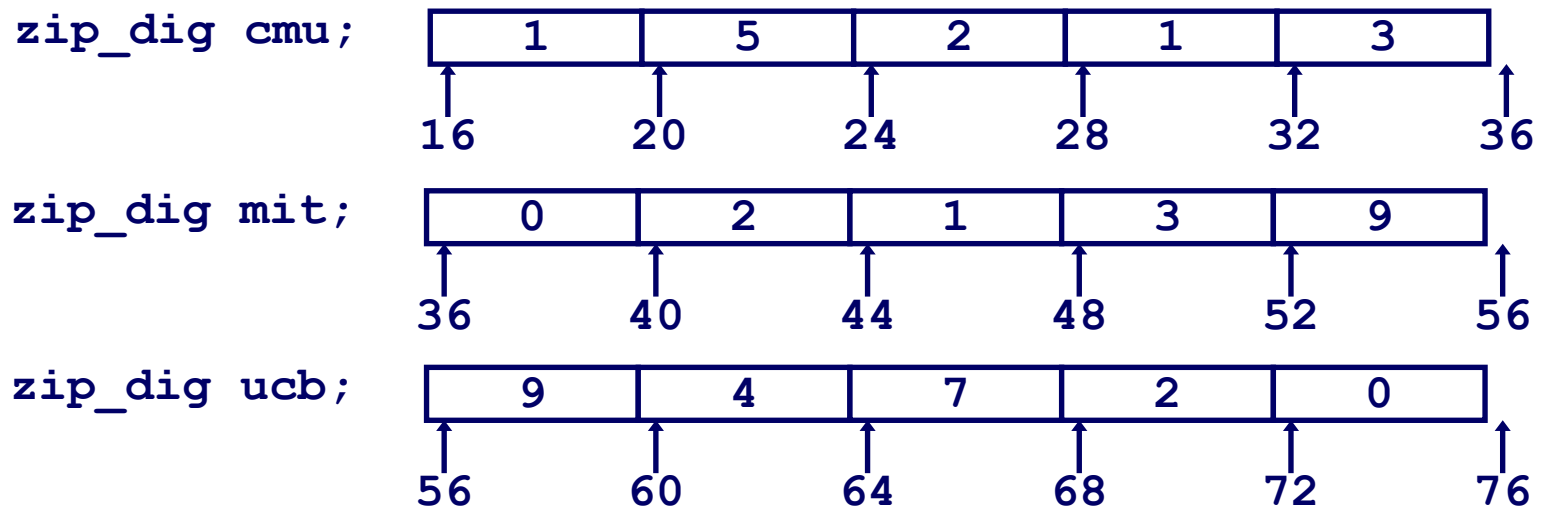
Reference	Type	Value
<code>val[4]</code>	<code>int</code>	3
<code>val</code>	<code>int *</code>	$x$
<code>val+1</code>	<code>int *</code>	$x + 4$
<code>&amp;val[2]</code>	<code>int *</code>	$x + 8$
<code>val[5]</code>	<code>int</code>	??
<code>*(val+1)</code>	<code>int</code>	5
<code>val + i</code>	<code>int *</code>	$x + 4 i$

# Array Example



```
#define ZLEN 5
typedef int zip_dig[ZLEN];

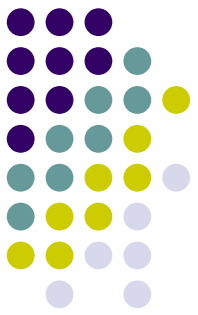
zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```



## Notes

- Declaration “zip\_dig cmu” equivalent to “int cmu[5]”
- Example arrays were allocated in successive 20 byte blocks
  - **Not guaranteed to happen in general**


# Array Accessing Example



```
int get_digit
(zip_dig z, int dig)
{
    return z[dig];
}
```

- Register `%rdi`: array starting address
- Register `%rsi`: array index
- Desired digit at: `%rdi + 4*%rsi`
- Use memory reference:
  - `(%rdi,%rsi,4)`

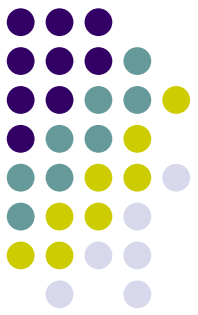
```
# %rdi = z
# %rsi = dig
movl (%rdi,%rsi,4), %eax # z[dig]
```



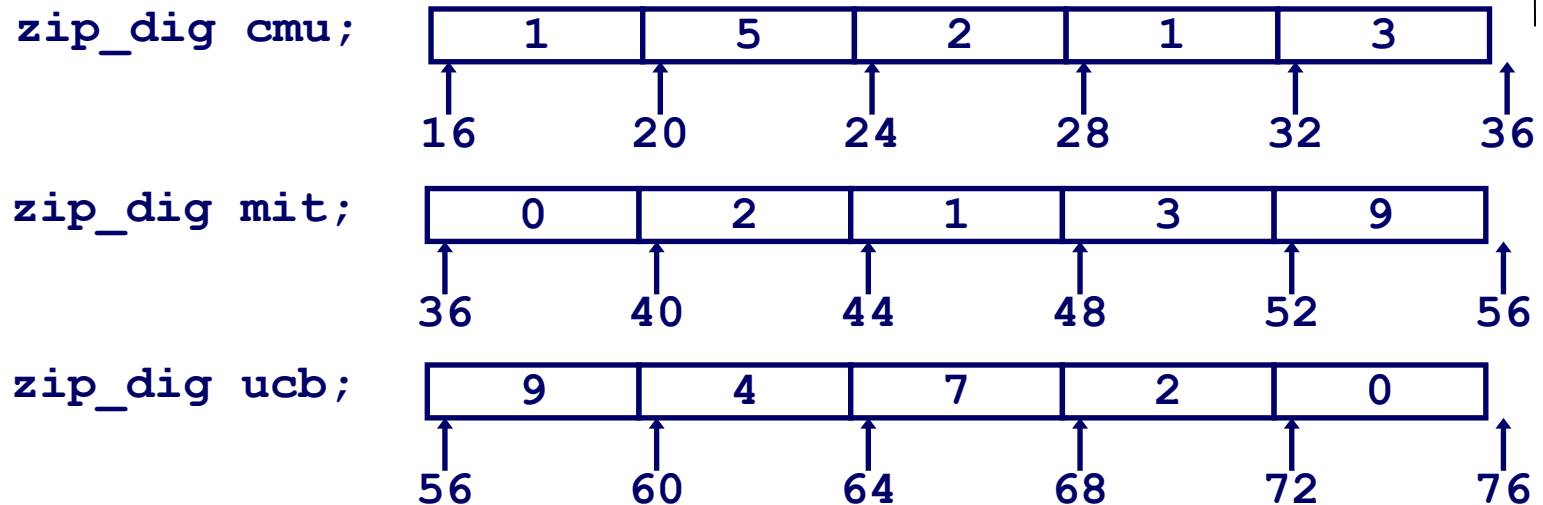
# Array Loop Example

```
void zincr(zip_dig z) {  
    size_t i;  
  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

```
# %rdi = z  
movl    $0, %eax          # i = 0  
jmp     .L3               # goto middle  
.L4:                       # loop:  
addl    $1, (%rdi,%rax,4) # z[i]++  
addq    $1, %rax          # i++  
.L3:                       # middle  
cmpq    $4, %rax          # i:4  
jbe     .L4               # if <=, goto loop  
rep; ret
```



# Referencing Examples



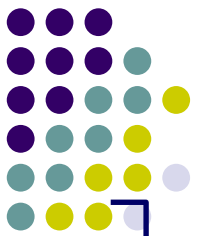
Code Does Not Do Any Bounds Checking!

Reference	Address	Value	Guaranteed?
<code>mit[3]</code>	$36 + 4 * 3 = 48$	3	<b>Yes</b>
<code>mit[5]</code>	$36 + 4 * 5 = 56$	9	<b>No</b>
<code>mit[-1]</code>	$36 + 4 * -1 = 32$	3	<b>No</b>
<code>cmu[15]</code>	$16 + 4 * 15 = 76$	??	<b>No</b>

- Out of range behavior implementation-dependent
  - No guaranteed relative allocation of different arrays




# Nested Array Allocation



## Declaration

$T \ A[R][C];$

- Array of data type  $T$
- $R$  rows,  $C$  columns
- Type  $T$  element requires  $K$  bytes



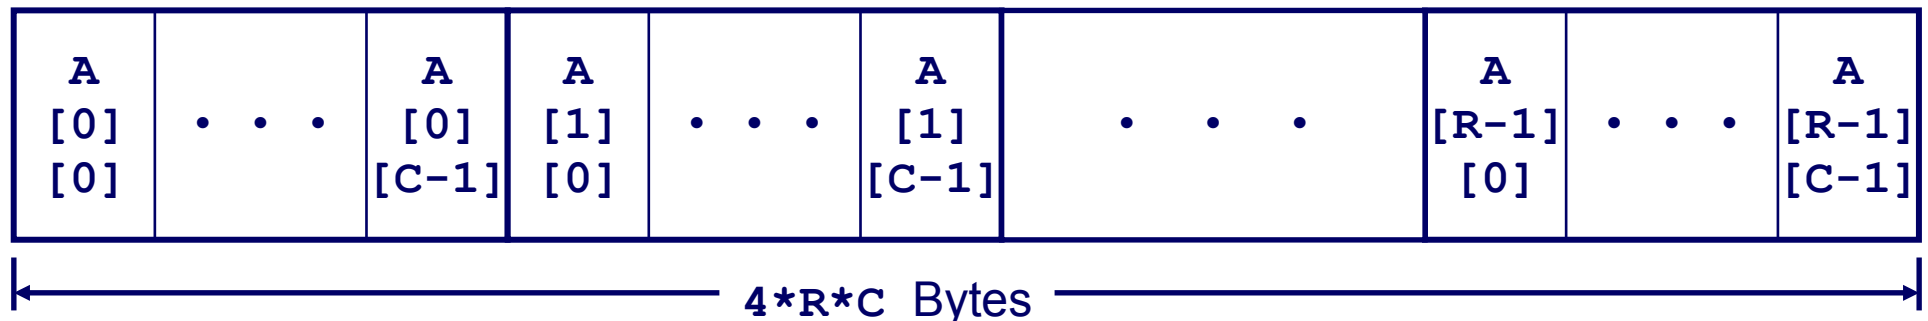
## Array Size

- $R * C * K$  bytes

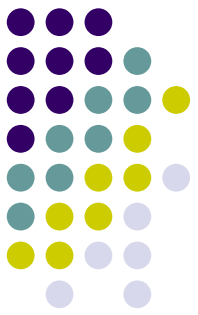
## Arrangement

- Row-Major Ordering

`int A[R][C];`

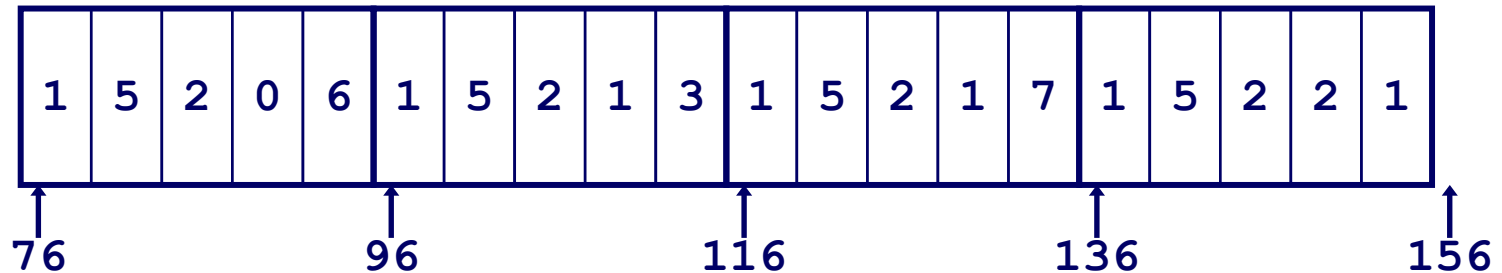


# Nested Array Example



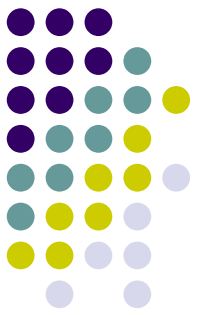
```
#define PCOUNT 4
zip_dig pgh[PCOUNT] =
    {{1, 5, 2, 0, 6},
     {1, 5, 2, 1, 3 },
     {1, 5, 2, 1, 7 },
     {1, 5, 2, 2, 1 }};
```

zip\_dig  
pgh[4];



- Declaration “zip\_dig pgh[4]” equivalent to “int pgh[4][5]”
- Variable `pgh` denotes array of 4 elements, allocated contiguously
  - Each element is an array of 5 `int`'s, allocated contiguously
- “Row-Major” ordering of all elements guaranteed

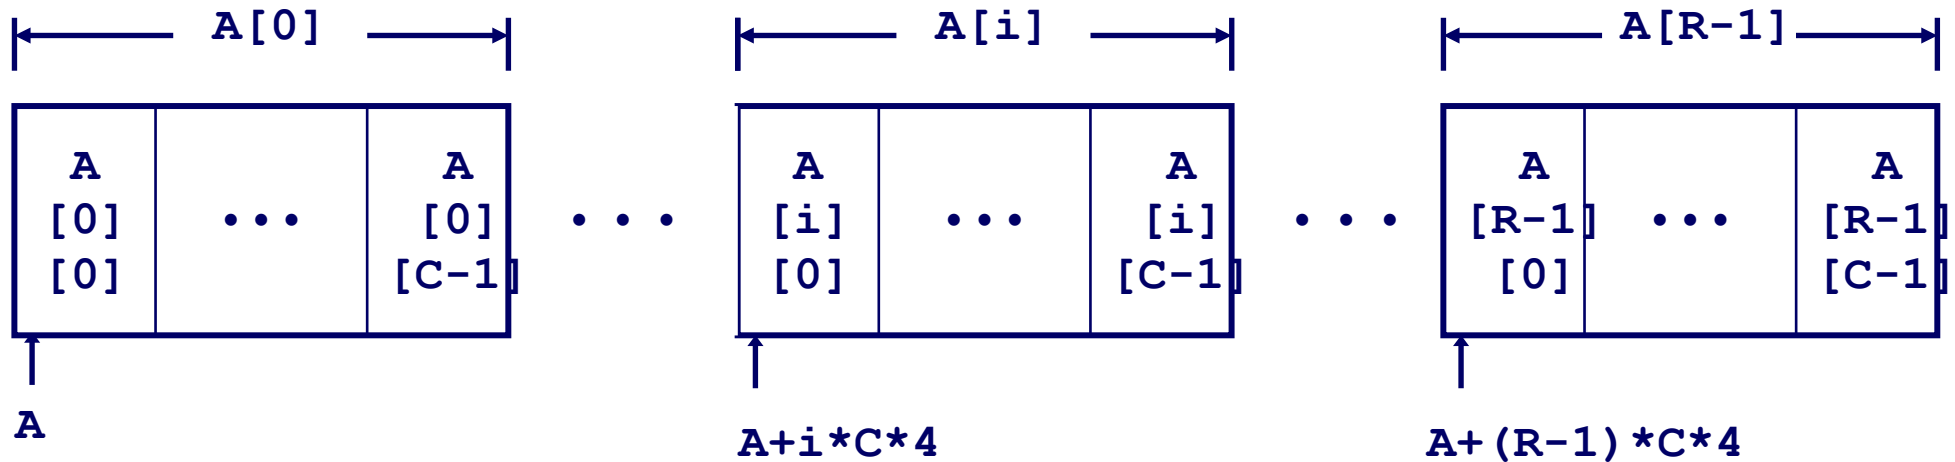
# Nested Array Row Access




## Row Vectors

- $A[i]$  is array of  $C$  elements
- Each element of type  $T$
- Starting address  $A + i * C * K$

```
int A[R][C];
```



# Nested Array Row Access Code



```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

## Row Vector


- `pgh[index]` is array of 5 int's
- Starting address `pgh+20*index`

## Code

- Computes and returns address
- Compute as `pgh + 4*(index+4*index)`


```
# %rdi = index
leaq (%rdi,%rdi,4),%rax # 5 * index
leaq pgh(,%rax,4),%rax  # pgh + (20 * index)
```

# Nested Array Element Access




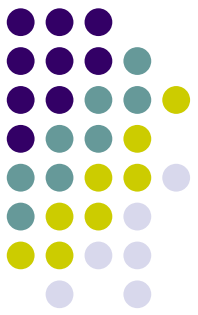
## Array Elements

- $A[i][j]$  is element of type  $T$
- Address  $A + (i * C + j) * K$



```
int A[R][C];
```





# Nested Array Element Access Code

## Array Elements

- `pgh[index][dig]` is `int`
- Address:  $pgh + 20 * index + 4 * dig = pgh + 4 * (5 * index + dig)$

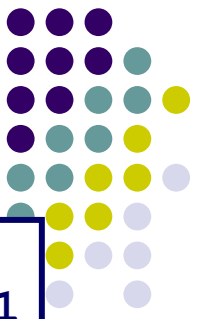
## Code

- Computes address as:  $pgh + 4 * dig + 4 * (index + 4 * index)$
- `movl` performs memory reference


```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

```
leaq    (%rdi,%rdi,4), %rax    # 5*index
addl    %rax, %rsi             # 5*index+dig
movl    pgh(,%rsi,4), %eax     # M[pgh + 4*(5*index+dig)]
```

# Referencing Examples



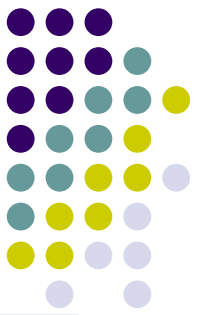
zip\_dig  
pgh[4];



Reference	Address	Value	Guaranteed?
Pgh[3][3]	$76 + 20 * 3 + 4 * 3 = 148$	2	Yes
Pgh[2][5]	$76 + 20 * 2 + 4 * 5 = 136$	1	Yes
Pgh[2][-1]	$76 + 20 * 2 + 4 * -1 = 112$	3	Yes
Pgh[4][-1]	$76 + 20 * 4 + 4 * -1 = 152$	1	Yes
Pgh[0][19]	$76 + 20 * 0 + 4 * 19 = 152$	1	Yes
Pgh[0][-1]	$76 + 20 * 0 + 4 * -1 = 72$	??	No

- Code does not do any bounds checking
- Ordering of elements within array guaranteed

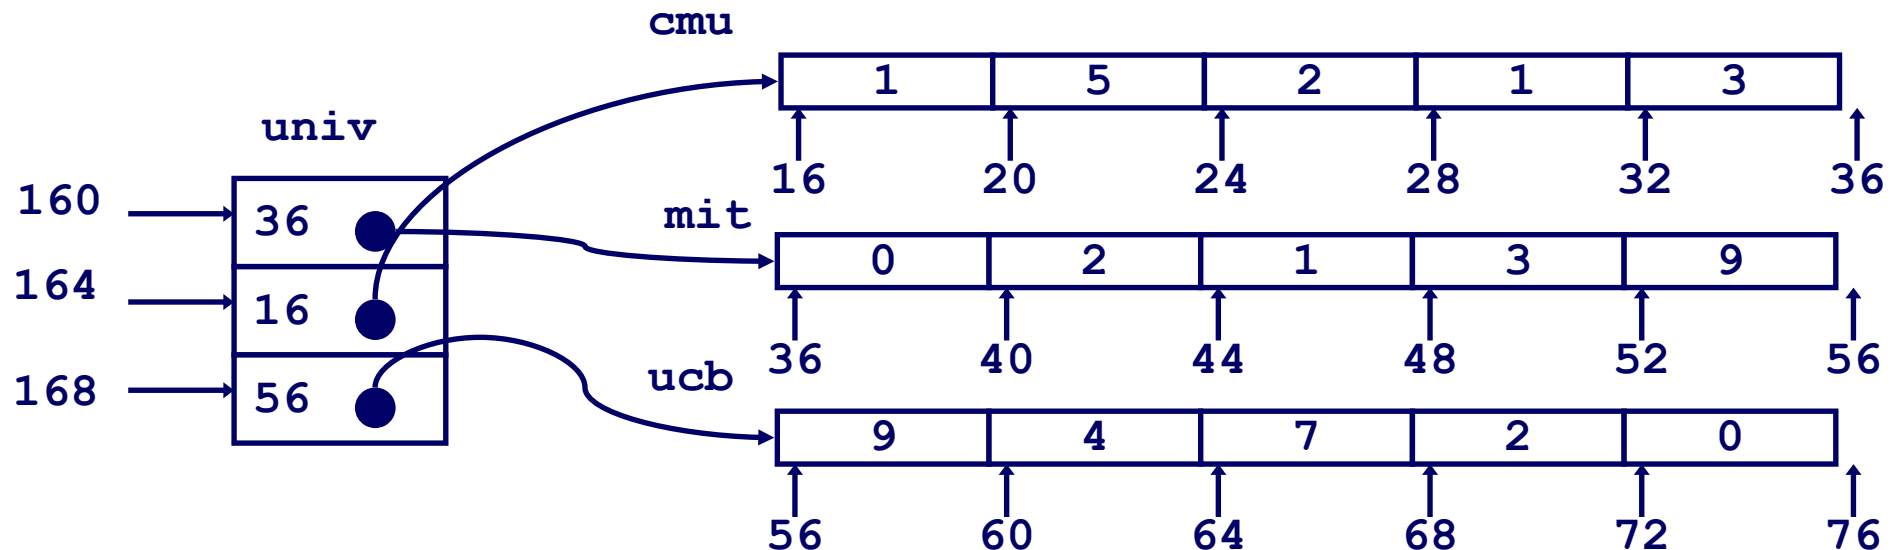
# Multi-Level Array Example



- Variable `univ` denotes array of 3 elements
- Each element is a pointer
  - 4 bytes
- Each pointer points to array of `int`'s

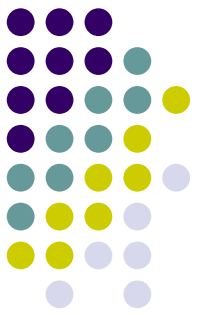
```
zip_dig cmu = { 1, 5, 2, 1, 3 };  
zip_dig mit = { 0, 2, 1, 3, 9 };  
zip_dig ucb = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3  
int *univ[UCOUNT] = {mit, cmu, ucb};
```





# Element Access in Multi-Level Array



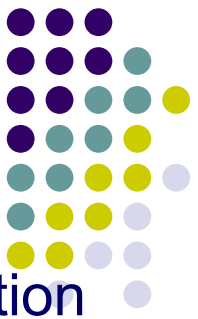
```
int get_univ_digit
(int index, int digit)
{
    return univ[index][digit];
}
```

## Computation

- Element access  $\text{Mem}[\text{Mem}[\text{univ} + 8 * \text{index}] + 4 * \text{digit}]$
- Must do two memory reads
  - First get pointer to row array
  - Then access element within array

```
salq    $2, %rsi           # 4*digit
addq    univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax        # return *p
ret
```

# Array Element Accesses




- Similar C references

## Nested Array

```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

- Element at

$\text{Mem}[\text{pgh} + 20 * \text{index} + 4 * \text{dig}]$




- Different address computation

## Multi-Level Array

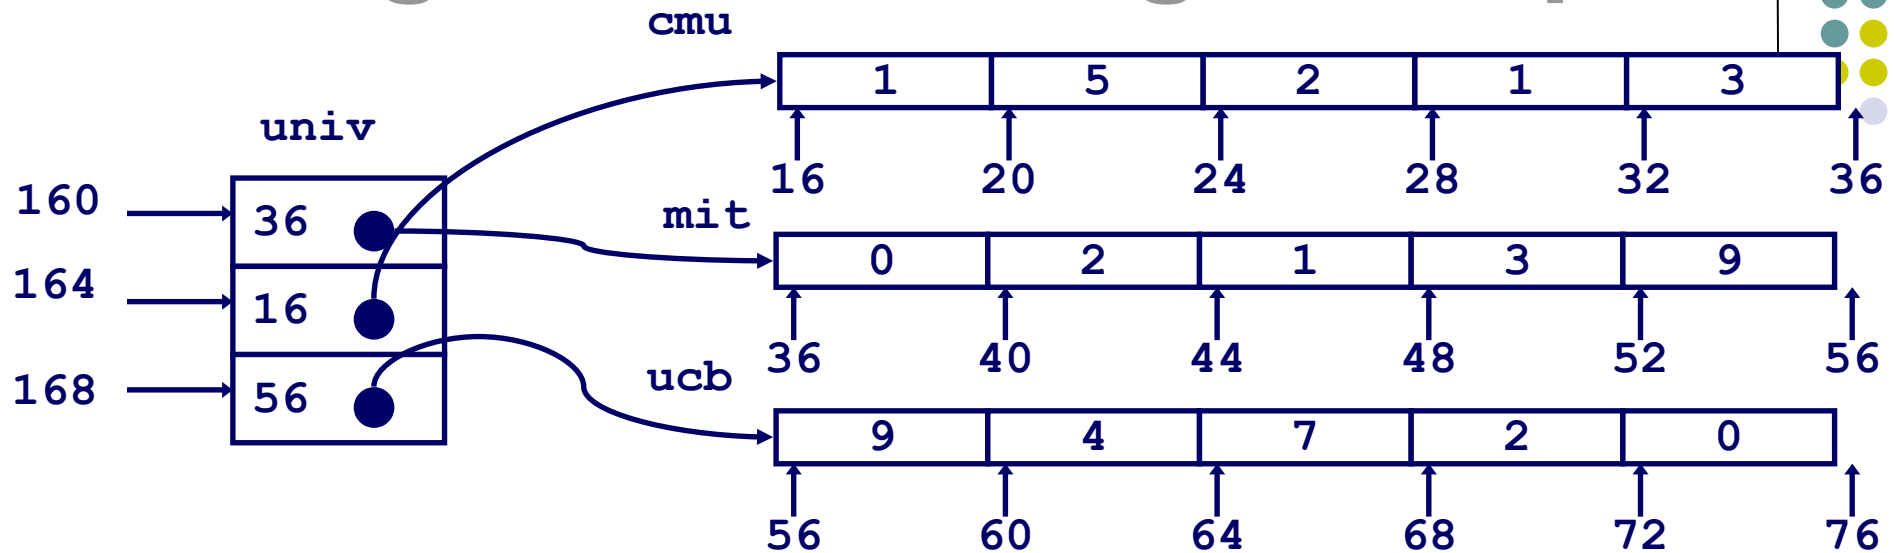
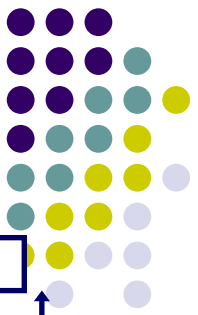
```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```

- Element at

$\text{Mem}[\text{Mem}[\text{univ} + 8 * \text{index}] + 4 * \text{dig}]$



# Strange Referencing Examples



Reference	Address	Value	Guaranteed?
<code>univ[2][3]</code>	$56 + 4 * 3 = 68$	2	<b>Yes</b>
<code>univ[1][5]</code>	$16 + 4 * 5 = 36$	0	<b>No</b>
<code>univ[2][-1]</code>	$56 + 4 * -1 = 52$	9	<b>No</b>
<code>univ[3][-1]</code>	??	??	<b>No</b>
<code>univ[1][12]</code>	$16 + 4 * 12 = 64$	7	<b>No</b>

- Code does not do any bounds checking
- Ordering of elements in different arrays not guaranteed



# N X N Matrix Code

## ■ Fixed dimensions

- Know value of N at compile time

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele(fix_matrix a,
            size_t i, size_t j)
{
    return a[i][j];
}
```

## ■ Variable dimensions, explicit indexing

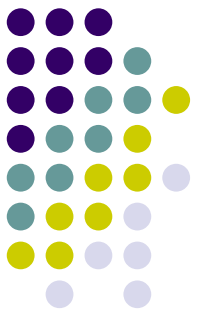
- Traditional way to implement dynamic arrays

```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element a[i][j] */
int vec_ele(size_t n, int *a,
            size_t i, size_t j)
{
    return a[IDX(n,i,j)];
}
```

## ■ Variable dimensions, implicit indexing

- Now supported by gcc
- ISO C99 standard
- Try at your own risk

```
/* Get element a[i][j] */
int var_ele(size_t n, int a[n][n],
            size_t i, size_t j) {
    return a[i][j];
}
```

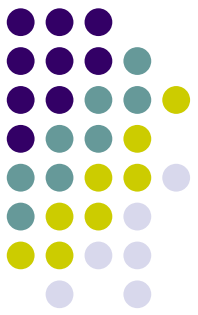


# 16 X 16 Matrix Access

- Address  $A + i * (C * K) + j * K$
- $C = 16, K = 4$

```
/* Get element a[i][j] */  
int fix_ele(fix_matrix a, size_t i, size_t j) {  
    return a[i][j];  
}
```

```
# a in %rdi, i in %rsi, j in %rdx  
salq    $6, %rsi          # 64*i  
addq    %rsi, %rdi         # a + 64*i  
movl    (%rdi,%rdx,4), %eax # M[a + 64*i + 4*j]  
ret
```



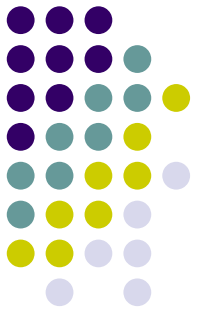
# n X n Matrix Access

- Address  $A + i * (C * K) + j * K$
- $C = n, K = 4$
- Must perform integer multiplication

```
/* Get element a[i][j] */  
int var_ele(size_t n, int a[n][n], size_t i, size_t j) {  
    return a[i][j];  
}
```

```
# n in %rdi, a in %rsi, i in %rdx, j in %rcx  
imulq    %rdx, %rdi          # n*i  
leaq     (%rsi,%rdi,4), %rax  # a + 4*n*i  
movl     (%rax,%rcx,4), %eax  # a + 4*n*i + 4*j  
ret
```

# Today



## ■ Arrays

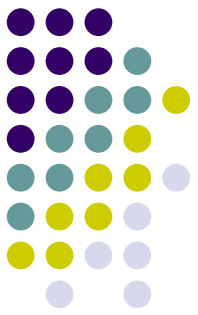
- One-dimensional
- Multi-dimensional (nested)
- Multi-level

## ■ Structures

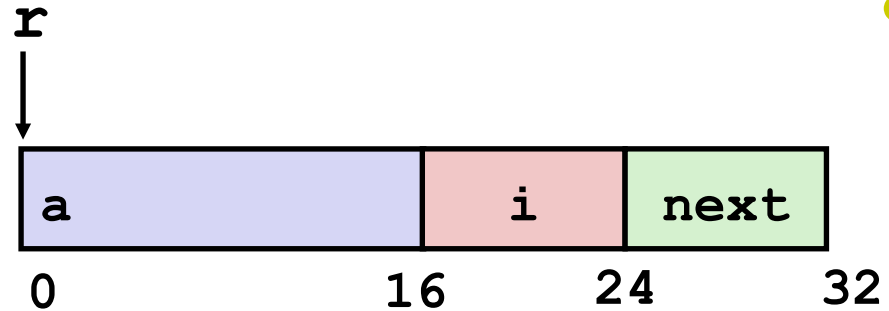
- Allocation
- Access
- Alignment

## ■ Optional: Floating Point

# Structure Representation

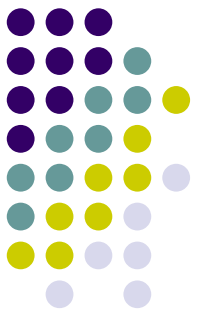


```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



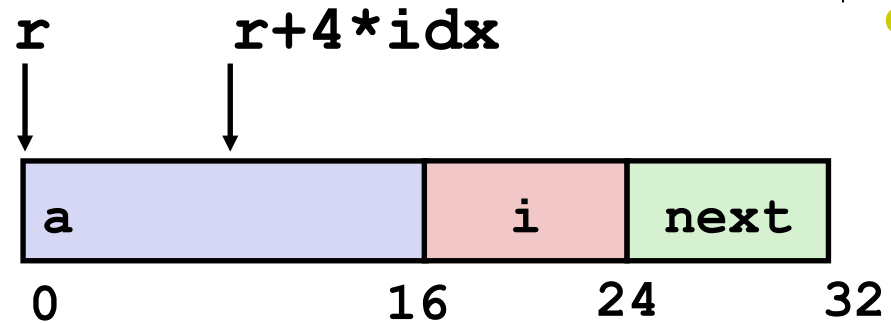
- **Structure represented as block of memory**
  - Big enough to hold all of the fields
- **Fields ordered according to declaration**
  - Even if another ordering could yield a more compact representation
- **Compiler determines overall size + positions of fields**
  - Machine-level program has no understanding of the structures in the source code





# Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



## ■ Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Compute as  $r + 4 * idx$

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```


```
# r in %rdi, idx in %rsi  
leaq (%rdi,%rsi,4), %rax  
ret
```

# Following Linked List

## ■ C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    int i;
    struct rec *next;
};
```



Element i

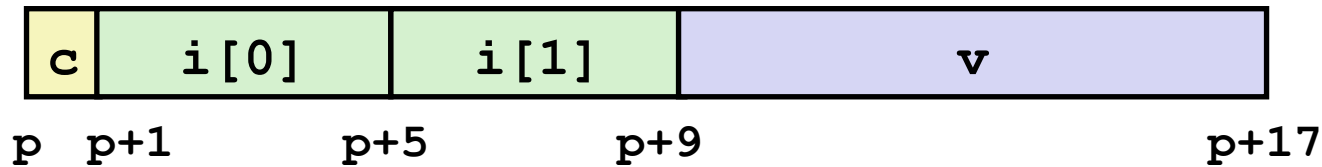
Register	Value
%rdi	r
%rsi	val

```
.L11:                                # loop:
    movslq    16(%rdi), %rax          # i = M[r+16]
    movl      %esi, (%rdi,%rax,4)     # M[r+4*i] = val
    movq      24(%rdi), %rdi         # r = M[r+24]
    testq     %rdi, %rdi             # Test r
    jne       .L11                  # if !=0 goto loop
```

# Structures & Alignment



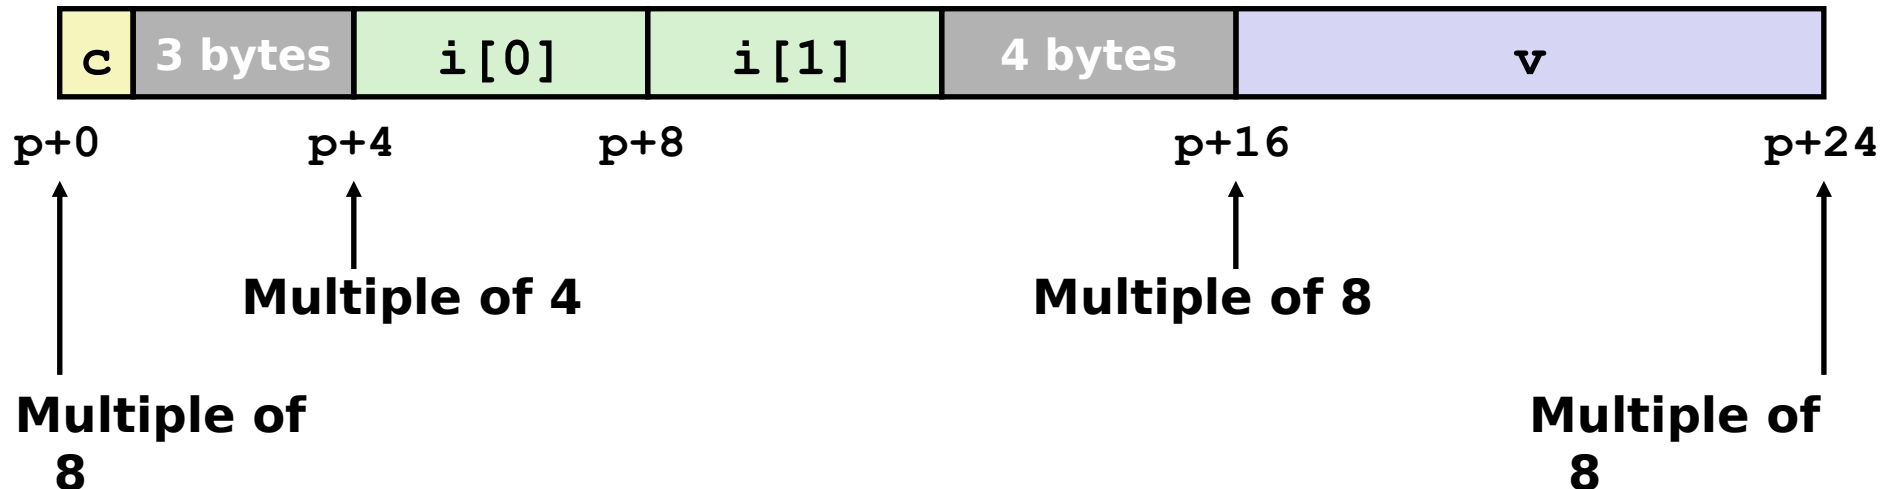
## ■ Unaligned Data

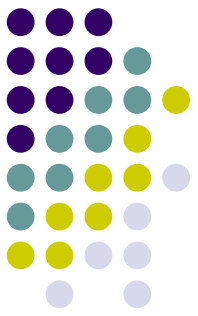


```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

## ■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K





# Alignment Principles

## ■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K
- Required on some machines; advised on x86-64

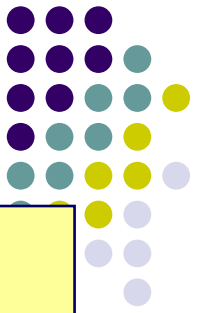
## ■ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
  - Inefficient to load or store datum that spans quad word boundaries
  - Virtual memory trickier when datum spans 2 pages

## ■ Compiler

- Inserts gaps in structure to ensure correct alignment of fields


# Linux vs. Windows (IA32 bits)



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

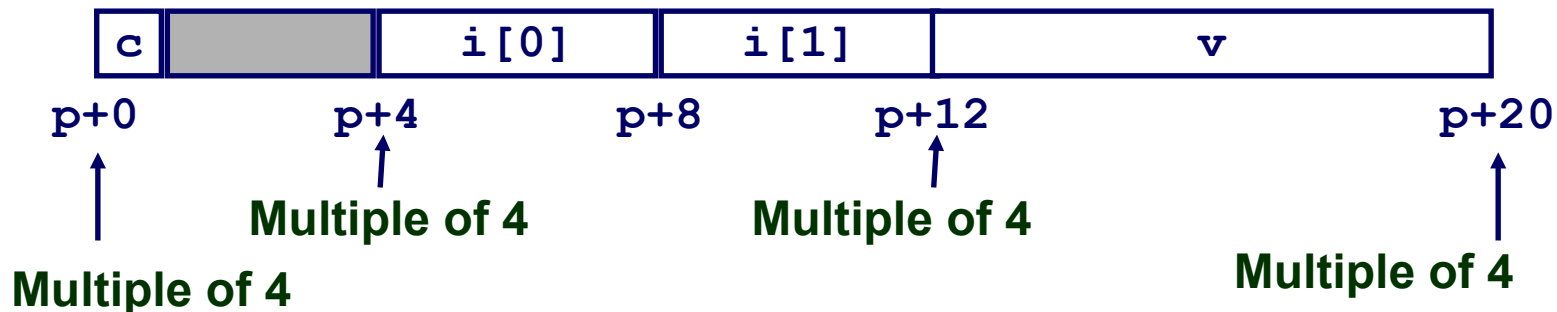
Windows (including Cygwin):


- $K = 8$ , due to `double` element



Linux:

- $K = 4$ ; `double` treated like a 4-byte data type





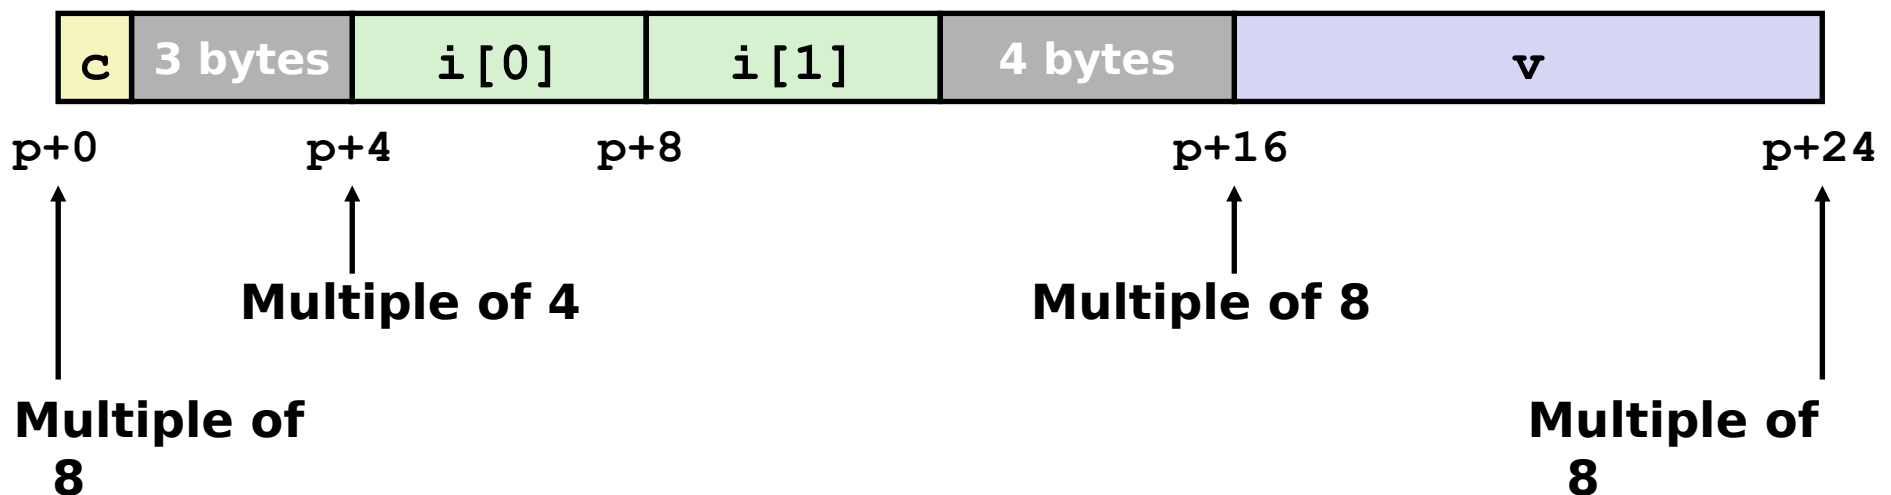
# Specific Cases of Alignment (x86-64)

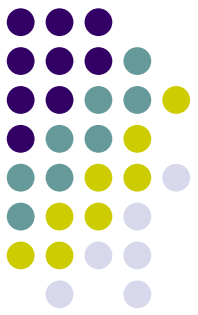
- **1 byte: `char`, ...**
  - no restrictions on address
- **2 bytes: `short`, ...**
  - lowest 1 bit of address must be  $0_2$
- **4 bytes: `int`, `float`, ...**
  - lowest 2 bits of address must be  $00_2$
- **8 bytes: `double`, `long`, `char *`, ...**
  - lowest 3 bits of address must be  $000_2$
- **16 bytes: `long double` (GCC on Linux)**
  - lowest 4 bits of address must be  $0000_2$

# Satisfying Alignment with Structures

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- **Within structure:**
  - Must satisfy each element's alignment requirement
- **Overall structure placement**
  - Each structure has alignment requirement  $K$ 
    - $K$  = Largest alignment of any element
  - Initial address & structure length must be multiples of  $K$
- **Example:**
  - $K = 8$ , due to double element

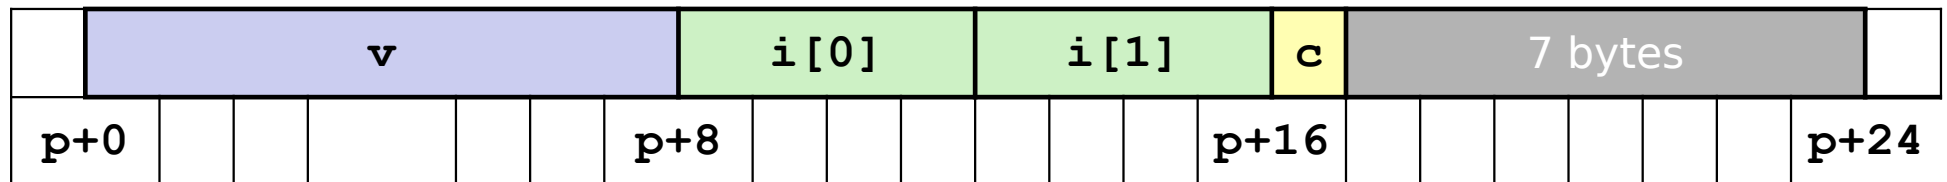




# Meeting Overall Alignment Requirement

- For largest alignment requirement  $K$
- Overall structure must be multiple of  $K$

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



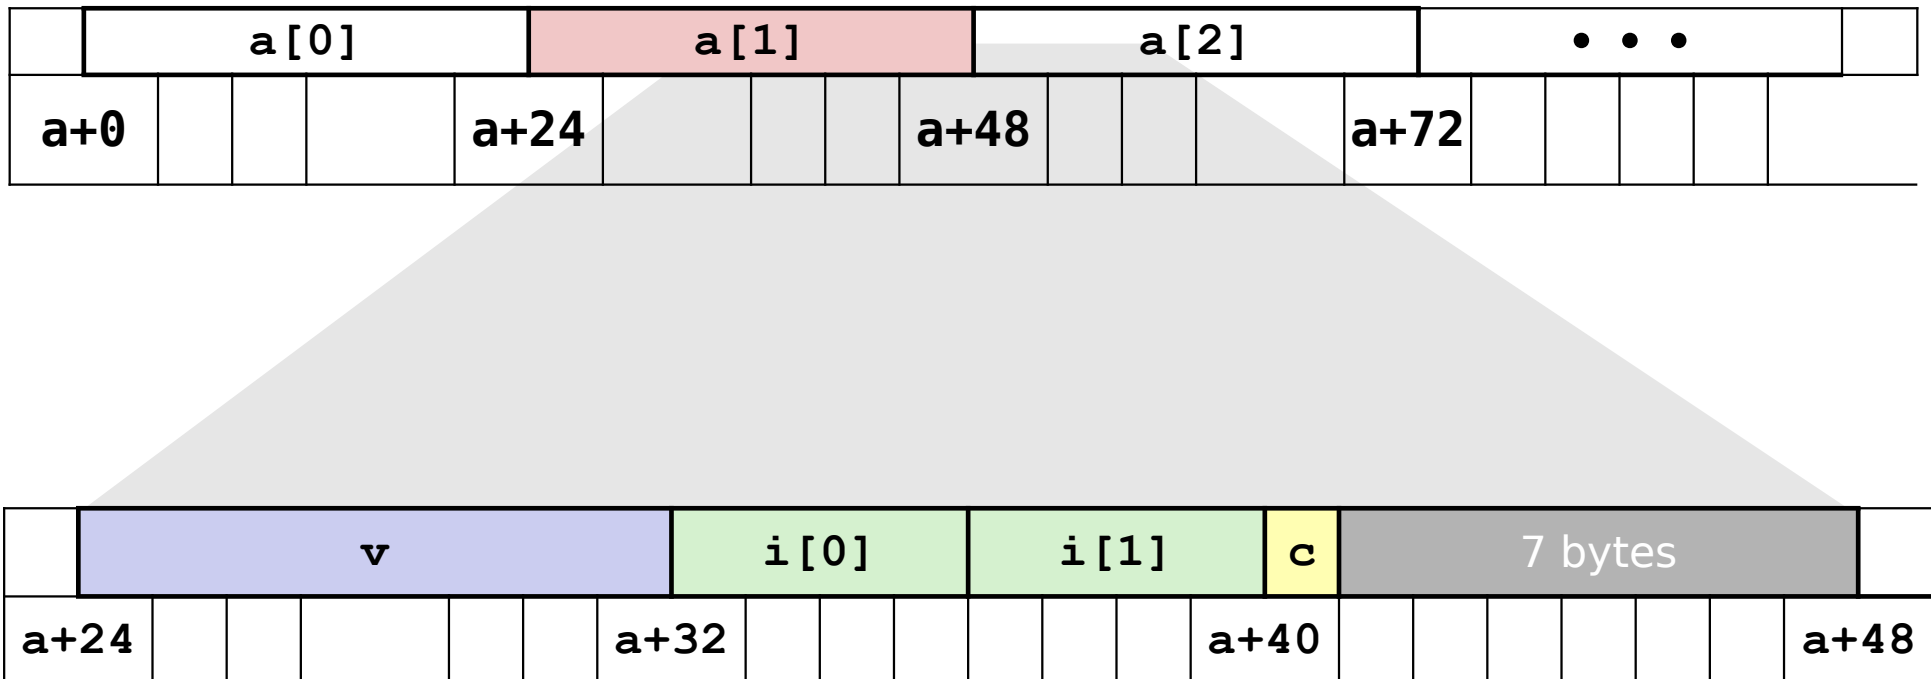
Multiple of  $K=8$



# Arrays of Structures

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```

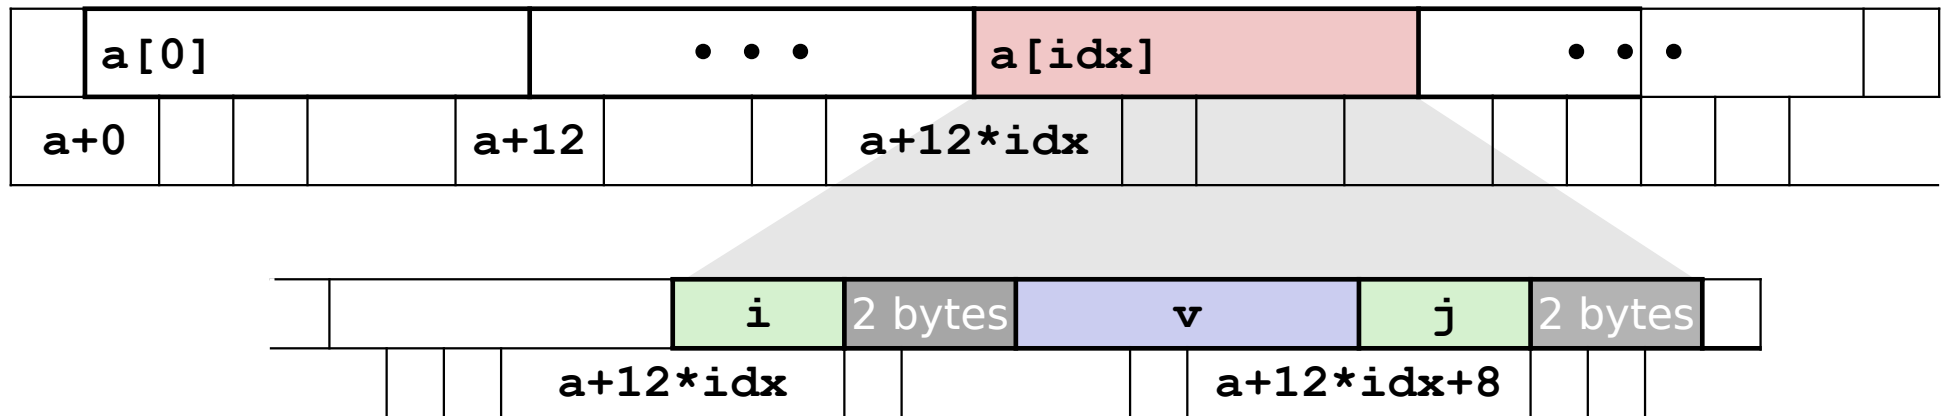
- Overall structure length multiple of K
- Satisfy alignment requirement for every element



# Accessing Array Elements

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```

- Compute array offset  $12 \cdot \text{idx}$ 
  - `sizeof(S3)`, including alignment spacers
- Element **j** is at offset 8 within structure
- Assembler gives offset **a + 8**
  - Resolved during linking



```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %rdi = idx  
leaq (%rdi,%rdi,2),%rax # 3*idx  
movzwl a+8(,%rax,4),%eax
```




# Saving Space

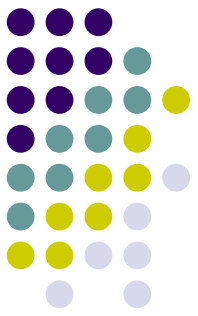
- Put large data types first
- Effect (K=4)

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```

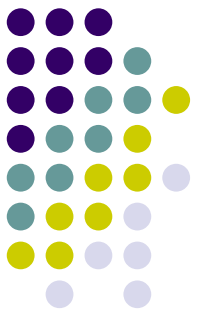


```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```





# שאלות?



# Optional Materials


# Background

## ■ History

- x87 FP
  - Legacy, very ugly
- SSE FP
  - Supported by Shark machines
  - Special case use of vector instructions
- AVX FP
  - Newest version
  - Similar to SSE
  - Documented in book

# Programming with SSE3

## XMM Registers: 16 total




■ Each 16 bytes


■ 16 single-byte integers



■ 8 16-bit integers



■ 4 32-bit integers




■ 4 single-precision floats



■ 2 double-precision floats



■ 1 single-precision float




■ 1 double-precision float



# Scalar & SIMD Operations


■ Scalar Operations: Single Precision

`addss %xmm0, %xmm1`




■ SIMD Operations: Single Precision

`addps %xmm0, %xmm1`

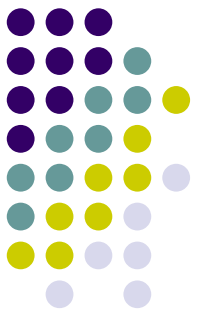


■ Scalar Operations: Double Precision

`addsd %xmm0, %xmm1`







# FP Basics

- Arguments passed in `%xmm0`, `%xmm1`, ...
- Result returned in `%xmm0`
- All XMM registers caller-saved


```
float fadd(float x, float y)
{
    return x + y;
}
```

```
double dadd(double x, double y)
{
    return x + y;
}
```

```
# x in %xmm0, y in %xmm1
addss    %xmm1, %xmm0
ret
```

```
# x in %xmm0, y in %xmm1
addsd    %xmm1, %xmm0
ret
```


# FP Memory Referencing



- Integer (and pointer) arguments passed in regular registers
- FP values passed in XMM registers
- Different mov instructions to move between XMM registers, and between memory and XMM registers

```
double dincr(double *p, double v)
{
    double x = *p;
    *p = x + v;
    return x;
}
```

```
# p in %rdi, v in %xmm0
movapd  %xmm0, %xmm1    # Copy v
movsd   (%rdi), %xmm0    # x = *p
addsd   %xmm0, %xmm1    # t = x + v
movsd   %xmm1, (%rdi)    # *p = t
ret
```



# Other Aspects of FP Code

- **Lots of instructions**
  - Different operations, different formats, ...
- **Floating-point comparisons**
  - Instructions `ucomiss` and `ucomisd`
  - Set condition codes CF, ZF, and PF
- **Using constant values**
  - Set XMM0 register to 0 with instruction `xorpd %xmm0, %xmm0`
  - Others loaded from memory

