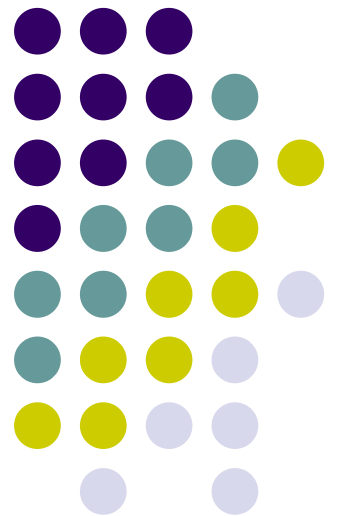



Computer Organization: A Programmer's Perspective

The Memory Hierarchy


Gal A. Kaminka
galk@cs.biu.ac.il






The CPU-Memory Gap

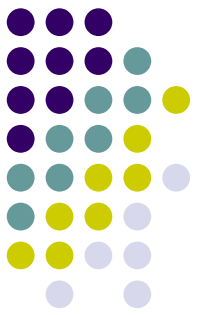
The gap widens between DRAM, disk, and CPU speeds.





Key Question: How to have fast, cheap memory?

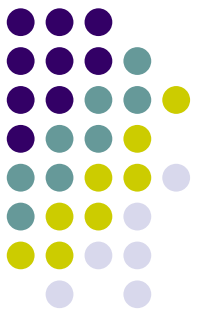
Principle of Locality



Programs tend to reuse data and instructions:

- near those they have used recently
- Or same as those they have used recently
- **Temporal locality:** Recently referenced items are likely to be referenced in the near future.
- **Spatial locality:** Items with nearby addresses tend to be referenced close together in time.

Example



```
sum = 0;
for (i = 0; i < n; i++)
    sum += a[i];
return sum;
```

Locality Example:

Data

Reference array elements in succession (stride-1 reference pattern):

Spatial locality

Reference `sum` each iteration:

Instructions

Temporal locality


Reference instructions in sequence:

Cycle through loop repeatedly:

Spatial locality

Temporal locality

Locality Qualitative Estimation



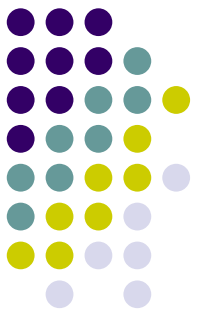
Question: Does this function have good locality?

```
int sumarray(int a[M][M])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < M; j++)
            sum += a[i][j];

    return sum
}
```

Locality Example



Question: Does this function have good locality?

```
int sumarray(int a[M][M])
{
    int i, j, sum = 0;

    for (j = 0; j < M; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];

    return sum
}
```

Important Skill for Professional Programmer:


Be able to look at code, get a qualitative sense of its locality

Memory Hierarchies



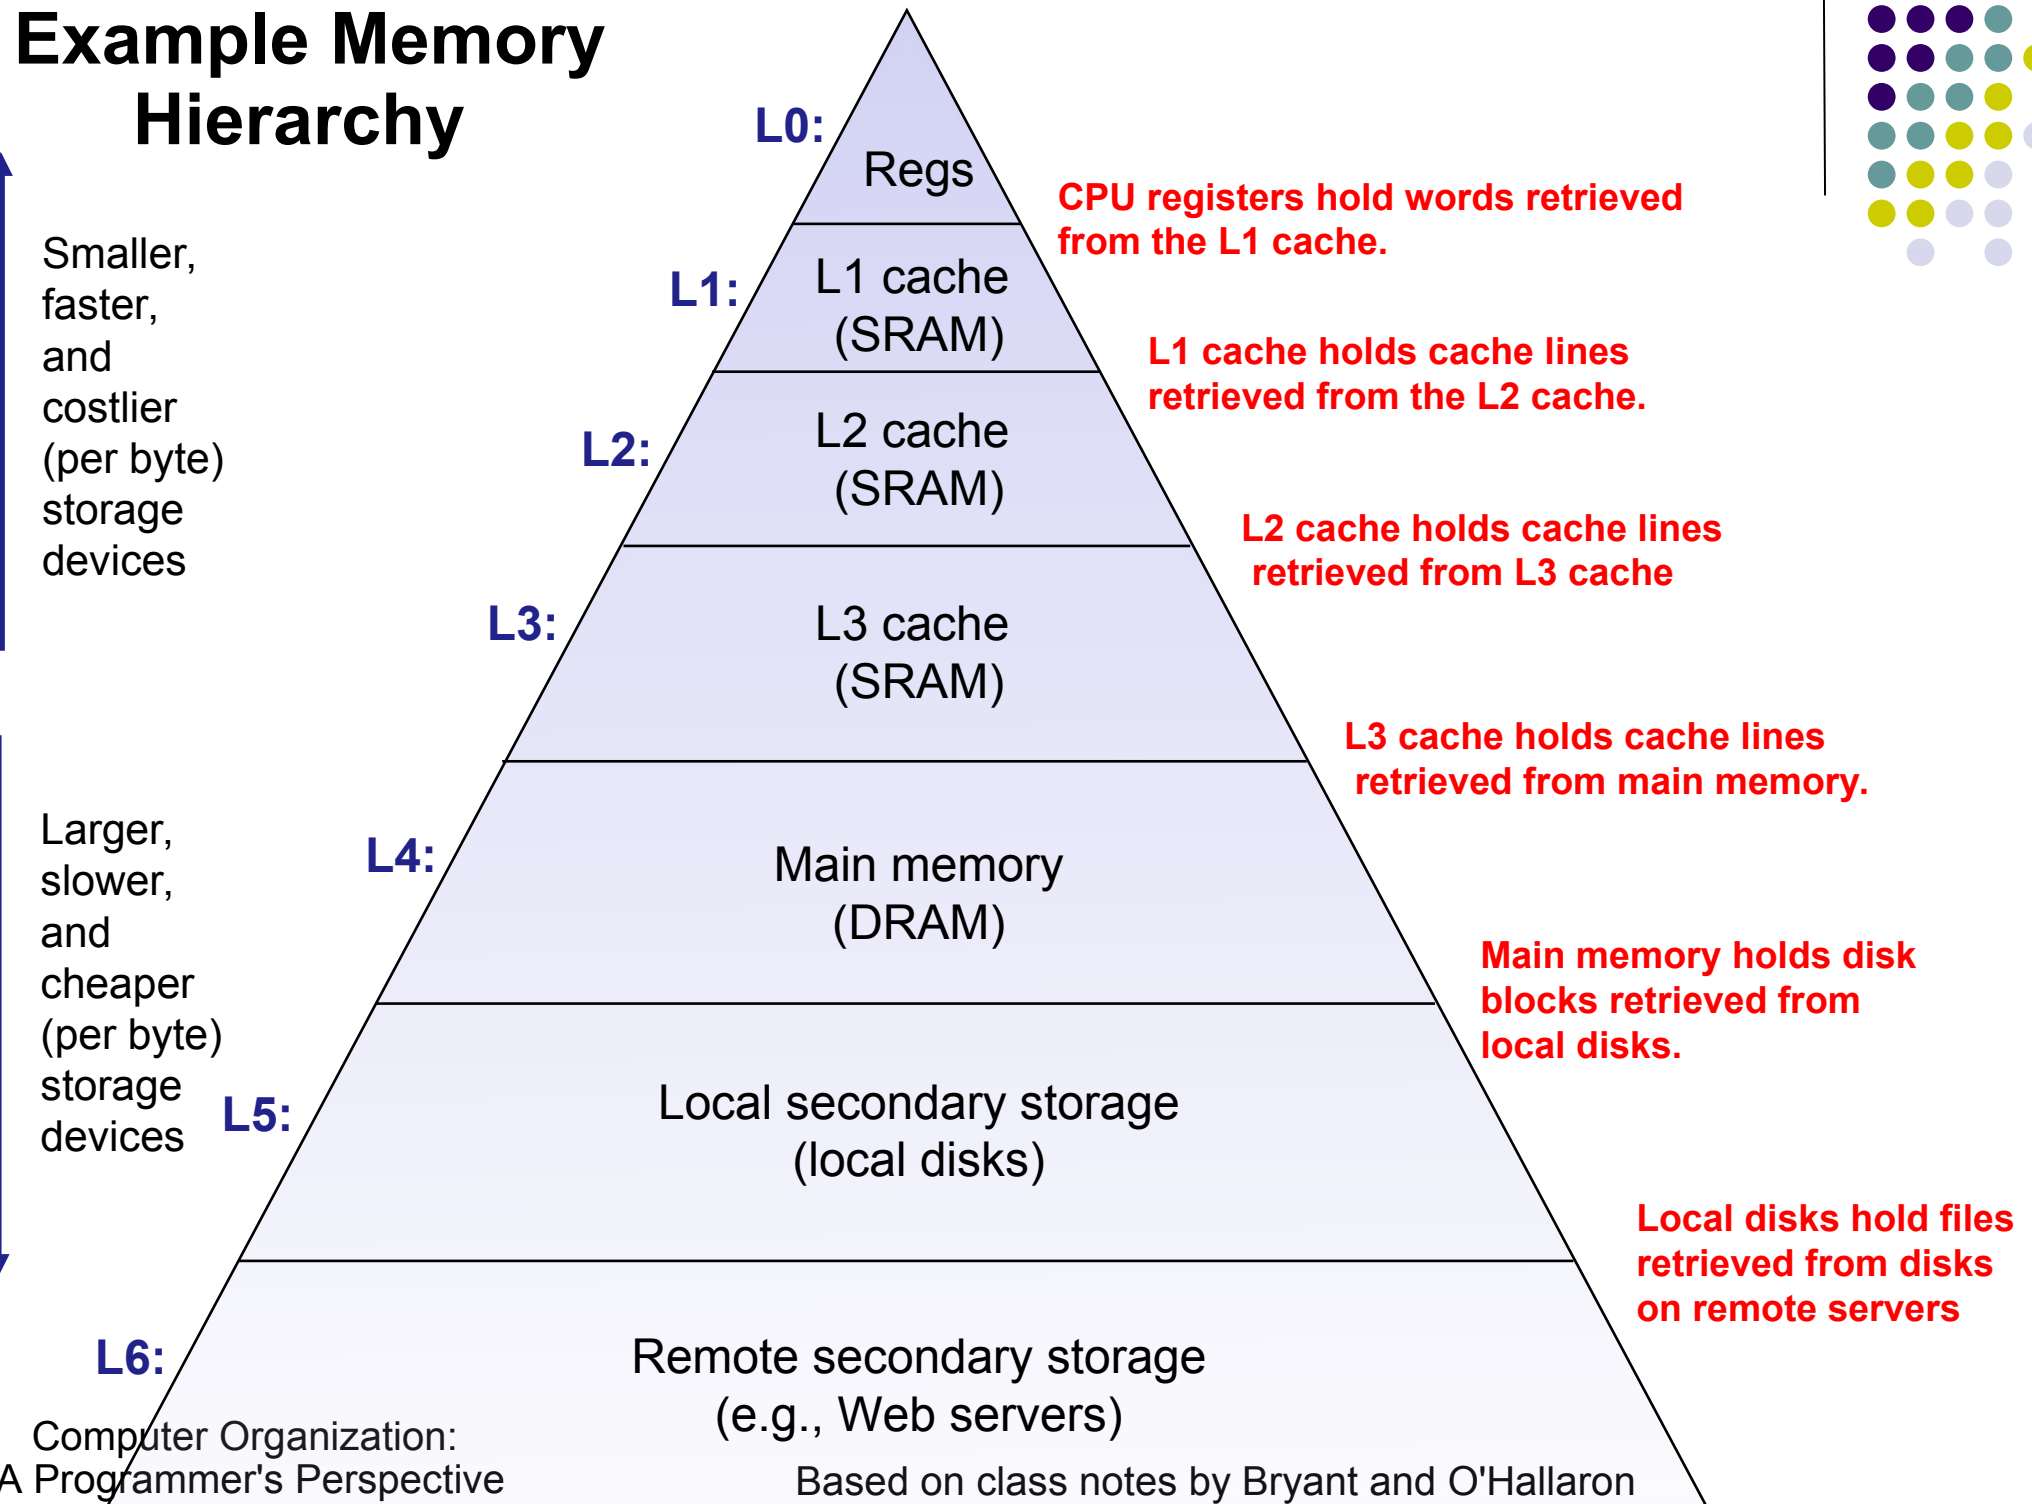
- Some fundamental and enduring properties of hardware and software:
 - Fast storage technologies cost more per byte and have less capacity.
 - The gap between CPU and main memory speed is widening.
 - Well-written programs tend to exhibit good locality.
- These properties complement each other beautifully.
- They suggest an approach for organizing memory and storage systems known as a memory hierarchy.

Example Memory Hierarchy



Smaller,
faster,
and
costlier
(per byte)
storage
devices

Larger,
slower,
and
cheaper
(per byte)
storage
devices




Caches




- **Cache:** A smaller, faster storage that acts as a staging area for a subset of the data in a larger, slower device.
- Fundamental idea of a memory hierarchy:
 - For each k , the faster, smaller device at level k serves as a cache for the larger, slower device at level $k+1$.
- Why do memory hierarchies work?
 - Programs tend to access the data at level k more often than they access the data at level $k+1$.
 - Thus, the storage at level $k+1$ can be slower, and thus larger and cheaper per bit.
 - **Net effect:** A large pool of memory that costs as much as the cheap storage near the bottom, but that serves data to programs at the rate of the fast storage near the top.

Caching in a Memory Hierarchy



Level k:




Smaller, faster, more expensive device at level k caches a subset of the blocks from level k+1



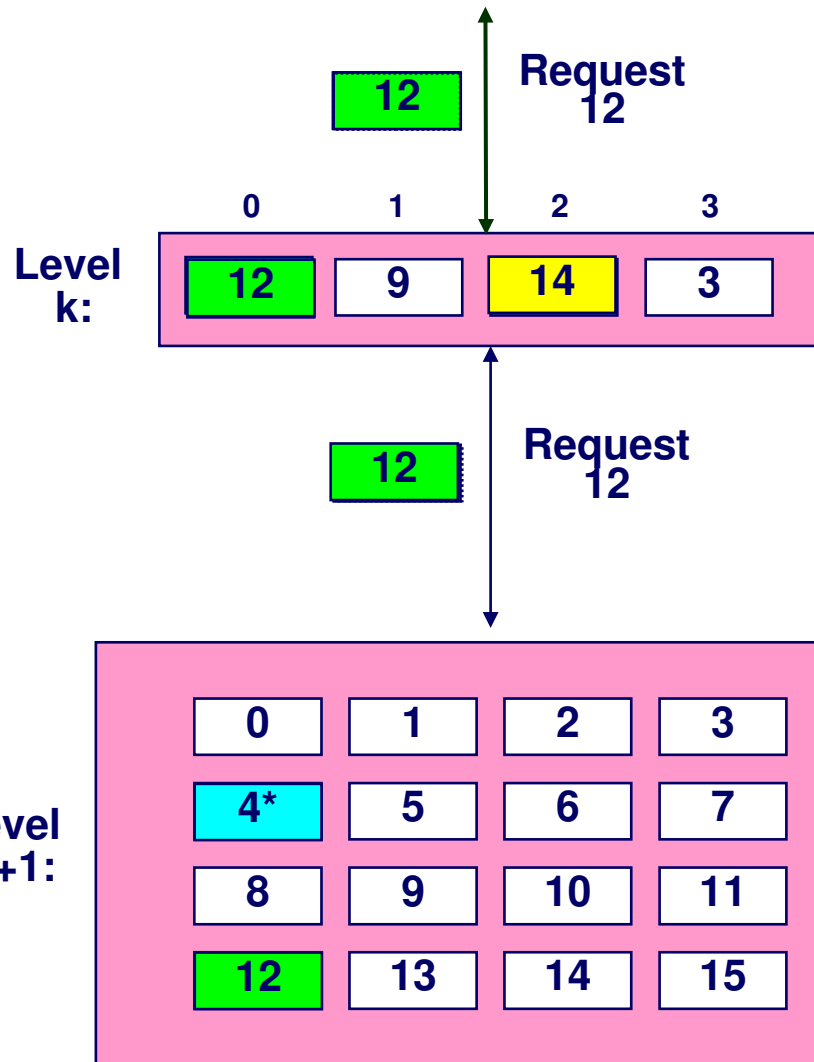
Data is copied between levels in block-sized transfer units

Level k+1:



Larger, slower, cheaper storage device at level k+1 is partitioned into blocks.

General Caching Concepts



Program needs object d, which is stored in some block b.

Cache hit

Program finds b in the cache at level k. E.g., block 14.

Cache miss

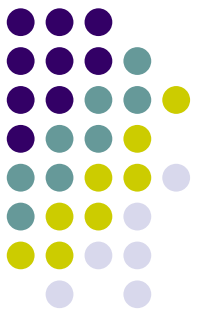
b is not at level k, so level k cache must fetch it from level k+1. E.g., block 12.

If level k cache is full, then some current block must be replaced (evicted). Which one is the “victim”?

Placement policy: where can the new block go? E.g., $b \bmod 4$

Replacement policy: which block should be evicted? E.g., LRU

General Caching Concepts



Types of cache misses:

- **Cold (compulsary) miss**
 - Cold misses occur because the cache is empty.
- **Conflict miss**
 - Most caches limit blocks at level $k+1$ to a small subset (sometimes a singleton) of the block positions at level k .
 - e.g. Block i at level $k+1$ is placed in block $(i \bmod 4)$ at level $k+1$.
 - Conflict misses occur when the level k cache is large enough, but multiple data objects all map to the same level k block.
 - E.g. Referencing blocks 0, 8, 0, 8, 0, 8, ... would miss every time.
- **Capacity miss**
 - Occurs when the set of active cache blocks (working set) is larger than the cache.

Cache Memories

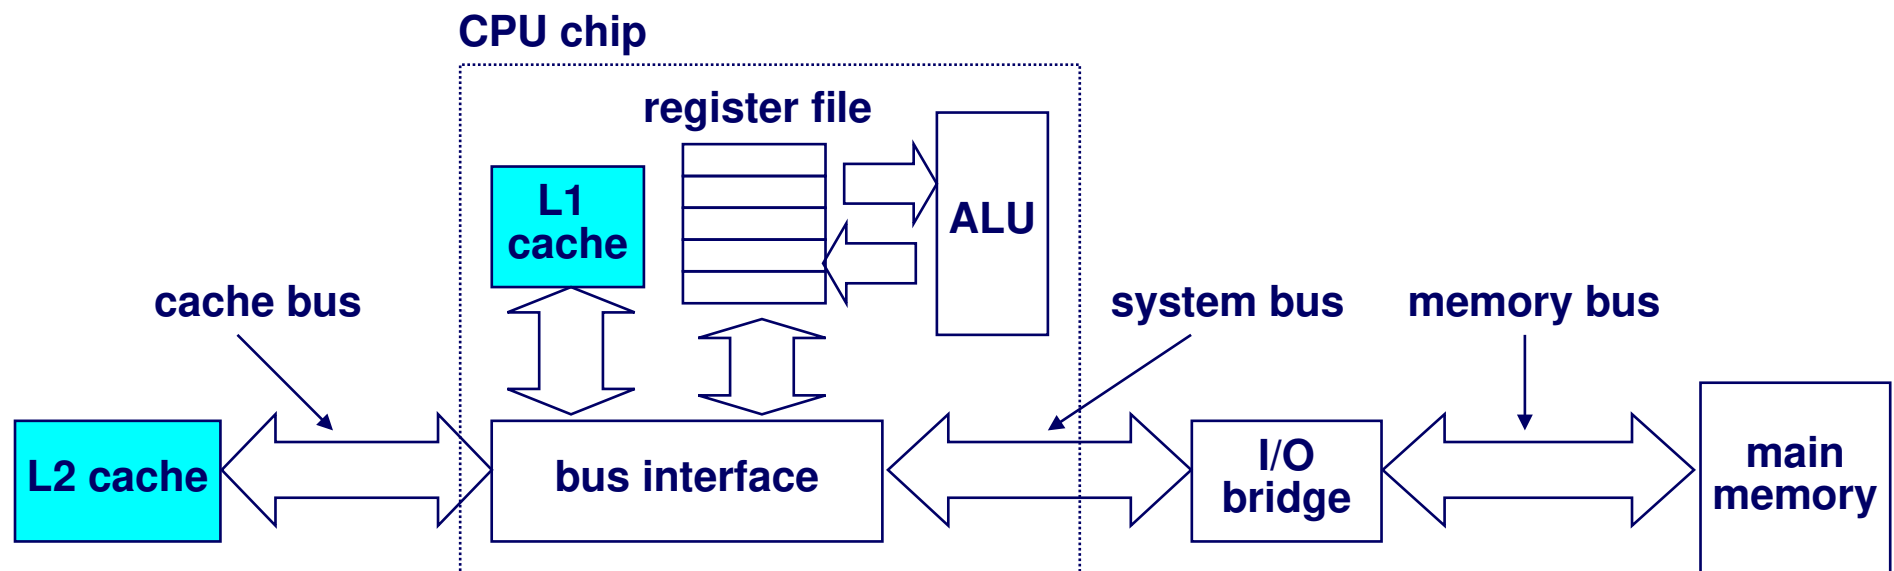



Cache memories are small, fast SRAM-based memories managed automatically in hardware.

Hold frequently accessed blocks of main memory

CPU looks first for data in L1, then in L2, then in main memory.

Typical bus structure:






General Organization of a Cache Memory



Cache is an array of sets.

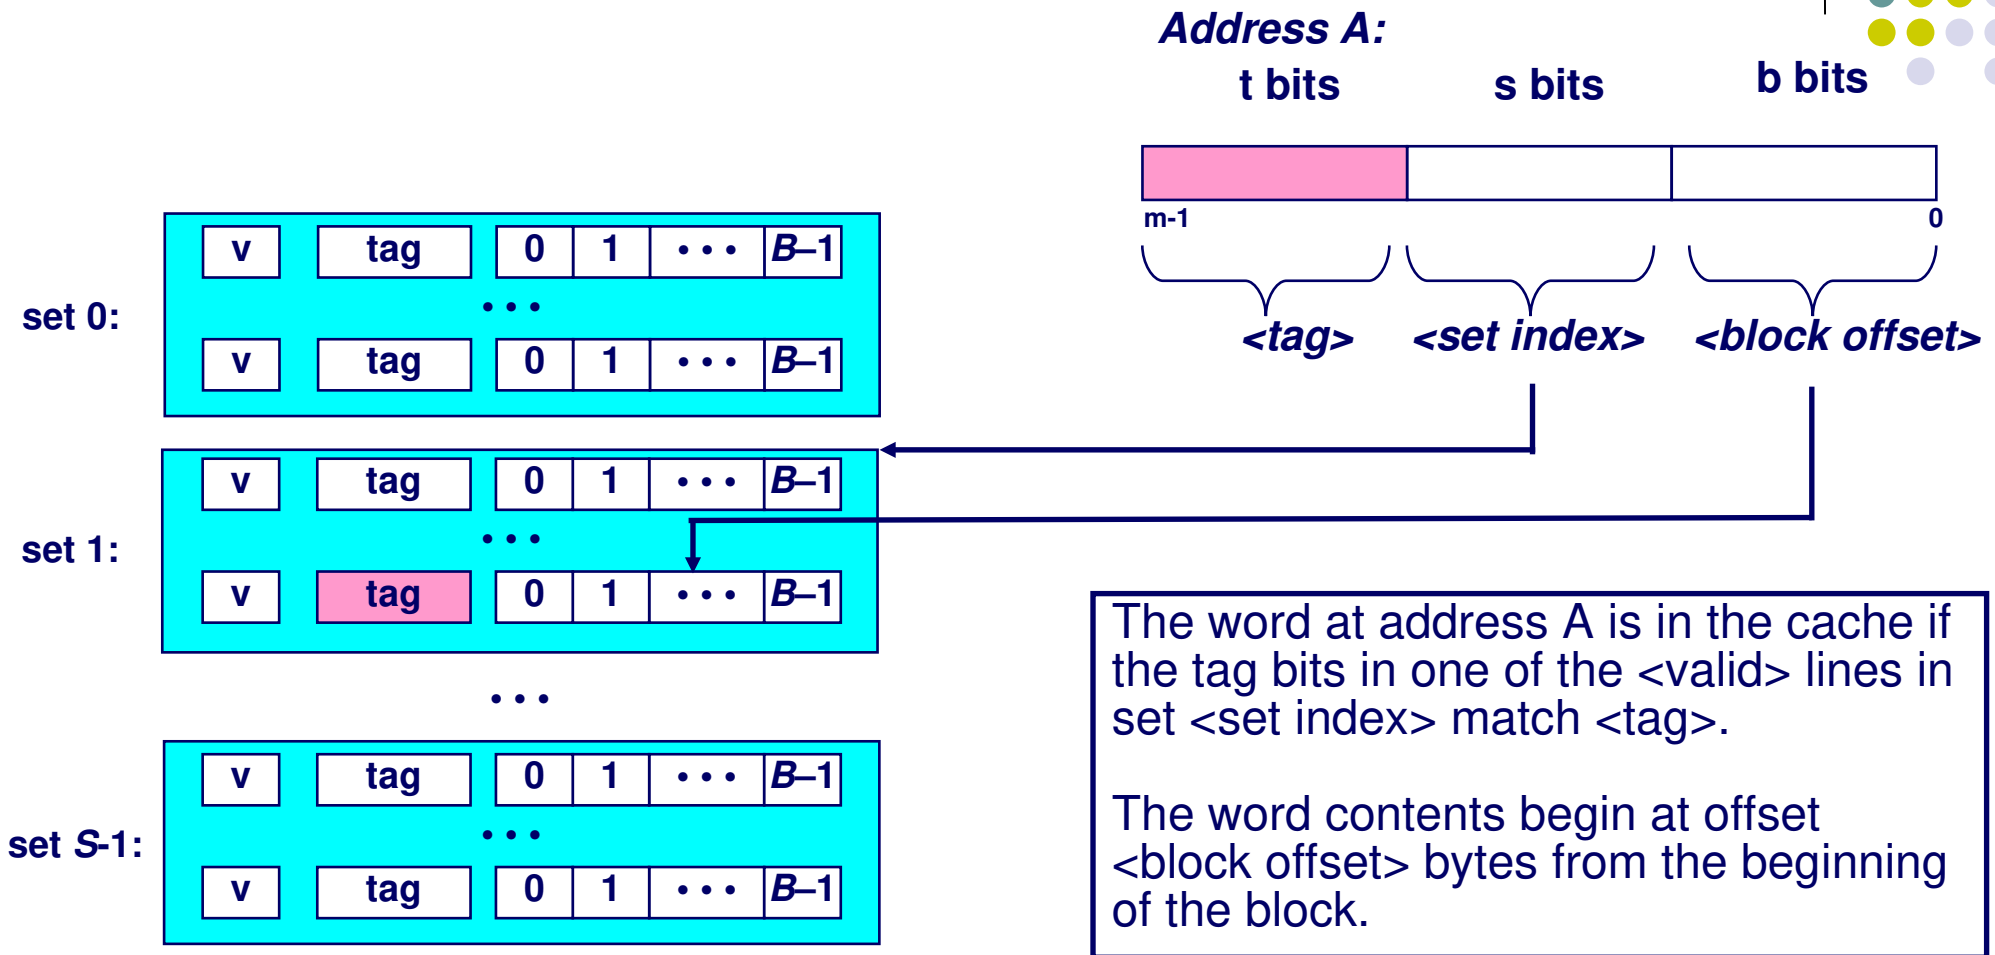

Each set contains one or more lines.

Each line holds a block of data.

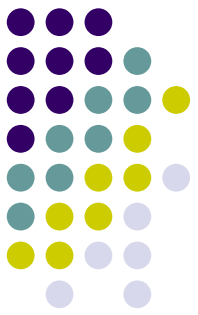


Cache size: $C = B \times E \times S$ data bytes

Addressing Caches




Direct-Mapped Cache

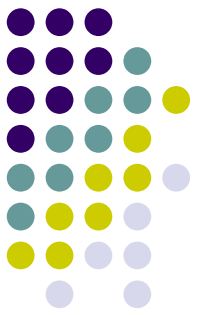


Simplest kind of cache

Characterized by exactly one line per set ($E=1$).

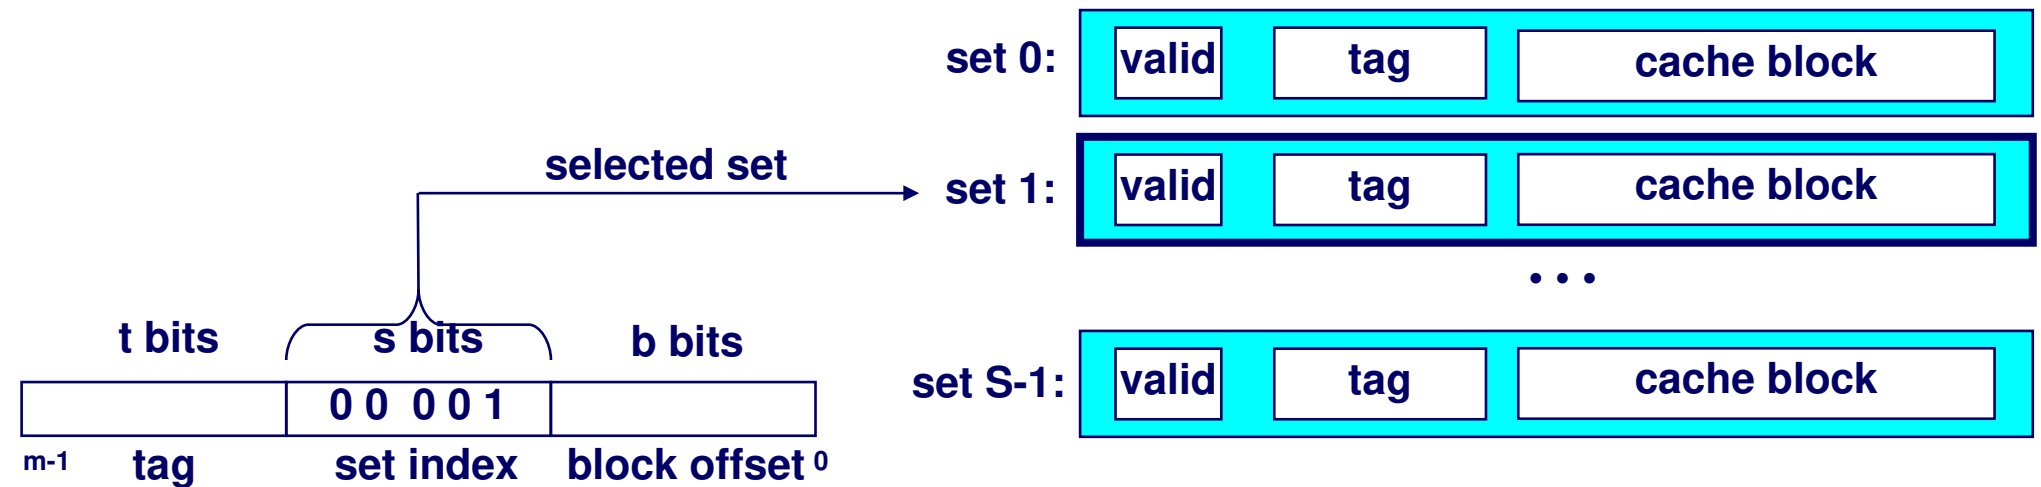


Accessing Direct-Mapped Caches

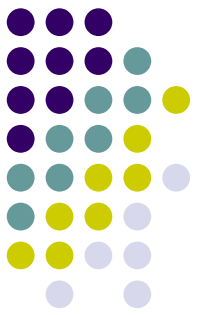


Set selection

Use the set index bits to determine the set of interest.




Accessing Direct-Mapped Caches



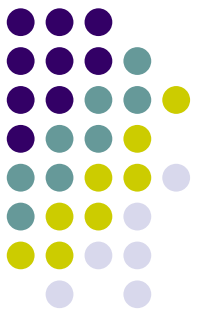
Line matching and word selection

Line matching: Find a valid line in the selected set with a matching tag

Word selection: Then extract the word (here 32bits made from 4 bytes $B_0..B_3$)



Direct-Mapped Cache Simulation



M=16 byte addresses, B=2 bytes/block,
S=4 sets, E=1 entry/set

t=1 s=2 b=1

x	xx	x
---	----	---

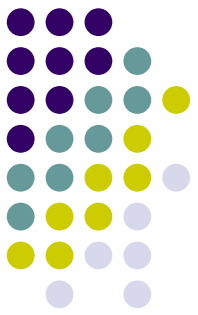
Address trace (reads):

0 [0000₂], 1 [0001₂], 13 [1101₂], 8 [1000₂], 0 [0000₂]

0 [0000₂] (*cold miss*)

	v	tag	data
(1)			

Direct-Mapped Cache Simulation



M=16 byte addresses, B=2 bytes/block,
S=4 sets, E=1 entry/set

t=1 s=2 b=1

x	xx	x
---	----	---

Address trace (reads):

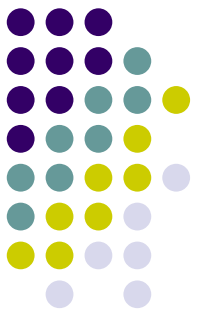
0 [0000₂], 1 [0001₂], 13 [1101₂], 8 [1000₂], 0 [0000₂]

0 [0000₂] (*cold miss*)

(1)

v	tag	data	
1	0	M[0]	M[1]

Direct-Mapped Cache Simulation



M=16 byte addresses, B=2 bytes/block,
S=4 sets, E=1 entry/set

t=1 s=2 b=1

x	xx	x
---	----	---

Address trace (reads):

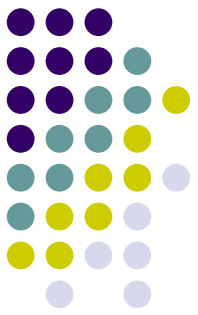
0 [0000₂], 1 [0001₂], 13 [1101₂], 8 [1000₂], 0 [0000₂]

1 [0001₂] (*hit!*)

v	tag	data	
1	0	M[0]	M[1]

(2)

Direct-Mapped Cache Simulation



M=16 byte addresses, B=2 bytes/block,
S=4 sets, E=1 entry/set

t=1 s=2 b=1

x	xx	x
---	----	---

Address trace (reads):

0 [0000₂], 1 [0001₂], **13 [1101₂]**, 8 [1000₂], 0 [0000₂]

(1)

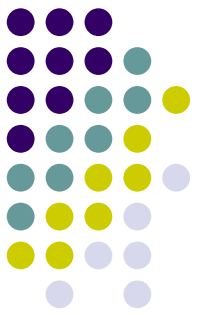
v	tag	data	
1	0	M[0]	M[1]

13 [1101₂] (*cold miss*)

(3)

v	tag	data	
1	0	M[0]	M[1]
1	1	M[12]	M[13]

Direct-Mapped Cache Simulation



M=16 byte addresses, B=2 bytes/block,
S=4 sets, E=1 entry/set

t=1 s=2 b=1

x	xx	x
---	----	---

Address trace (reads):

0 [0000₂], 1 [0001₂], 13 [1101₂], 8 [1000₂], 0 [0000₂]

(1)

v	tag	data
1	0	M[0] M[1]

13 [1101₂] (*cold miss*)

(3)

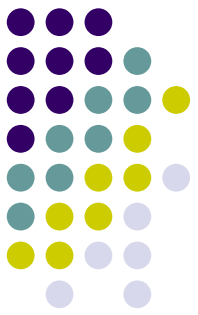
v	tag	data
1	0	M[0] M[1]
1	1	M[12] M[13]

8 [1000₂] (*conflict miss*)

(4)

v	tag	data
1	1	M[8] M[9]
1	1	M[12] M[13]

Direct-Mapped Cache Simulation



M=16 byte addresses, B=2 bytes/block,
S=4 sets, E=1 entry/set

t=1 s=2 b=1

x	xx	x
---	----	---

Address trace (reads):

0 [0000₂], 1 [0001₂], 13 [1101₂], 8 [1000₂], 0 [0000₂]

(1)

v	tag	data
1	0	M[0] M[1]

13 [1101₂] (*miss*)

(3)

v	tag	data
1	0	M[0] M[1]
1	1	M[12] M[13]

8 [1000₂] (*conflict miss*)

(4)

v	tag	data
1	1	M[8] M[9]
1	1	M[12] M[13]

0 [0000₂] (*conflict miss*)


(5)

v	tag	data
1	0	M[0] M[1]
1	1	M[12] M[13]

Why Use Middle Bits as Index?



4-line Cache



High-Order Bit Indexing

Adjacent memory lines would map to same cache entry


Poor use of spatial locality

Middle-Order Bit Indexing


Consecutive memory lines map to different cache lines

Can hold C-byte region of address space in cache at one time


High-Order Bit Indexing



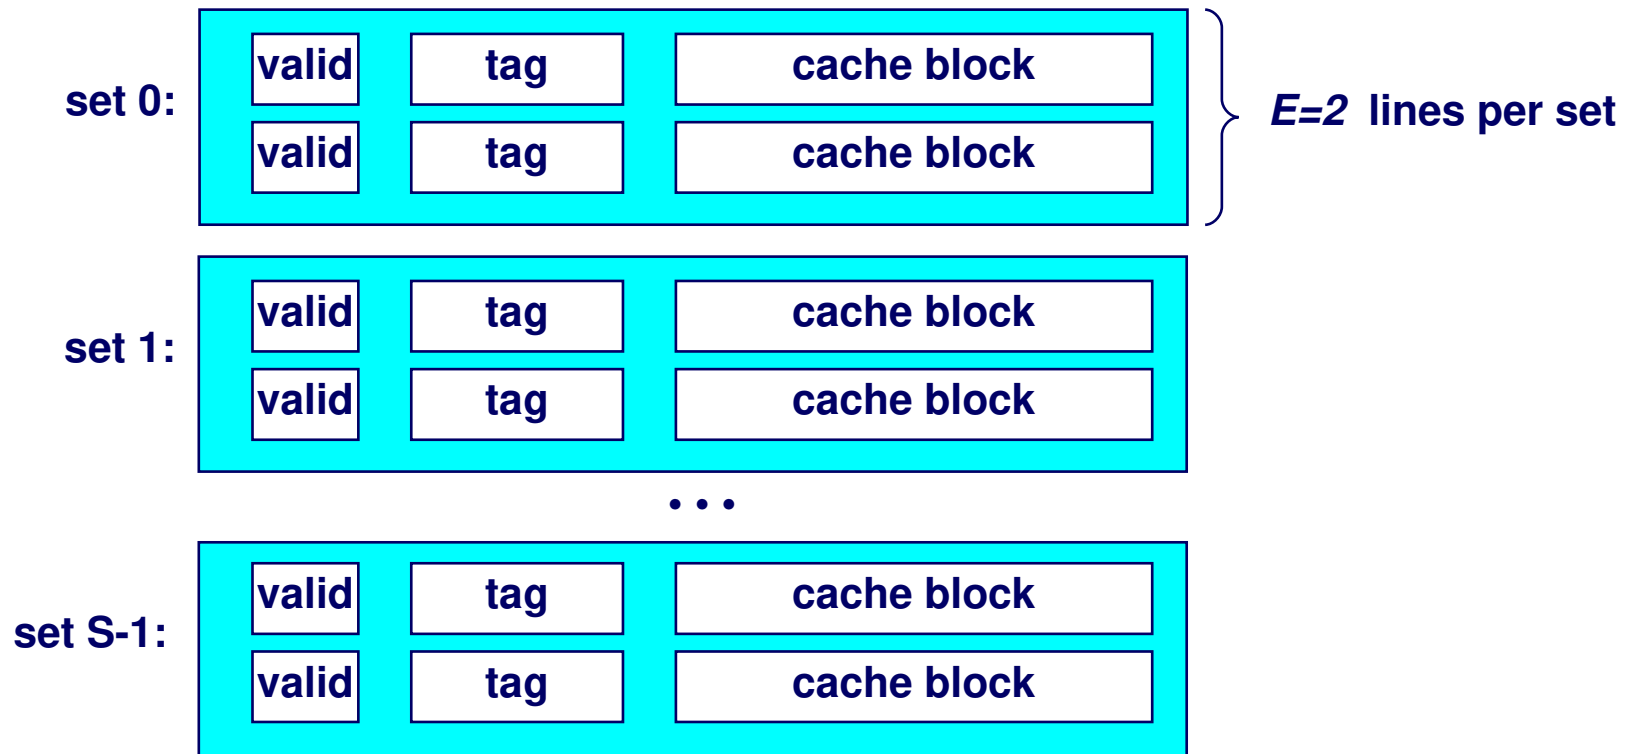
Middle-Order Bit Indexing



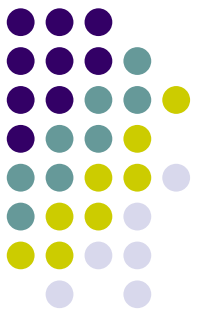
Set Associative Caches



Characterized by more than one line per set

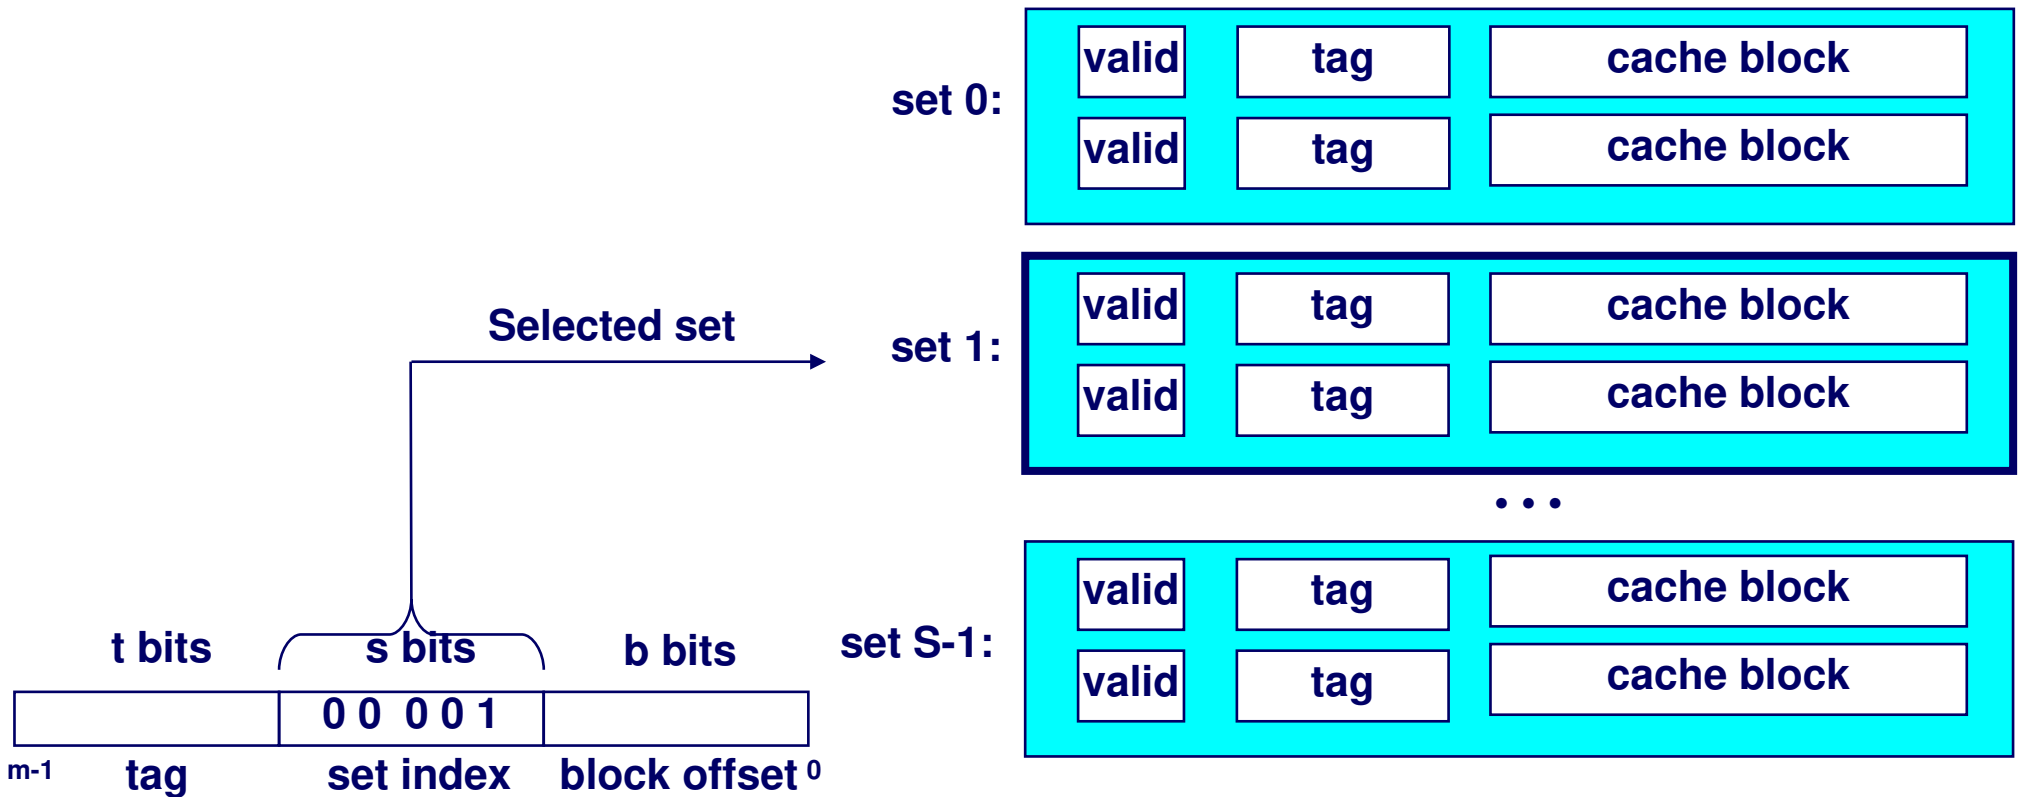


Accessing Set Associative Caches

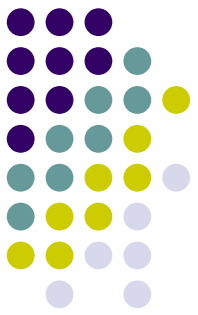


Set selection

identical to direct-mapped cache

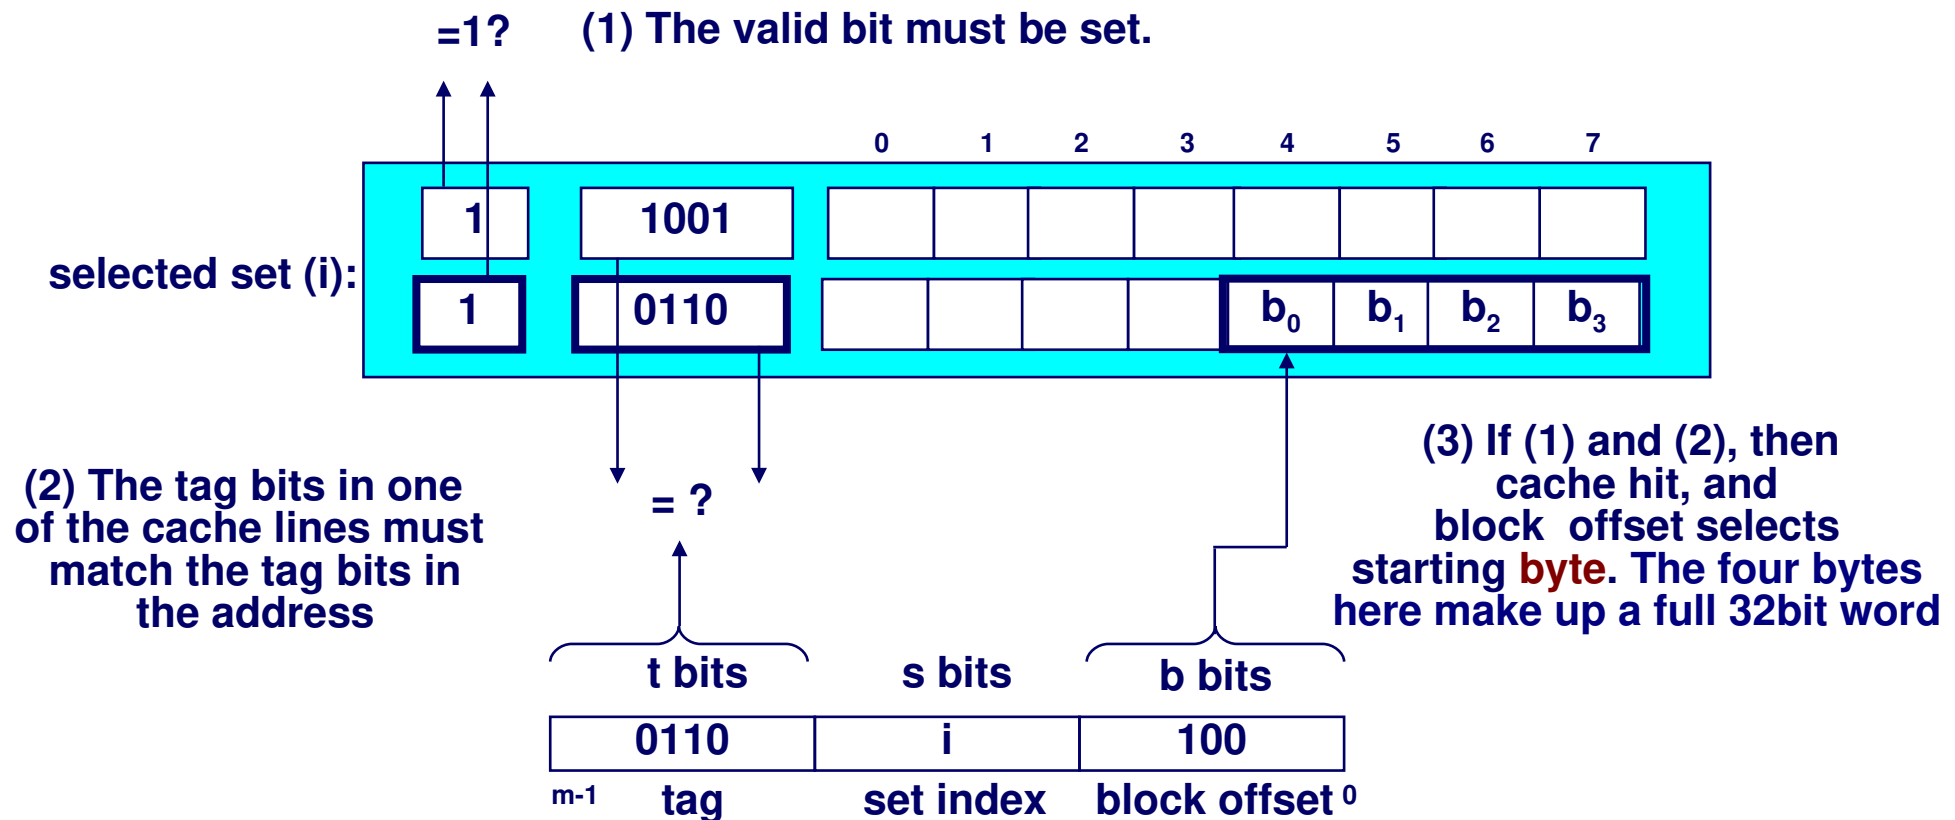


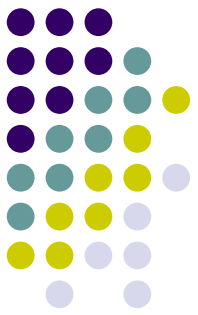
Accessing Set Associative Caches



Line matching and word selection

must compare the tag in each valid line in the selected set.





What about writes?

■ Multiple copies of data exist:

- L1, L2, L3, Main Memory, Disk

■ What to do on a write-hit?

- **Write-through** (write immediately to one level down)
- **Write-back** (defer write to one level down until replacement of line)
 - Need a dirty bit (line different from memory or not)

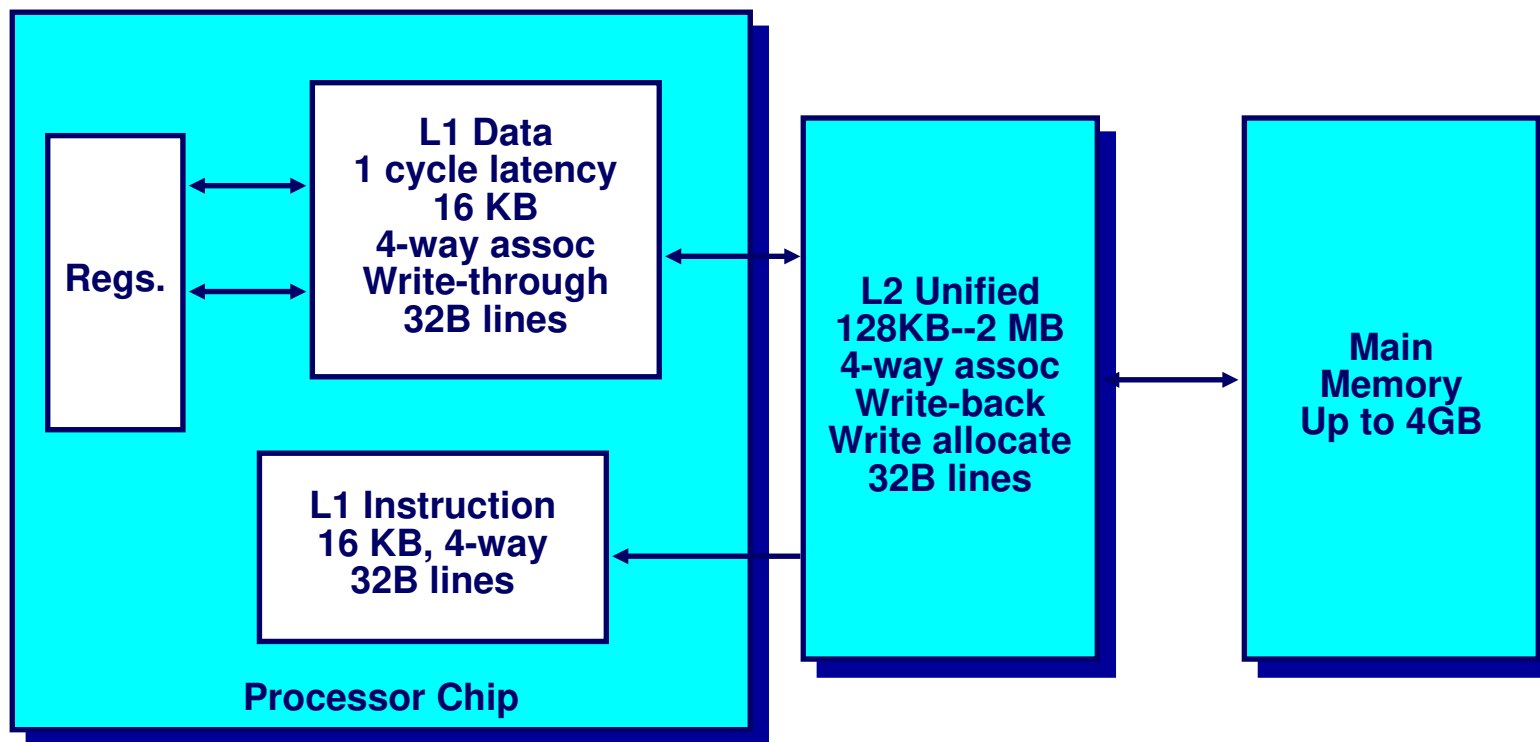

■ What to do on a write-miss?

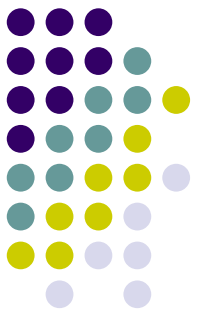
- **Write-allocate** (load into cache, update line in cache)
 - Good if more writes to the location follow
- **No-write-allocate** (writes straight to one level down, does not load into cache)

■ Typical

- Write-through + No-write-allocate
- **Write-back + Write-allocate**

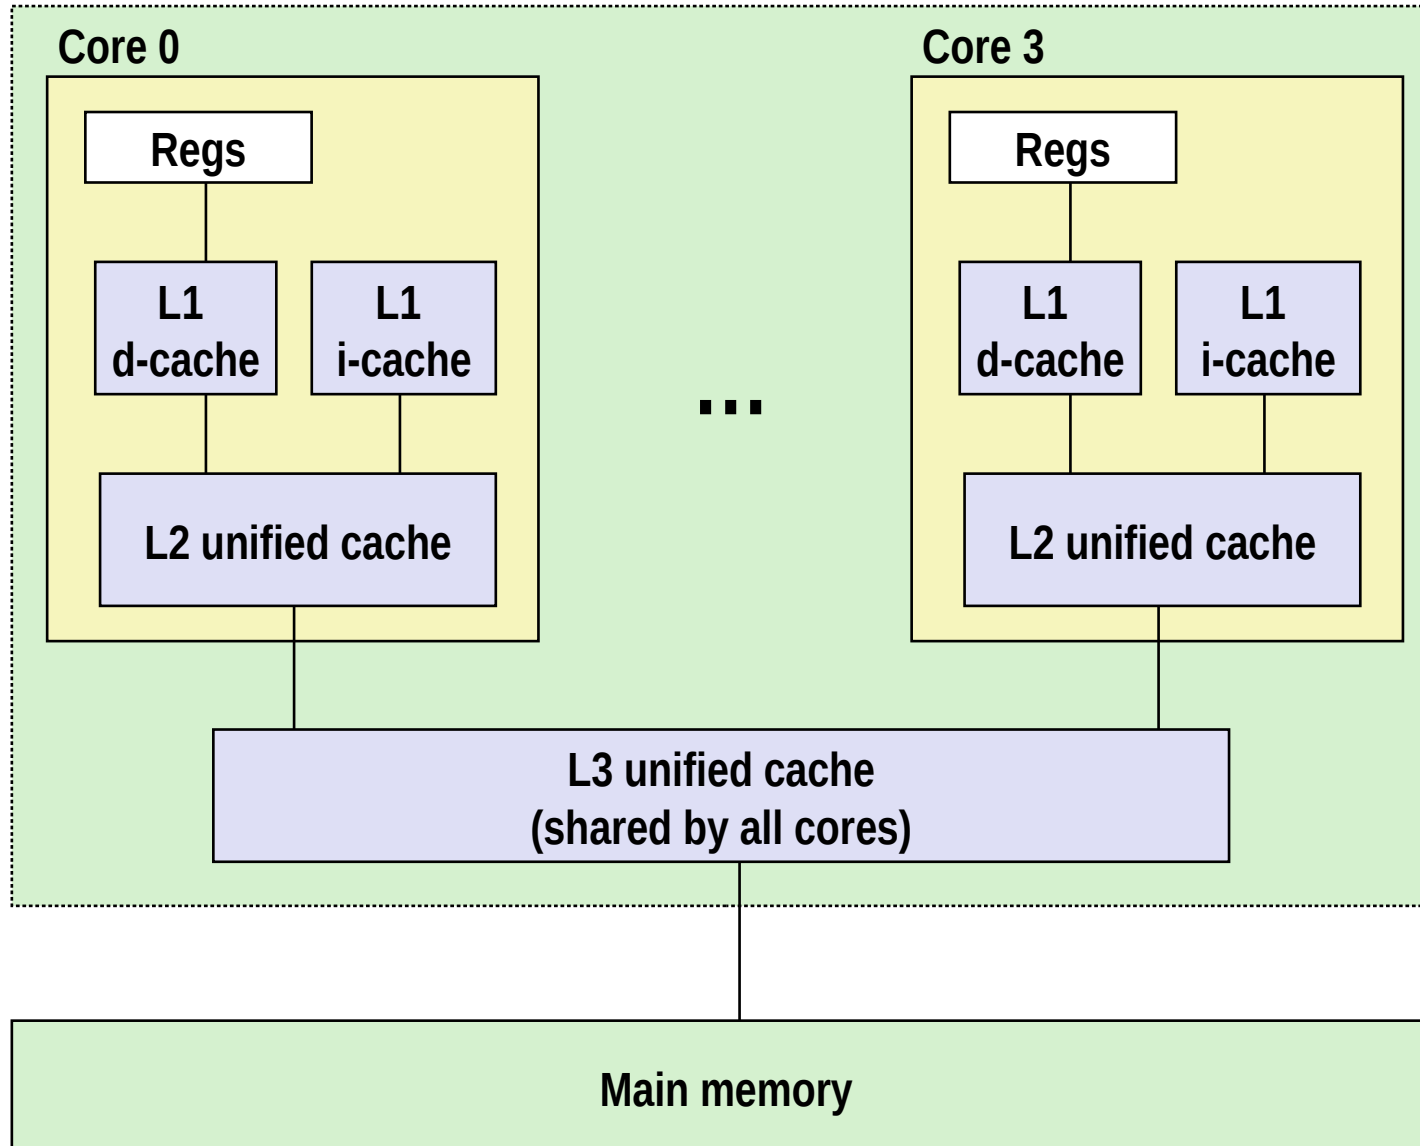
Intel Pentium Cache Hierarchy





Intel Core i7 Cache Hierarchy

Processor package



L1 i-cache and d-cache:

32 KB, 8-way,
Access: 4 cycles

L2 unified cache:

256 KB, 8-way,
Access: 10 cycles

L3 unified cache:

8 MB, 16-way,
Access: 40-75 cycles

Block size: 64 bytes
for all caches.

