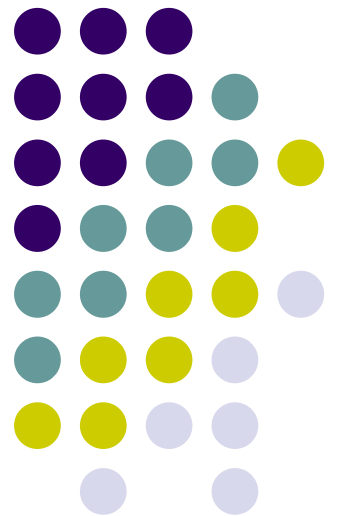



Computer Organization: A Programmer's Perspective

Profiling

Gal A. Kaminka
galk@cs.biu.ac.il



Profiling: Performance Analysis



Performance Analysis (“Profiling”)

Understanding the run-time behavior of programs

What parts are executed, when, for how long

What parts require improvement, optimization


Profiling Programs

Tools of the trade

Granularity of profiling (modules, functions, instructions, ...)

Performance measurement

Components of Performance



Run time

How long does it take to compute?

Memory

How much memory does it take?

These are issues that affect the above components:


Input/Output (I/O)

How much access to external devices, services?

System calls

Parallelization of tasks

Measurement Challenge



How Much Time Does Program X Require?

CPU time

- How many total seconds are used when executing X?

- Measure used for most applications

- Small dependence on other system activities

Actual (“Wall-Clock”) Time

- How many seconds elapse between start and completion of X?

- Depends on system load, I/O times, etc.

Confounding Factors

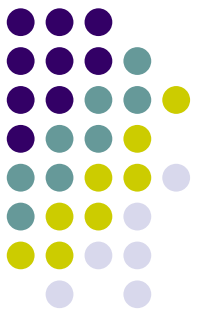
- How does time get measured?

- Many processes share computing resources

 - Transient effects when switching from one process to another

 - The effects of alternating among processes become noticeable


“Time” on a Computer System



real (wall clock) time

 = **user time** (*time executing instructions in the user process*)

 = **system time** (*time executing instructions in kernel on behalf of user process*)


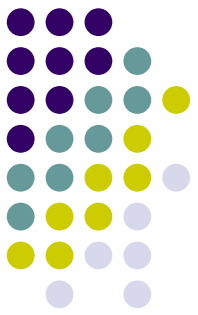
 = **some other user's time** (*time executing instructions in different user's process*)

 +  +  = **real (wall clock) time**

We will use the word “time” to refer to user time.

 **cumulative user time**

Activity Periods: Light Load



Most of the time spent executing one process

Periodic interrupts every 10ms

Interval timer

Keep system from executing one process to exclusion of others

Other interrupts


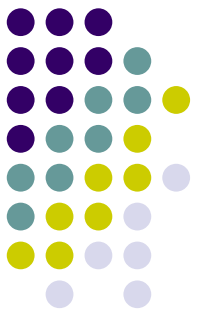
Due to I/O activity

Inactivity periods

System time spent processing interrupts

~250,000 clock cycles

Activity Periods: Heavy Load




Sharing processor with one other active process

From perspective of this process, system appears to be “inactive” for ~50% of the time

Other process is executing

Interval Counting



OS Measures Run-times Using Interval Timer

Maintain 2 counts per process

- User time


- System time

Each time get timer interrupt, increment counter for executing process


- User time if running in user mode

- System time if running in kernel mode

Interval Counting Example




(a) Interval Timings



$$A \quad 110u + 40s$$

$$B \quad 70u + 30s$$

(b) Actual Times




$$A \quad 120.0u + 33.3s$$

$$B \quad 73.3u + 23.3s$$

Unix time Command

(here, timing a “make osevent” command)



```
> time make osevent
```

```
gcc -O2 -Wall -g -march=i486 -c clock.c
```

```
gcc -O2 -Wall -g -march=i486 -c options.c
```

```
gcc -O2 -Wall -g -march=i486 -c load.c
```

```
gcc -O2 -Wall -g -march=i486 -o osevent osevent.c . . .
```

```
0.820u 0.300s 0:01.32 84.8% 0+0k 0+0io 4049pf+0w
```

```
>
```

0.82 seconds user time

82 timer intervals

0.30 seconds system time

30 timer intervals

1.32 seconds wall time

84.8% of total was used running these processes

$(0.82+0.3)/1.32 = .848$

Unix `time` Command

(here, timing a “make osevent” command)



```
> time make osevent
gcc -O2 -Wall -g -march=i486 -c clock.c
gcc -O2 -Wall -g -march=i486 -c options.c
gcc -O2 -Wall -g -march=i486 -c load.c
gcc -O2 -Wall -g -march=i486 -o osevent osevent.c . . .
0.820u 0.300s 0:01.32 84.8% 0+0k 0+0io 4049pf+0w
>
```

0.82 seconds user time

82 timer intervals

0.30 seconds system time

30 timer intervals



1.32 seconds wall time

84.8% of total was used running these processes

$(0.82+0.3)/1.32 = .848$

`time` tells us where the CPU time is spent:
Our code, the system (I/O), or elsewhere

Accuracy of Interval Counting



Minimum

Maximum

Computed time = 70ms

Min Actual = $60 + \epsilon$

Max Actual = $80 - \epsilon$

Average Case Analysis

Over/underestimates tend to balance out


As long as total run time is sufficiently large

Min run time ~1 second

100 timer intervals

Consistently miss 4% overhead due to timer interrupts

The 90/10 Rule of Thumb




90% of execution time is in 10% of code

Lesson:


Find the 10% that really count!
Let the compiler worry about the rest

Important:

First make program work correctly
Make sure easy to maintain
Then optimize




Priority
depends on
project size,
scope,
maturity



Profiling Within Our Code

Profiling at sub-program level



We can measure execution time at:

All functions of a program

- Flat statistics


- Call context statistics

Specific function

- Specific instructions, operations

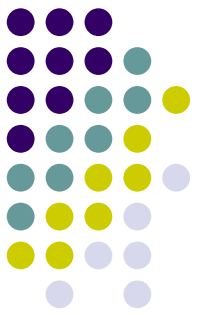
- Memory use, system calls, etc.

Profiling modules (functions)




- Profilers: Tools used to measure run-time performance
 - Frequency and duration of function execution
 - Memory use
- Input:
 - Events (object creation/deletion, thread state, method calls)
 - Instruction counts (how many CPU instructions ran), clock cycles
 - Typically (sampled counts, not accurate)
 - Counters (frequency and duration)
 - Instrumentation
- Output:
 - Execution trace
 - [Statistics](#)

Profile types




- Flat: time spent in each function
 - % time (out of total running time)
 - # of calls made to this function
 - Average, maximum, minimum execution-time per call
 - - Self
 - - Including descendants in call graph
- Call-graph: performance depending on call stack
 - e.g., duration depending on whom was caller, what was passed

Examples of profilers

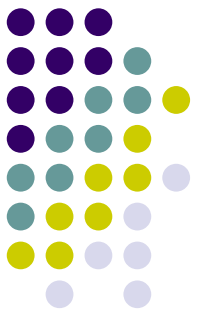


- `gprof` (compiler-assisted instrumentation)
 - Compile (and link) with “-pg” flag
 - During run-time, program will create file “gmon.out”
 - “`gprof <exec> > report.txt`” will generate report
 - Flat and call-graph run-time profiling
- `valgrind` (run-time instrumentation)
 - run “`valgrind <exec>`”, get a report
 - Several tools:
 - Memory leak checker, other memory bugs
 - Cache use profiling
 - Heap memory profiling (who is allocating memory)



Examples of profilers

- cProfile (python)
 - e.g., “`python -m cProfile -o prg.prof prg.py`”
 - Flat and call-graph run-time profiling
- Py-spy (python)
 - e.g., “`py-spy record -o output.svg -pid pid`”
- Visualvm, JDK Mission Control, glowroot (Java)
 - Also many built into different IDEs
- Every professional programmer needs to know profilers!



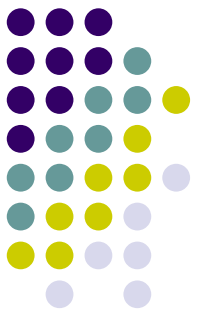
Example *flat* gprof output

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
57.50	0.23	0.23				main
17.50	0.30	0.07	3	23.33	23.33	count4()
7.50	0.33	0.03	2	15.00	38.33	count3()
7.50	0.36	0.03	1	30.00	30.00	count1()
5.00	0.38	0.02	1	20.00	20.00	count()
5.00	0.40	0.02	1	20.00	58.33	count2()

% time spent out of total

Cumulative and self seconds spent

of calls, msec per call (self and total)



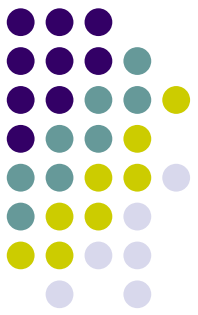
Example *flat* gprof output

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
57.50	0.23	0.23				main
17.50	0.30	0.07	3	23.33	23.33	count4()
7.50	0.33	0.03	2	15.00	38.33	count3()
7.50	0.36	0.03	1	30.00	30.00	count1()
5.00	0.38	0.02	1	20.00	20.00	count()
5.00	0.40	0.02	1	20.00	58.33	count2()

% time spent out of total

Cumulative and self seconds spent

of calls, msec per call (self and total)



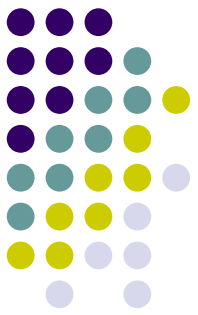
Example *flat* gprof output

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
57.50	0.23	0.23				main
17.50	0.30	0.07	3	23.33	23.33	count4()
7.50	0.33	0.03	2	15.00	38.33	count3()
7.50	0.36	0.03	1	30.00	30.00	count1()
5.00	0.38	0.02	1	20.00	20.00	count()
5.00	0.40	0.02	1	20.00	58.33	count2()

% time spent out of total

Cumulative and self seconds spent

of calls, msec per call (self and total)



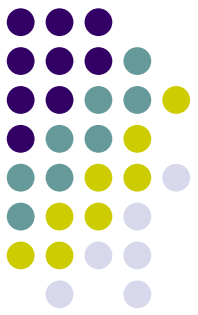
Example *flat* gprof output

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
57.50	0.23	0.23				main
17.50	0.30	0.07	3	23.33	23.33	count4()
7.50	0.33	0.03	2	15.00	38.33	count3()
7.50	0.36	0.03	1	30.00	30.00	count1()
5.00	0.38	0.02	1	20.00	20.00	count()
5.00	0.40	0.02	1	20.00	58.33	count2()

% time spent out of total

Cumulative and self seconds spent

of calls, msec per call (self and total)



Example *flat* gprof output

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
57.50	0.23	0.23				main
17.50	0.30	0.07	3	23.33	23.33	count4()
7.50	0.33	0.03	2	15.00	38.33	count3()
7.50	0.36	0.03	1	30.00	30.00	count1()
5.00	0.38	0.02	1	20.00	20.00	count()
5.00	0.40	0.02	1	20.00	58.33	count2()


% time spent out of total

Cumulative and self seconds spent

of calls, msec per call (self and total)

Example *flat* cProfile output

[stackoverflow.com/questions/582336/how-can-you-profile-a-python-script]



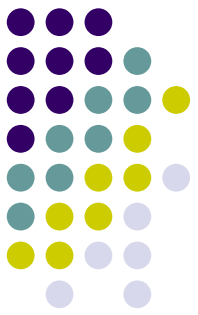
1007 function calls in 0.061 CPU seconds

ncalls	totttime	percall	cumtime	percall	file:line#(function)
1	0.000	0.000	0.061	0.061	<string>:1 (<module>)
1000	0.051	0.000	0.051	0.000	euler048.py:2 (<lambda>)
1	0.005	0.005	0.061	0.061	euler048.py:2 (<module>)
1	0.000	0.000	0.061	0.061	{execfile}
1	0.002	0.002	0.053	0.053	{map}
1	0.000	0.000	0.000	0.000	{method 'disable' ...}
objects}					
1	0.000	0.000	0.000	0.000	{range}
1	0.003	0.003	0.003	0.003	{sum}

See: [Python Profiling \(Amjith Ramanujam on youtube\)](#)

- Shows also GUI tools and more tools, how to use, etc.
- Remember others exist

Code Profiling Example




- Task: Count n -gram frequencies in text document
 - Sorted list of words (1-gram) from most frequent to least
 - Also pairs (2-gram)
- Information retrieval, natural language processing

- Data Set
 - Collected works of Shakespeare
 - 946,596 total words, 26,596 unique

Shakespeare's most frequent words

29,801	the
27,529	and
21,029	I
20,957	to
18,514	of
15,370	a
14,010	you
12,936	my
11,722	in
11,519	that

Code Profiling



Augment Executable Program with Timing Functions

Computes (approximate) amount of time spent in each function

Also maintains counter for each function indicating number of times called

Using

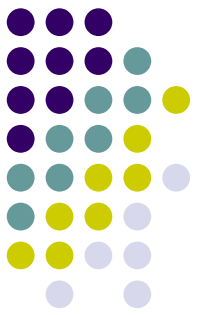
```
gcc -O2 -pg prog. -o prog
./prog
```

Executes in normal fashion, but also generates file
`gmon.out`

```
gprof prog
```

Generates profile information based on `gmon.out`

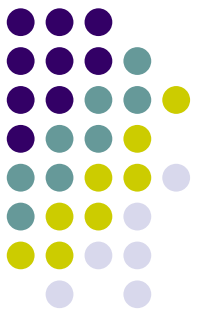
Implementation



- Steps
 - Convert strings to lowercase
 - Apply hash function
 - Read words and insert into hash table
 - Mostly list operations
 - Maintain counter for each unique word
 - Sort results
- Initial implementation
 - Sort: insertion sort(?)
 - List insertion: at end of list (via recursive call)
 - Hash: sum of characters in word, modulo (%) table size
 - Convert to lower:

```
for (i = 0; i < strlen(s); i++)  
    if (s[i] >= 'A' && s[i] <= 'Z')    s[i] -= ('A' - 'a');
```

Profiling Results



%	cumulative	self		self	total	
time	seconds	seconds	calls	ms/call	ms/call	name
86.60	8.21	8.21	1	8210.00	8210.00	sort_words
5.80	8.76	0.55	946596	0.00	0.00	lower1
4.75	9.21	0.45	946596	0.00	0.00	find_ele_rec
1.27	9.33	0.12	946596	0.00	0.00	h_add

Call Statistics



Number of calls and cumulative time for each function

Performance Limiter

Using inefficient sorting algorithm



Single call uses 87% of CPU time

Code Optimizations





- First step:
 - Use more efficient sorting function
 - Library function `qsort`

Code Optimizations



- First step:
 - Use more efficient sorting function
 - Library function `qsort`
- Now list operations main issue



Further Optimizations



Improve list operations:

- Iteration (loop) instead of recursion
- Iter First: insert elements into first place of linked list
 - Causes code to slow down
- Iter Last: insert elements at end of list
 - Much better. Why?

Further Optimizations





Improve list operations:

- Iteration (loop) instead of recursion
- Iter First: insert elements into first place of linked list
 - Causes code to slow down
- Iter Last: insert elements at end of list
 - Much better. Why?

Tend to place most common words at front of list



Further Optimizations



Hashing

- Big table: Increase number of hash buckets
- Better hash: Use more sophisticated hash function

Further Optimizations




Lower

- Move strlen out of loop:

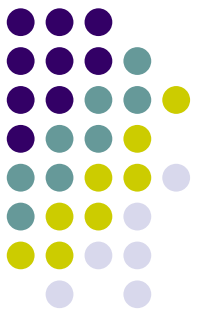
```
len = strlen(s)
for (i = 0; i < len; i++)
    if (s[i] >= 'A' && s[i] <= 'Z')    s[i] -= ('A' - 'a');
```

Implementation matters (even on fast machines)



- For 1-gram (single words), from 9.3 to 0.5 (X ~20 speedup)
 - This was on an old 32 bit machine
 - Does it really matter?
- On i7, 16GB, bought early in 2020:
 - 1-gram speedup: from 0.26 to 0.02 (~13)
 - 2-gram speedup: from 238.14 to 0.15 (~1587.6)
- Profiling is standard, common practice
 - Among professionals
 - Very powerful tools

Observations




Benefits

- Helps identify performance bottlenecks
- Especially within complex system with many components

Limitations

- Only shows performance for data tested
 - e.g., linear lower did not show big gain, since words are short
 - Quadratic inefficiency could remain lurking in code
- Timing mechanism fairly crude
 - Only works for programs that run for > 3 seconds

Do some self-studying!



Both gprof and valgrind are extremely powerful tools

Many options, many features

Take the time to study them

They can save you many many hours of work


www.valgrind.org

“man gprof”, “man valgrind”

Google for tutorials

....

Things NOT to do



Reduce number of lines in code, write unreadable code

There is no direct relation between # lines and execution time

Compare:

```
for (int i=0; i<5; i++) a[i] = i;
```

to:

```
a[0]=0; a[1]=1; a[2]=2; a[3]=3; a[4]=4
```

Optimize as you go

Invest effort without thought

Often makes code difficult to maintain, re-use

Guess at where to spend effort

Instead: Measure! **Profile!** Don't guess.

